StarPU Internal Handbook

for StarPU 1.3.9

This manual documents the internal usage of StarPU version 1.3.9. Its contents was last updated on 21 October 2021.

Contents

Chapter 1

Introduction

1.1   Motivation

Chapter 2

StarPU Core

2.1 StarPU Core Entities

TODO

2.1.1 Overview

Execution entities:

- **worker**: A worker (see Workers, Workers and Scheduling Contexts) entity is a CPU thread created by StarPU to manage one computing unit. The computing unit can be a local CPU core, an accelerator or GPU device, or — on the master side when running in master-slave distributed mode — a remote slave computing node. It is responsible for querying scheduling policies for tasks to execute.

- **sched_context**: A scheduling context (see Scheduling Contexts, Workers and Scheduling Contexts) is a logical set of workers governed by an instance of a scheduling policy. It defines the computing units to which the scheduling policy instance may assign work entities.

- **driver**: A driver is the set of hardware-dependent routines used by a worker to initialize its associated computing unit, execute work entities on it, and finalize the computing unit usage at the end of the session.

Work entities:

- **task**: A task is a high level work request submitted to StarPU by the application, or internally by StarPU itself.

- **job**: A job is a low level view of a work request. It is not exposed to the application. A job structure may be shared among several task structures in the case of a parallel task.

Data entities:

- **data handle**: A data handle is a high-level, application opaque object designating a piece of data currently registered to the StarPU data management layer. Internally, it is a _starpu_data_state structure.

- **data replicate**: A data replicate is a low-level object designating one copy of a piece of data registered to StarPU as a data handle, residing in one memory node managed by StarPU. It is not exposed to the application.

2.1.2 Workers

A **worker** is a CPU thread created by StarPU. Its role is to manage one computing unit. This computing unit can be a local CPU core, in which case, the worker thread manages the actual CPU core to which it is assigned; or it can be a computing device such as a GPU or an accelerator (or even a remote computing node when StarPU is running in distributed master-slave mode.) When a worker manages a computing device, the CPU core to which the worker's thread is by default exclusively assigned to the device management work and does not participate to computation.

2.1.2.1 States

Scheduling operations related state

While a worker is conducting a scheduling operations, e.g. the worker is in the process of selecting a new task to execute, flag state_sched_op_pending is set to `!0`, otherwise it is set to `0`.

While state_sched_op_pending is !0, the following exhaustive list of operations on that workers are restricted in the stated way:

- adding the worker to a context is not allowed;

- removing the worker from a context is not allowed;

- adding the worker to a parallel task team is not allowed;

- removing the worker from a parallel task team is not allowed;

- querying state information about the worker is only allowed while `state_relax_refcnt > 0`;

  - in particular, querying whether the worker is blocked on a parallel team entry is only allowed while `state_relax_refcnt > 0`.

Entering and leaving the state_sched_op_pending state is done through calls to _starpu_worker_enter_sched_op() and _starpu_worker_leave_sched_op() respectively (see these functions in use in functions _starpu_get_worker←_task() and _starpu_get_multi_worker_task()). These calls ensure that any pending conflicting operation deferred while the worker was in the state_sched_op_pending state is performed in an orderly manner.

Scheduling contexts related states

Flag `state_changing_ctx_notice` is set to `!0` when a thread is about to add the worker to a scheduling context or remove it from a scheduling context, and is currently waiting for a safe window to do so, until the targeted worker is not in a scheduling operation or parallel task operation anymore. This flag set to `!0` will also prevent the targeted worker to attempt a fresh scheduling operation or parallel task operation to avoid starving conditions. However, a scheduling operation that was already in progress before the notice is allowed to complete.

Flag `state_changing_ctx_waiting` is set to `!0` when a scheduling context worker addition or removal involving the targeted worker is about to occur and the worker is currently performing a scheduling operation to tell the targeted worker that the initiator thread is waiting for the scheduling operation to complete and should be woken up upon completion.

Relaxed synchronization related states

Any StarPU worker may participate to scheduling operations, and in this process, may be forced to observe state information from other workers. A StarPU worker thread may therefore be observed by any thread, even other StarPU workers. Since workers may observe each other in any order, it is not possible to rely exclusively on the `sched_mutex` of each worker to protect the observation of worker state flags by other workers, because worker A observing worker B would involve locking workers in (A B) sequence, while worker B observing worker A would involve locking workers in (B A) sequence, leading to lock inversion deadlocks.

In consequence, no thread must hold more than one worker's sched_mutex at any time. Instead, workers implement a relaxed locking scheme based on the `state_relax_refcnt` counter, itself protected by the worker's sched←_mutex. When `state_relax_refcnt > 0`, the targeted worker state flags may be observed, otherwise the thread attempting the observation must repeatedly wait on the targeted worker's `sched_cond` condition until `state_relax_refcnt > 0`.

The relaxed mode, while on, can actually be seen as a transactional consistency model, where concurrent accesses are authorized and potential conflicts are resolved after the fact. When the relaxed mode is off, the consistency model becomes a mutual exclusion model, where the sched_mutex of the worker must be held in order to access or change the worker state.

Parallel tasks related states

When a worker is scheduled to participate to the execution of a parallel task, it must wait for the whole team of workers participating to the execution of this task to be ready. While the worker waits for its teammates, it is not available to run other tasks or perform other operations. Such a waiting operation can therefore not start while conflicting operations such as scheduling operations and scheduling context resizing involving the worker are ongoing. Conversely these operations and other may query weather the worker is blocked on a parallel task entry with starpu_worker_is_blocked_in_parallel().

The starpu_worker_is_blocked_in_parallel() function is allowed to proceed while and only while `state_relax↩`
`_refcnt > 0`. Due to the relaxed worker locking scheme, the `state_blocked_in_parallel` flag of the
targeted worker may change after it has been observed by an observer thread. In consequence, flag `state_↩`
`blocked_in_parallel_observed` of the targeted worker is set to `1` by the observer immediately after the
observation to "taint" the targeted worker. The targeted worker will clear the `state_blocked_in_parallel↩`
`_observed` flag tainting and defer the processing of parallel task related requests until a full scheduling operation
shot completes without the `state_blocked_in_parallel_observed` flag being tainted again. The pur-
pose of this tainting flag is to prevent parallel task operations to be started immediately after the observation of a
transient scheduling state.

Worker's management of parallel tasks is governed by the following set of state flags and counters:

- `state_blocked_in_parallel`: set to `!0` while the worker is currently blocked on a parallel task;

- `state_blocked_in_parallel_observed`: set to `!0` to taint the worker when a thread has ob-
  served the state_blocked_in_parallel flag of this worker while its `state_relax_refcnt` state counter
  was >0. Any pending request to add or remove the worker from a parallel task team will be deferred until a
  whole scheduling operation shot completes without being tainted again.

- `state_block_in_parallel_req`: set to `!0` when a thread is waiting on a request for the worker to
  be added to a parallel task team. Must be protected by the worker's `sched_mutex`.

- `state_block_in_parallel_ack`: set to `!0` by the worker when acknowledging a request for being
  added to a parallel task team. Must be protected by the worker's `sched_mutex`.

- `state_unblock_in_parallel_req`: set to `!0` when a thread is waiting on a request for the worker
  to be removed from a parallel task team. Must be protected by the worker's `sched_mutex`.

- `state_unblock_in_parallel_ack`: set to `!0` by the worker when acknowledging a request for
  being removed from a parallel task team. Must be protected by the worker's `sched_mutex`.

- `block_in_parallel_ref_count`: counts the number of consecutive pending requests to enter par-
  allel task teams. Only the first of a train of requests for entering parallel task teams triggers the transition of
  the `state_block_in_parallel_req` flag from `0` to `1`. Only the last of a train of requests to leave a
  parallel task team triggers the transition of flag `state_unblock_in_parallel_req` from `0` to `1`. Must
  be protected by the worker's `sched_mutex`.

2.1.2.2 Operations

Entry point

All the operations of a worker are handled in an iterative fashion, either by the application code on a thread launched
by the application, or automatically by StarPU on a device-dependent CPU thread launched by StarPU. Whether a
worker's operation cycle is managed automatically or not is controlled per session by the field `not_launched↩`
`_drivers` of the `starpu_conf` struct, and is decided in _starpu_launch_drivers() function.

When managed automatically, cycles of operations for a worker are handled by the corresponding driver specific
`_starpu_<DRV>_worker()` function, where `DRV` is a driver name such as cpu (`_starpu_cpu_worker`)
or cuda (`_starpu_cuda_worker`), for instance. Otherwise, the application must supply a thread which will
repeatedly call starpu_driver_run_once() for the corresponding worker.

In both cases, control is then transferred to _starpu_cpu_driver_run_once() (or the corresponding driver specific
func). The cycle of operations typically includes, at least, the following operations:

- **task scheduling**

- **parallel task team build-up**

- **task input processing**

- **data transfer processing**

- **task execution**

When the worker cycles are handled by StarPU automatically, the iterative operation processing ends when the
`running` field of `_starpu_config` becomes false. This field should not be read directly, instead it should be
read through the [_starpu_machine_is_running()](#) function.

Task scheduling

If the worker does not yet have a queued task, it calls _starpu_get_worker_task() to try and obtain a task. This may involve scheduling operations such as stealing a queued but not yet executed task from another worker. The operation may not necessarily succeed if no tasks are ready and/or suitable to run on the worker's computing unit.

Parallel task team build-up

If the worker has a task ready to run and the corresponding job has a size >1, then the task is a parallel job and the worker must synchronize with the other workers participating to the parallel execution of the job to assign a unique rank for each worker. The synchronization is done through the job's `sync_mutex` mutex.

Task input processing

Before the task can be executed, its input data must be made available on a memory node reachable by the worker's computing unit. To do so, the worker calls _starpu_fetch_task_input()

Data transfer processing

The worker makes pending data transfers (involving memory node(s) that it is driving) progress, with a call to __←
starpu_datawizard_progress(),

Task execution

Once the worker has a pending task assigned and the input data for that task are available in the memory node reachable by the worker's computing unit, the worker calls _starpu_cpu_driver_execute_task() (or the corresponding driver specific function) to proceed to the execution of the task.

2.1.3 Scheduling Contexts

A scheduling context is a logical set of workers governed by an instance of a scheduling policy. Tasks submitted to a given scheduling context are confined to the computing units governed by the workers belonging to this scheduling context at the time they get scheduled.

A scheduling context is identified by an unsigned integer identifier between `0` and `STARPU_NMAX_SCHED_CT←`
`XS - 1`. The `STARPU_NMAX_SCHED_CTXS` identifier value is reserved to indicated an unallocated, invalid or deleted scheduling context.

Accesses to the scheduling context structure are governed by a multiple-readers/single-writer lock (`rwlock` field). Changes to the structure contents, additions or removals of workers, statistics updates, all must be done with proper exclusive write access.

2.1.4 Workers and Scheduling Contexts

A worker can be assigned to one or more **scheduling contexts**. It exclusively receives tasks submitted to the scheduling context(s) it is currently assigned at the time such tasks are scheduled. A worker may add itself to or remove itself from a scheduling context.

Locking and synchronization rules between workers and scheduling contexts

A thread currently holding a worker sched_mutex must not attempt to acquire a scheduling context rwlock, neither for writing nor for reading. Such an attempt constitutes a lock inversion and may result in a deadlock.

A worker currently in a scheduling operation must enter the relaxed state before attempting to acquire a scheduling context rwlock, either for reading or for writing.

When the set of workers assigned to a scheduling context is about to be modified, all the workers in the union between the workers belonging to the scheduling context before the change and the workers expected to belong to the scheduling context after the change must be notified using the notify_workers_about_changing_ctx_pending() function prior to the update. After the update, all the workers in that same union must be notified for the update completion with a call to notify_workers_about_changing_ctx_done().

The function notify_workers_about_changing_ctx_pending() places every worker passed in argument in a state compatible with changing the scheduling context assignment of that worker, possibly blocking until that worker leaves incompatible states such as a pending scheduling operation. If the caller of `notify_workers_about←`
`_changing_ctx_pending()` is itself a worker included in the set of workers passed in argument, it does not notify itself, with the assumption that the worker is already calling `notify_workers_about_changing←`
`_ctx_pending()` from a state compatible with a scheduling context assignment update. Once a worker has

been notified about a scheduling context change pending, it cannot proceed with incompatible operations such as a scheduling operation until it receives a notification that the context update operation is complete.

2.1.5 Drivers

Each driver defines a set of routines depending on some specific hardware. These routines include hardware discovery/initialization, task execution, device memory management and data transfers.
While most hardware dependent routines are in source files located in the `/src/drivers` subdirectory of the StarPU tree, some can be found elsewhere in the tree such as `src/datawizard/malloc.c` for memory allocation routines or the subdirectories of `src/datawizard/interfaces/` for data transfer routines.
The driver ABI defined in the _starpu_driver_ops structure includes the following operations:

- .init: initialize a driver instance for the calling worker managing a hardware computing unit compatible with this driver.

- .run_once: perform a single driver progress cycle for the calling worker (see Operations).

- .deinit: deinitialize the driver instance for the calling worker

- .run: executes the following sequence automatically: call .init, repeatedly call .run_once until the function _starpu_machine_is_running() returns false, call .deinit.

The source code common to all drivers is shared in `src/drivers/driver_common/driver_↵common.[ch]`. This file includes services such as grabbing a new task to execute on a worker, managing statistics accounting on job startup and completion and updating the worker status

2.1.5.1 Master/Slave Drivers

A subset of the drivers corresponds to drivers managing computing units in master/slave mode, that is, drivers involving a local master instance managing one or more remote slave instances on the targeted device(s). This includes devices such as discrete manycore accelerators (e.g. Intel's Knight Corners board, for instance), or pseudo devices such as a cluster of cpu nodes driver through StarPU's MPI master/slave mode. A driver instance on the master side is named the **source**, while a driver instances on the slave side is named the **sink**.
A significant part of the work realized on the source and sink sides of master/slave drivers is identical among all master/slave drivers, due to the similarities in the software pattern. Therefore, many routines are shared among all these drivers in the `src/drivers/mp_common` subdirectory. In particular, a set of default commands to be used between sources and sinks is defined, assuming the availability of some communication channel between them (see enum _starpu_mp_command)
TODO

2.1.6 Tasks and Jobs

TODO

2.1.7 Data

TODO

Chapter 3

Module Index

3.1 Modules

Here is a list of all modules:

Chapter 4

Module Documentation

4.1 Workers

Data Structures

- struct _starpu_worker
- struct _starpu_combined_worker
- struct _starpu_worker_set
- struct _starpu_machine_topology
- struct _starpu_machine_config
- struct _starpu_machine_config.bindid_workers

Macros

- #define **STARPU_MAX_PIPELINE**
- #define **starpu_worker_get_count**
- #define **starpu_worker_get_id**
- #define **_starpu_worker_get_id_check**(f, l)
- #define **starpu_worker_relax_on**
- #define **starpu_worker_relax_off**
- #define **starpu_worker_get_relax_state**

Enumerations

- enum **initialization** { **UNINITIALIZED**, **CHANGING**, **INITIALIZED** }

Functions

- void _starpu_set_argc_argv (int ∗argc, char ∗∗∗argv)
- int ∗ **_starpu_get_argc** ()
- char ∗∗∗ **_starpu_get_argv** ()
- void _starpu_conf_check_environment (struct starpu_conf ∗conf)
- void _starpu_may_pause (void)
- static unsigned _starpu_machine_is_running (void)
- void _starpu_worker_init (struct _starpu_worker ∗workerarg, struct _starpu_machine_config ∗pconfig)
- uint32_t _starpu_worker_exists (struct starpu_task ∗)
- uint32_t _starpu_can_submit_cuda_task (void)
- uint32_t _starpu_can_submit_cpu_task (void)
- uint32_t _starpu_can_submit_opencl_task (void)
- unsigned _starpu_worker_can_block (unsigned memnode, struct _starpu_worker ∗worker)
- void _starpu_block_worker (int workerid, starpu_pthread_cond_t ∗cond, starpu_pthread_mutex_t ∗mutex)
- void _starpu_driver_start (struct _starpu_worker ∗worker, unsigned fut_key, unsigned sync)
- void _starpu_worker_start (struct _starpu_worker ∗worker, unsigned fut_key, unsigned sync)

- static unsigned **_starpu_worker_get_count** (void)
- static void _starpu_set_local_worker_key (struct _starpu_worker ∗worker)
- static struct _starpu_worker ∗ _starpu_get_local_worker_key (void)
- static void _starpu_set_local_worker_set_key (struct _starpu_worker_set ∗worker)
- static struct _starpu_worker_set ∗ _starpu_get_local_worker_set_key (void)
- static struct _starpu_worker ∗ _starpu_get_worker_struct (unsigned id)
- static struct _starpu_sched_ctx ∗ _starpu_get_sched_ctx_struct (unsigned id)
- struct _starpu_combined_worker ∗ **_starpu_get_combined_worker_struct** (unsigned id)
- static struct _starpu_machine_config ∗ _starpu_get_machine_config (void)
- static int _starpu_get_disable_kernels (void)
- static enum _starpu_worker_status _starpu_worker_get_status (int workerid)
- static void _starpu_worker_set_status (int workerid, enum _starpu_worker_status status)
- static struct _starpu_sched_ctx ∗ _starpu_get_initial_sched_ctx (void)
- int **starpu_worker_get_nids_by_type** (enum starpu_worker_archtype type, int ∗workerids, int maxsize)
- int starpu_worker_get_nids_ctx_free_by_type (enum starpu_worker_archtype type, int ∗workerids, int max-size)
- static unsigned **_starpu_worker_mutex_is_sched_mutex** (int workerid, starpu_pthread_mutex_t ∗mutex)
- static int **_starpu_worker_get_nsched_ctxs** (int workerid)
- static unsigned _starpu_get_nsched_ctxs (void)
- static int _starpu_worker_get_id (void)
- static unsigned __starpu_worker_get_id_check (const char ∗f, int l)
- enum starpu_node_kind **_starpu_worker_get_node_kind** (enum starpu_worker_archtype type)
- void **_starpu_worker_set_stream_ctx** (unsigned workerid, struct _starpu_sched_ctx ∗sched_ctx)
- struct _starpu_sched_ctx ∗ **_starpu_worker_get_ctx_stream** (unsigned stream_workerid)
- static void _starpu_worker_request_blocking_in_parallel (struct _starpu_worker ∗const worker)
- static void _starpu_worker_request_unblocking_in_parallel (struct _starpu_worker ∗const worker)
- static void _starpu_worker_process_block_in_parallel_requests (struct _starpu_worker ∗const worker)
- static void _starpu_worker_enter_sched_op (struct _starpu_worker ∗const worker)
- void _starpu_worker_apply_deferred_ctx_changes (void)
- static void **_starpu_worker_leave_sched_op** (struct _starpu_worker ∗const worker)
- static int **_starpu_worker_sched_op_pending** (void)
- static void _starpu_worker_enter_changing_ctx_op (struct _starpu_worker ∗const worker)
- static void _starpu_worker_leave_changing_ctx_op (struct _starpu_worker ∗const worker)
- static void _starpu_worker_relax_on (void)
- static void _starpu_worker_relax_on_locked (struct _starpu_worker ∗worker)
- static void **_starpu_worker_relax_off** (void)
- static void **_starpu_worker_relax_off_locked** (void)
- static int **_starpu_worker_get_relax_state** (void)
- static void _starpu_worker_lock (int workerid)
- static int **_starpu_worker_trylock** (int workerid)
- static void **_starpu_worker_unlock** (int workerid)
- static void **_starpu_worker_lock_self** (void)
- static void **_starpu_worker_unlock_self** (void)
- static int **_starpu_wake_worker_relax** (int workerid)
- int **starpu_wake_worker_relax_light** (int workerid)
- void _starpu_worker_refuse_task (struct _starpu_worker ∗worker, struct starpu_task ∗task)

Variables

- int **_starpu_worker_parallel_blocks**
- struct _starpu_machine_config _starpu_config **STARPU_ATTRIBUTE_INTERNAL**

4.1.1 Detailed Description

4.1.2 Data Structure Documentation

4.1.2.1 struct _starpu_worker

This is initialized by _starpu_worker_init()

Data Fields

struct _starpu_machine_config ∗	config	
starpu_pthread_mutex_t	mutex	
enum starpu_worker_archtype	arch	what is the type of worker ?
uint32_t	worker_mask	what is the type of worker ?
struct starpu_perfmodel_arch	perf_arch	in case there are different models of the same arch
starpu_pthread_t	worker_thread	the thread which runs the worker
unsigned	devid	which cpu/gpu/etc is controlled by the worker ?
unsigned	subworkerid	which sub-worker this one is for the cpu/gpu
int	bindid	which cpu is the driver bound to ? (logical index)
int	workerid	uniquely identify the worker among all processing units types
int	combined_workerid	combined worker currently using this worker
int	current_rank	current rank in case the worker is used in a parallel fashion
int	worker_size	size of the worker in case we use a combined worker
starpu_pthread_cond_t	started_cond	indicate when the worker is ready
starpu_pthread_cond_t	ready_cond	indicate when the worker is ready
unsigned	memory_node	which memory node is the worker associated with ?
unsigned	numa_memory_node	which numa memory node is the worker associated with? (logical index)
starpu_pthread_cond_t	sched_cond	condition variable used for passive waiting operations on worker STARPU_PTHREAD_COND_B←↩ ROADCAST must be used instead of STARPU_PTHREAD_COND←↩ _SIGNAL, since the condition is shared for multiple purpose
starpu_pthread_mutex_t	sched_mutex	mutex protecting sched_cond
unsigned	state_relax_refcnt	mark scheduling sections where other workers can safely access the worker state
unsigned	state_sched_op_pending	a task pop is ongoing even though sched_mutex may temporarily be unlocked
unsigned	state_changing_ctx_waiting	a thread is waiting for operations such as pop to complete before acquiring sched_mutex and modifying the worker ctx
unsigned	state_changing_ctx_notice	the worker ctx is about to change or being changed, wait for flag to be cleared before starting new scheduling operations
unsigned	state_blocked_in_parallel	worker is currently blocked on a parallel section

Data Fields

unsigned	state_blocked_in_parallel_observed	the blocked state of the worker has been observed by another worker during a relaxed section
unsigned	state_block_in_parallel_req	a request for state transition from unblocked to blocked is pending
unsigned	state_block_in_parallel_ack	a block request has been honored
unsigned	state_unblock_in_parallel_req	a request for state transition from blocked to unblocked is pending
unsigned	state_unblock_in_parallel_ack	an unblock request has been honored
unsigned	block_in_parallel_ref_count	cumulative blocking depth • =0 worker unblocked • >0 worker blocked • transition from 0 to 1 triggers a block_req • transition from 1 to 0 triggers a unblock_req
starpu_pthread_t	thread_changing_ctx	thread currently changing a sched_ctx containing the worker
struct _starpu_ctx_change_list	ctx_change_list	list of deferred context changes when the current thread is a worker, _and_ this worker is in a scheduling operation, new ctx changes are queued to this list for subsequent processing once worker completes the ongoing scheduling operation
struct starpu_task_list	local_tasks	this queue contains tasks that have been explicitly submitted to that queue
struct starpu_task **	local_ordered_tasks	this queue contains tasks that have been explicitly submitted to that queue with an explicit order
unsigned	local_ordered_tasks_size	this records the size of local_ordered_tasks
unsigned	current_ordered_task	this records the index (within local_ordered_tasks) of the next ordered task to be executed
unsigned	current_ordered_task_order	this records the order of the next ordered task to be executed
struct starpu_task *	current_task	task currently executed by this worker (non-pipelined version)
struct starpu_task *	current_tasks[STARPU_MAX_PIPELINE]	tasks currently executed by this worker (pipelined version)
starpu_pthread_wait_t	wait	
struct timespec	cl_start	Codelet start time of the task currently running
struct timespec	cl_end	Codelet end time of the last task running
unsigned char	first_task	Index of first task in the pipeline

Data Fields

unsigned char	ntasks	number of tasks in the pipeline
unsigned char	pipeline_length	number of tasks to be put in the pipeline
unsigned char	pipeline_stuck	whether a task prevents us from pipelining
struct _starpu_worker_set ∗	set	in case this worker belongs to a set
unsigned	worker_is_running	
unsigned	worker_is_initialized	
enum _starpu_worker_status	status	what is the worker doing now ? (eg. CALLBACK)
unsigned	state_keep_awake	!0 if a task has been pushed to the worker and the task has not yet been seen by the worker, the worker should no go to sleep before processing this task
char	name[128]	
char	short_name[32]	
unsigned	run_by_starpu	Is this run by StarPU or directly by the application ?
struct _starpu_driver_ops ∗	driver_ops	
struct _starpu_sched_ctx_list ∗	sched_ctx_list	
int	tmp_sched_ctx	
unsigned	nsched_ctxs	the no of contexts a worker belongs to
struct _starpu_barrier_counter	tasks_barrier	wait for the tasks submitted
unsigned	has_prev_init	had already been inited in another ctx
unsigned	removed_from_ctx[STARPU_NMAX_SCHED_CTXS+1]	
unsigned	spinning_backoff	number of cycles to pause when spinning
unsigned	nb_buffers_transferred	number of piece of data already send to worker
unsigned	nb_buffers_totransfer	number of piece of data already send to worker
struct starpu_task ∗	task_transferring	The buffers of this task are being sent
unsigned	shares_tasks_lists[STARPU_NMAX_SCHED_CTXS+1]	indicate whether the workers shares tasks lists with other workers in this case when removing him from a context it disapears instantly
unsigned	poped_in_ctx[STARPU_NMAX_SCHED_CTXS+1]	boolean to chose the next ctx a worker will pop into
unsigned	reverse_phase[2]	boolean indicating at which moment we checked all ctxs and change phase for the booleab poped_in_ctx one for each of the 2 priorities
unsigned	pop_ctx_priority	indicate which priority of ctx is currently active: the values are 0 or 1

Data Fields

unsigned	is_slave_somewhere	bool to indicate if the worker is slave in a ctx
struct _starpu_sched_ctx *	stream_ctx	
hwloc_bitmap_t	hwloc_cpu_set	
hwloc_obj_t	hwloc_obj	
char	padding[STARPU_CACHELINE_SIZE]	Keep this last, to make sure to separate worker data in separate cache lines.

4.1.2.2 struct _starpu_combined_worker

Data Fields

struct starpu_perfmodel_arch	perf_arch	in case there are different models of the same arch
uint32_t	worker_mask	what is the type of workers ?
int	worker_size	
unsigned	memory_node	which memory node is associated that worker to ?
int	combined_workerid[STARPU_NMAXWORKERS]	
hwloc_bitmap_t	hwloc_cpu_set	
char	padding[STARPU_CACHELINE_SIZE]	Keep this last, to make sure to separate worker data in separate cache lines.

4.1.2.3 struct _starpu_worker_set

in case a single CPU worker may control multiple accelerators

Data Fields

starpu_pthread_mutex_t	mutex	
starpu_pthread_t	worker_thread	the thread which runs the worker
unsigned	nworkers	
unsigned	started	Only one thread for the whole set
void *	retval	
struct _starpu_worker *	workers	
starpu_pthread_cond_t	ready_cond	indicate when the set is ready
unsigned	set_is_initialized	

4.1.2.4 struct _starpu_machine_topology

Data Fields

unsigned	nworkers	Total number of workers.
unsigned	ncombinedworkers	Total number of combined workers.
unsigned	nsched_ctxs	
hwloc_topology_t	hwtopology	Topology as detected by hwloc.
struct starpu_tree *	tree	custom hwloc tree

Data Fields

unsigned	nhwcpus		Total number of CPU cores, as detected by the topology code. May be different from the actual number of CPU workers.
unsigned	nhwpus		Total number of PUs (i.e. threads), as detected by the topology code. May be different from the actual number of PU workers.
unsigned	nhwcudagpus		Total number of CUDA devices, as detected. May be different from the actual number of CUDA workers.
unsigned	nhwopenclgpus		Total number of OpenCL devices, as detected. May be different from the actual number of OpenCL workers.
unsigned	nhwmpi		Total number of MPI nodes, as detected. May be different from the actual number of node workers.
unsigned	ncpus		Actual number of CPU workers used by StarPU.
unsigned	ncudagpus		Actual number of CUDA GPUs used by StarPU.
unsigned	nworkerpercuda		
int	cuda_th_per_stream		
int	cuda_th_per_dev		
unsigned	nopenclgpus		Actual number of OpenCL workers used by StarPU.
unsigned	nmpidevices		Actual number of MPI workers used by StarPU.
unsigned	nhwmpidevices		
unsigned	nhwmpicores[STARPU_MAXMPIDEVS]		Each MPI node has its set of cores.
unsigned	nmpicores[STARPU_MAXMPIDEVS]		
unsigned	nhwmicdevices		Topology of MP nodes (MIC) as well as necessary objects to communicate with them.
unsigned	nmicdevices		
unsigned	nhwmiccores[STARPU_MAXMICDEVS]		Each MIC node has its set of cores.
unsigned	nmiccores[STARPU_MAXMICDEVS]		
unsigned	workers_bindid[STARPU_NMAXWORKERS]		Indicates the successive logical PU identifier that should be used to bind the workers. It is either filled according to the user's explicit parameters (from starpu_conf) or according to the STARPU_WORKERS_CPUID env. variable. Otherwise, a round-robin policy is used to distributed the workers over the cores.
unsigned	workers_cuda_gpuid[STARPU_NMAXWORKERS]		Indicates the successive CUDA identifier that should be used by the CUDA driver. It is either filled according to the user's explicit parameters (from starpu_conf) or according to the STARPU_WORKERS_CUDAID env. variable. Otherwise, they are taken in ID order.

Data Fields

unsigned	workers_opencl_gpuid[STARPU_NMAXWORKERS]	Indicates the successive OpenCL identifier that should be used by the OpenCL driver. It is either filled according to the user's explicit parameters (from starpu_conf) or according to the STARPU_WORKERS_OPENCLID env. variable. Otherwise, they are taken in ID order.
unsigned	workers_mpi_ms_deviceid[STARPU_NMAXWORKERS]	unsigned workers_mic_↩ deviceid[STARPU_NMAXWORKERS];

4.1.2.5 struct _starpu_machine_config

Data Fields

struct _starpu_machine_topology	topology	
int	cpu_depth	
int	pu_depth	
int	current_bindid	Where to bind next worker ?
char	currently_bound[STARPU_NMAXWORKERS]	
char	currently_shared[STARPU_NMAXWORKERS]	
int	current_cuda_gpuid	Which GPU(s) do we use for CUDA ?
int	current_opencl_gpuid	Which GPU(s) do we use for OpenCL ?
int	current_mic_deviceid	Which MIC do we use?
int	current_mpi_deviceid	Which MPI do we use?
int	cpus_nodeid	Memory node for cpus, if only one
int	cuda_nodeid	Memory node for CUDA, if only one
int	opencl_nodeid	Memory node for OpenCL, if only one
int	mic_nodeid	Memory node for MIC, if only one
int	mpi_nodeid	Memory node for MPI, if only one
char	padding1[STARPU_CACHELINE_SIZE]	Separate out previous variables from per-worker data.
struct _starpu_worker	workers[STARPU_NMAXWORKERS]	Basic workers : each of this worker is running its own driver and can be combined with other basic workers.
struct _starpu_combined_worker	combined_workers[STARPU_NMAX_COMBINEDWORKERS]	Combined workers : these worker are a combination of basic workers that can run parallel tasks together.
starpu_pthread_mutex_t	submitted_mutex	
char	padding2[STARPU_CACHELINE_SIZE]	Separate out previous mutex from the rest of the data.
struct _starpu_machine_config	bindid_workers	Translation table from bindid to worker IDs
unsigned	nbindid	size of bindid_workers

Data Fields

	uint32_t	worker_mask	This bitmask indicates which kinds of worker are available. For instance it is possible to test if there is a CUDA worker with the result of (worker_mask & STARPU_CUDA).
struct starpu_conf	conf	either the user given configuration passed to starpu_init or a default configuration
unsigned	running	this flag is set until the runtime is stopped
int	disable_kernels	
int	pause_depth	Number of calls to starpu_pause() - calls to starpu_resume(). When >0, StarPU should pause.
struct _starpu_sched_ctx	sched_ctxs[STARPU_NMAX_SCHED_CTXS+1]	The sched ctx of the current instance of starpu
unsigned	submitting	this flag is set until the application is finished submitting tasks
int	watchdog_ok	

4.1.2.6   struct _starpu_machine_config.bindid_workers

Translation table from bindid to worker IDs

Data Fields

int *	workerids	
unsigned	nworkers	size of workerids

4.1.3   Function Documentation

4.1.3.1   _starpu_set_argc_argv()

```
void _starpu_set_argc_argv (
          int * argc,
          char *** argv )
```
Three functions to manage argv, argc

4.1.3.2   _starpu_conf_check_environment()

```
void _starpu_conf_check_environment (
          struct starpu_conf * conf )
```
Fill conf with environment variables

4.1.3.3   _starpu_may_pause()

```
void _starpu_may_pause (
          void  )
```
Called by the driver when it is ready to pause

4.1.3.4   _starpu_machine_is_running()

```
static unsigned _starpu_machine_is_running (
            void  )  [inline], [static]
```
Has starpu_shutdown already been called ?

4.1.3.5   _starpu_worker_init()

```
void _starpu_worker_init (
            struct _starpu_worker * workerarg,
            struct _starpu_machine_config * pconfig )
```
initialise a worker

4.1.3.6   _starpu_worker_exists()

```
uint32_t _starpu_worker_exists (
            struct starpu_task *  )
```
Check if there is a worker that may execute the task.

4.1.3.7   _starpu_can_submit_cuda_task()

```
uint32_t _starpu_can_submit_cuda_task (
            void  )
```
Is there a worker that can execute CUDA code ?

4.1.3.8   _starpu_can_submit_cpu_task()

```
uint32_t _starpu_can_submit_cpu_task (
            void  )
```
Is there a worker that can execute CPU code ?

4.1.3.9   _starpu_can_submit_opencl_task()

```
uint32_t _starpu_can_submit_opencl_task (
            void  )
```
Is there a worker that can execute OpenCL code ?

4.1.3.10   _starpu_worker_can_block()

```
unsigned _starpu_worker_can_block (
            unsigned memnode,
            struct _starpu_worker * worker )
```
Check whether there is anything that the worker should do instead of sleeping (waiting on something to happen).

4.1.3.11   _starpu_block_worker()

```
void _starpu_block_worker (
            int workerid,
            starpu_pthread_cond_t * cond,
            starpu_pthread_mutex_t * mutex )
```
This function must be called to block a worker. It puts the worker in a sleeping state until there is some event that forces the worker to wake up.

4.1.3.12   _starpu_driver_start()

```
void _starpu_driver_start (
            struct _starpu_worker * worker,
            unsigned fut_key,
            unsigned sync )
```
This function initializes the current driver for the given worker

4.1.3.13   _starpu_worker_start()

```
void _starpu_worker_start (
            struct _starpu_worker * worker,
            unsigned fut_key,
            unsigned sync )
```
This function initializes the current thread for the given worker

4.1.3.14   _starpu_set_local_worker_key()

```
static void _starpu_set_local_worker_key (
            struct _starpu_worker * worker ) [inline], [static]
```
The _starpu_worker structure describes all the state of a StarPU worker. This function sets the pthread key which stores a pointer to this structure.

4.1.3.15   _starpu_get_local_worker_key()

```
static struct _starpu_worker* _starpu_get_local_worker_key (
            void ) [static]
```
Returns the _starpu_worker structure that describes the state of the current worker.

4.1.3.16   _starpu_set_local_worker_set_key()

```
static void _starpu_set_local_worker_set_key (
            struct _starpu_worker_set * worker ) [inline], [static]
```
The _starpu_worker_set structure describes all the state of a StarPU worker_set. This function sets the pthread key which stores a pointer to this structure.

4.1.3.17   _starpu_get_local_worker_set_key()

```
static struct _starpu_worker_set* _starpu_get_local_worker_set_key (
            void ) [static]
```
Returns the _starpu_worker_set structure that describes the state of the current worker_set.

4.1.3.18   _starpu_get_worker_struct()

```
static struct _starpu_worker* _starpu_get_worker_struct (
            unsigned id ) [static]
```
Returns the _starpu_worker structure that describes the state of the specified worker.

4.1.3.19   _starpu_get_sched_ctx_struct()

```
static struct _starpu_sched_ctx* _starpu_get_sched_ctx_struct (
            unsigned id ) [static]
```
Returns the starpu_sched_ctx structure that describes the state of the specified ctx

4.1.3.20   _starpu_get_machine_config()

```
static struct _starpu_machine_config* _starpu_get_machine_config (
            void ) [static]
```
Returns the structure that describes the overall machine configuration (eg. all workers and topology).

4.1.3.21   _starpu_get_disable_kernels()

```
static int _starpu_get_disable_kernels (
            void ) [inline], [static]
```
Return whether kernels should be run ($<=0$) or not (>0)

4.1.3.22 _starpu_worker_get_status()

```
static enum _starpu_worker_status _starpu_worker_get_status (
            int workerid ) [inline], [static]
```
Retrieve the status which indicates what the worker is currently doing.

4.1.3.23 _starpu_worker_set_status()

```
static void _starpu_worker_set_status (
            int workerid,
            enum _starpu_worker_status status ) [inline], [static]
```
Change the status of the worker which indicates what the worker is currently doing (eg. executing a callback).

4.1.3.24 _starpu_get_initial_sched_ctx()

```
static struct _starpu_sched_ctx* _starpu_get_initial_sched_ctx (
            void ) [static]
```
We keep an initial sched ctx which might be used in case no other ctx is available

4.1.3.25 starpu_worker_get_nids_ctx_free_by_type()

```
int starpu_worker_get_nids_ctx_free_by_type (
            enum starpu_worker_archtype type,
            int * workerids,
            int maxsize )
```
returns workers not belonging to any context, be careful no mutex is used, the list might not be updated

4.1.3.26 _starpu_get_nsched_ctxs()

```
static unsigned _starpu_get_nsched_ctxs (
            void ) [inline], [static]
```
Get the total number of sched_ctxs created till now

4.1.3.27 _starpu_worker_get_id()

```
static int _starpu_worker_get_id (
            void ) [inline], [static]
```
Inlined version when building the core.

4.1.3.28 __starpu_worker_get_id_check()

```
static unsigned __starpu_worker_get_id_check (
            const char * f,
            int l ) [inline], [static]
```
Similar behaviour to starpu_worker_get_id() but fails when called from outside a worker This returns an unsigned object on purpose, so that the caller is sure to get a positive value

4.1.3.29 _starpu_worker_request_blocking_in_parallel()

```
static void _starpu_worker_request_blocking_in_parallel (
            struct _starpu_worker *const worker ) [inline], [static]
```
Send a request to the worker to block, before a parallel task is about to begin.
Must be called with worker's sched_mutex held.

4.1.3.30 _starpu_worker_request_unblocking_in_parallel()

```
static void _starpu_worker_request_unblocking_in_parallel (
            struct _starpu_worker *const worker ) [inline], [static]
```
Send a request to the worker to unblock, after a parallel task is complete.
Must be called with worker's sched_mutex held.

4.1.3.31 _starpu_worker_process_block_in_parallel_requests()

```
static void _starpu_worker_process_block_in_parallel_requests (
            struct _starpu_worker *const worker ) [inline], [static]
```
Called by the the worker to process incoming requests to block or unblock on parallel task boundaries.
Must be called with worker's sched_mutex held.

4.1.3.32 _starpu_worker_enter_sched_op()

```
static void _starpu_worker_enter_sched_op (
            struct _starpu_worker *const worker ) [inline], [static]
```
Mark the beginning of a scheduling operation by the worker. No worker blocking operations on parallel tasks and no scheduling context change operations must be performed on contexts containing the worker, on contexts about to add the worker and on contexts about to remove the worker, while the scheduling operation is in process. The sched mutex of the worker may only be acquired permanently by another thread when no scheduling operation is in process, or when a scheduling operation is in process _and_ worker->state_relax_refcnt!=0. If a scheduling operation is in process _and_ worker->state_relax_refcnt==0, a thread other than the worker must wait on condition worker->sched_cond for worker->state_relax_refcnt!=0 to become true, before acquiring the worker sched mutex permanently.
Must be called with worker's sched_mutex held.

4.1.3.33 _starpu_worker_apply_deferred_ctx_changes()

```
void _starpu_worker_apply_deferred_ctx_changes (
            void )
```
Mark the end of a scheduling operation by the worker.
Must be called with worker's sched_mutex held.

4.1.3.34 _starpu_worker_enter_changing_ctx_op()

```
static void _starpu_worker_enter_changing_ctx_op (
            struct _starpu_worker *const worker ) [inline], [static]
```
Must be called before altering a context related to the worker whether about adding the worker to a context, removing it from a context or modifying the set of workers of a context of which the worker is a member, to mark the beginning of a context change operation. The sched mutex of the worker must be held before calling this function.
Must be called with worker's sched_mutex held.

4.1.3.35 _starpu_worker_leave_changing_ctx_op()

```
static void _starpu_worker_leave_changing_ctx_op (
            struct _starpu_worker *const worker ) [inline], [static]
```
Mark the end of a context change operation.
Must be called with worker's sched_mutex held.

4.1.3.36 _starpu_worker_relax_on()

```
static void _starpu_worker_relax_on (
            void ) [inline], [static]
```
Temporarily allow other worker to access current worker state, when still scheduling, but the scheduling has not yet been made or is already done

4.1.3.37 _starpu_worker_relax_on_locked()

```
static void _starpu_worker_relax_on_locked (
            struct _starpu_worker * worker ) [inline], [static]
```
Same, but with current worker mutex already held

4.1.3.38 _starpu_worker_lock()

```
static void _starpu_worker_lock (
            int workerid ) [inline], [static]
```
lock a worker for observing contents
notes:

- if the observed worker is not in state_relax_refcnt, the function block until the state is reached

4.1.3.39 _starpu_worker_refuse_task()

```
void _starpu_worker_refuse_task (
            struct _starpu_worker * worker,
            struct starpu_task * task )
```
Allow a worker pulling a task it cannot execute to properly refuse it and send it back to the scheduler.

Chapter 5

File Index

5.1 File List

Here is a list of all documented files with brief descriptions:

Chapter 6

StarPU File Documentation

6.1   barrier.h File Reference

```
#include <starpu_thread.h>
```

Data Structures

- struct _starpu_barrier

Functions

- int **_starpu_barrier_init** (struct _starpu_barrier ∗barrier, int count)
- int **_starpu_barrier_destroy** (struct _starpu_barrier ∗barrier)
- int **_starpu_barrier_wait** (struct _starpu_barrier ∗barrier)

6.1.1   Data Structure Documentation

6.1.1.1   struct _starpu_barrier

Data Fields

unsigned	count	
unsigned	reached_start	
unsigned	reached_exit	
double	reached_flops	
starpu_pthread_mutex_t	mutex	
starpu_pthread_mutex_t	mutex_exit	
starpu_pthread_cond_t	cond	

6.2   barrier_counter.h File Reference

```
#include <common/utils.h>
#include <common/barrier.h>
```

Data Structures

- struct _starpu_barrier_counter

Functions

- int **_starpu_barrier_counter_init** (struct _starpu_barrier_counter *barrier_c, unsigned count)
- int **_starpu_barrier_counter_destroy** (struct _starpu_barrier_counter *barrier_c)
- int **_starpu_barrier_counter_wait_for_empty_counter** (struct _starpu_barrier_counter *barrier_c)
- int **_starpu_barrier_counter_wait_until_counter_reaches_down_to_n** (struct _starpu_barrier_counter *barrier_c, unsigned n)
- int **_starpu_barrier_counter_wait_until_counter_reaches_up_to_n** (struct _starpu_barrier_counter *barrier_c, unsigned n)
- int **_starpu_barrier_counter_wait_for_full_counter** (struct _starpu_barrier_counter *barrier_c)
- int **_starpu_barrier_counter_decrement_until_empty_counter** (struct _starpu_barrier_counter *barrier↩_c, double flops)
- int **_starpu_barrier_counter_increment_until_full_counter** (struct _starpu_barrier_counter *barrier_↩c, double flops)
- int **_starpu_barrier_counter_increment** (struct _starpu_barrier_counter *barrier_c, double flops)
- int **_starpu_barrier_counter_check** (struct _starpu_barrier_counter *barrier_c)
- int **_starpu_barrier_counter_get_reached_start** (struct _starpu_barrier_counter *barrier_c)
- double **_starpu_barrier_counter_get_reached_flops** (struct _starpu_barrier_counter *barrier_c)

6.2.1 Data Structure Documentation

6.2.1.1 struct _starpu_barrier_counter

Data Fields

struct _starpu_barrier	barrier	
unsigned	min_threshold	
unsigned	max_threshold	
starpu_pthread_cond_t	cond2	

6.3 bound.h File Reference

```
#include <starpu.h>
#include <starpu_bound.h>
#include <core/jobs.h>
```

Functions

- void _starpu_bound_record (struct _starpu_job *j)
- void _starpu_bound_tag_dep (starpu_tag_t id, starpu_tag_t dep_id)
- void _starpu_bound_task_dep (struct _starpu_job *j, struct _starpu_job *dep_j)
- void _starpu_bound_job_id_dep (starpu_data_handle_t handle, struct _starpu_job *dep_j, unsigned long job_id)
- void starpu_bound_clear (void)

Variables

- int _starpu_bound_recording

6.3.1 Function Documentation

6.3.1.1 _starpu_bound_record()

```
void _starpu_bound_record (
            struct _starpu_job * j )
```
Record task for bound computation

6.3.1.2 _starpu_bound_tag_dep()

```
void _starpu_bound_tag_dep (
            starpu_tag_t id,
            starpu_tag_t dep_id )
```
Record tag dependency: id depends on dep_id

6.3.1.3 _starpu_bound_task_dep()

```
void _starpu_bound_task_dep (
            struct _starpu_job * j,
            struct _starpu_job * dep_j )
```
Record task dependency: j depends on dep_j

6.3.1.4 _starpu_bound_job_id_dep()

```
void _starpu_bound_job_id_dep (
            starpu_data_handle_t handle,
            struct _starpu_job * dep_j,
            unsigned long job_id )
```
Record job id dependency: j depends on job_id

6.3.1.5 starpu_bound_clear()

```
void starpu_bound_clear (
            void  )
```
Clear recording

6.3.2 Variable Documentation

6.3.2.1 _starpu_bound_recording

```
int _starpu_bound_recording
```
Are we recording?

6.4 cg.h File Reference

```
#include <starpu.h>
#include <common/config.h>
```

Data Structures

- struct _starpu_cg_list
- struct _starpu_cg
- union _starpu_cg.succ
- struct _starpu_cg.succ.succ_apps

Macros

- #define STARPU_DYNAMIC_DEPS_SIZE

Typedefs

- typedef [struct](#) _starpu_notify_job_start_data **_starpu_notify_job_start_data**

Enumerations

- enum **_starpu_cg_type** { **STARPU_CG_APPS**, **STARPU_CG_TAG**, **STARPU_CG_TASK** }

Functions

- void **_starpu_notify_dependencies** ([struct _starpu_job](#) ∗j)
- void **_starpu_job_notify_start** ([struct _starpu_job](#) ∗j, [struct](#) starpu_perfmodel_arch ∗perf_arch)
- void **_starpu_job_notify_ready_soon** ([struct _starpu_job](#) ∗j, _starpu_notify_job_start_data ∗data)
- void **_starpu_cg_list_init** ([struct _starpu_cg_list](#) ∗list)
- void **_starpu_cg_list_deinit** ([struct _starpu_cg_list](#) ∗list)
- int **_starpu_add_successor_to_cg_list** ([struct _starpu_cg_list](#) ∗successors, [struct _starpu_cg](#) ∗cg)
- int **_starpu_list_task_successors_in_cg_list** ([struct _starpu_cg_list](#) ∗successors, unsigned ndeps, [struct](#) starpu_task ∗task_array[ ])
- int **_starpu_list_task_scheduled_successors_in_cg_list** ([struct _starpu_cg_list](#) ∗successors, unsigned ndeps, [struct](#) starpu_task ∗task_array[ ])
- int **_starpu_list_tag_successors_in_cg_list** ([struct _starpu_cg_list](#) ∗successors, unsigned ndeps, starpu_tag_t tag_array[ ])
- void **_starpu_notify_cg** (void ∗pred, [struct _starpu_cg](#) ∗cg)
- void **_starpu_notify_cg_list** (void ∗pred, [struct _starpu_cg_list](#) ∗successors)
- void **_starpu_notify_job_start_cg_list** (void ∗pred, [struct _starpu_cg_list](#) ∗successors, _starpu_notify_↩
job_start_data ∗data)
- void **_starpu_notify_task_dependencies** ([struct _starpu_job](#) ∗j)
- void **_starpu_notify_job_start_tasks** ([struct _starpu_job](#) ∗j, _starpu_notify_job_start_data ∗data)

6.4.1 Data Structure Documentation

6.4.1.1 struct _starpu_cg_list

Completion Group list, records both the number of expected notifications before the completion can start, and the list of successors when the completion is finished.

Data Fields

[struct _starpu_spinlock](#)	lock	Protects atomicity of the list and the terminated flag
unsigned	ndeps	Number of notifications to be waited for
unsigned	ndeps_completed	
unsigned	terminated	Whether the completion is finished. For restartable/restarted tasks, only the first iteration is taken into account here.
unsigned	nsuccs	List of successors
unsigned	succ_list_size	How many allocated items in succ
[struct _starpu_cg](#) ∗∗	succ	

6.4.1.2 struct _starpu_cg

Completion Group

Data Fields

unsigned	ntags	number of tags depended on
unsigned	remaining	number of remaining tags
enum _starpu_cg_type	cg_type	

Data Fields

union _starpu_cg	succ	

6.4.1.3 union _starpu_cg.succ

Data Fields

struct _starpu_tag *	tag	STARPU_CG_TAG
struct _starpu_job *	job	STARPU_CG_TASK
succ	succ_apps	STARPU_CG_APPS in case this completion group is related to an application, we have to explicitly wake the waiting thread instead of reschedule the corresponding task

6.4.1.4 struct _starpu_cg.succ.succ_apps

STARPU_CG_APPS in case this completion group is related to an application, we have to explicitly wake the waiting thread instead of reschedule the corresponding task

Data Fields

unsigned	completed	
starpu_pthread_mutex_t	cg_mutex	
starpu_pthread_cond_t	cg_cond	

6.4.2 Macro Definition Documentation

6.4.2.1 STARPU_DYNAMIC_DEPS_SIZE

```
#define STARPU_DYNAMIC_DEPS_SIZE
```
we do not necessarily want to allocate room for 256 dependencies, but we want to handle the few situation where there are a lot of dependencies as well

6.5 coherency.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/starpu_spinlock.h>
#include <common/rwlock.h>
#include <common/timing.h>
#include <common/fxt.h>
#include <common/list.h>
#include <datawizard/interfaces/data_interface.h>
#include <datawizard/datastats.h>
#include <datawizard/memstats.h>
#include <datawizard/data_request.h>
```

Data Structures

- struct _starpu_data_replicate

- struct _starpu_jobid_list
- struct _starpu_task_wrapper_list
- struct _starpu_task_wrapper_dlist
- struct _starpu_data_state

Typedefs

- typedef void(∗ **_starpu_data_handle_unregister_hook**) (starpu_data_handle_t)

Enumerations

- enum **_starpu_cache_state** { **STARPU_OWNER**, **STARPU_SHARED**, **STARPU_INVALID** }

Functions

- int _starpu_fetch_data_on_node (starpu_data_handle_t handle, int node, struct _starpu_data_replicate ∗replicate, enum starpu_data_access_mode mode, unsigned detached, enum _starpu_is_prefetch is_← prefetch, unsigned async, void(∗callback_func)(void ∗), void ∗callback_arg, int prio, const char ∗origin)
- void _starpu_release_data_on_node (struct _starpu_data_state ∗state, uint32_t default_wt_mask, struct _← starpu_data_replicate ∗replicate)
- void **_starpu_update_data_state** (starpu_data_handle_t handle, struct _starpu_data_replicate ∗requesting← _replicate, enum starpu_data_access_mode mode)
- uint32_t **_starpu_get_data_refcnt** (struct _starpu_data_state ∗state, unsigned node)
- size_t **_starpu_data_get_size** (starpu_data_handle_t handle)
- size_t **_starpu_data_get_alloc_size** (starpu_data_handle_t handle)
- uint32_t **_starpu_data_get_footprint** (starpu_data_handle_t handle)
- void **__starpu_push_task_output** (struct _starpu_job ∗j)
- void _starpu_push_task_output (struct _starpu_job ∗j)
- void **_starpu_release_nowhere_task_output** (struct _starpu_job ∗j)
- STARPU_ATTRIBUTE_WARN_UNUSED_RESULT int **_starpu_fetch_task_input** (struct starpu_task ∗task, struct _starpu_job ∗j, int async)
- void **_starpu_fetch_task_input_tail** (struct starpu_task ∗task, struct _starpu_job ∗j, struct _starpu_worker ∗worker)
- void **_starpu_fetch_nowhere_task_input** (struct _starpu_job ∗j)
- int **_starpu_select_src_node** (struct _starpu_data_state ∗state, unsigned destination)
- int **_starpu_determine_request_path** (starpu_data_handle_t handle, int src_node, int dst_node, enum starpu_data_access_mode mode, int max_len, unsigned ∗src_nodes, unsigned ∗dst_nodes, unsigned ∗handling_nodes, unsigned write_invalidation)
- struct _starpu_data_request ∗ _starpu_create_request_to_fetch_data (starpu_data_handle_t handle, struct _starpu_data_replicate ∗dst_replicate, enum starpu_data_access_mode mode, enum _starpu_is_prefetch is_prefetch, unsigned async, void(∗callback_func)(void ∗), void ∗callback_arg, int prio, const char ∗origin)
- void **_starpu_redux_init_data_replicate** (starpu_data_handle_t handle, struct _starpu_data_replicate ∗replicate, int workerid)
- void **_starpu_data_start_reduction_mode** (starpu_data_handle_t handle)
- void **_starpu_data_end_reduction_mode** (starpu_data_handle_t handle)
- void **_starpu_data_end_reduction_mode_terminate** (starpu_data_handle_t handle)
- void **_starpu_data_set_unregister_hook** (starpu_data_handle_t handle, _starpu_data_handle_← unregister_hook func)

Variables

- int **_starpu_has_not_important_data**

6.5.1 Data Structure Documentation

6.5.1.1 struct _starpu_data_replicate

this should contain the information relative to a given data replicate

Data Fields

starpu_data_handle_t	handle	
void ∗	data_interface	describe the actual data layout, as manipulated by data interfaces in ∗_interface.c
int	refcnt	How many requests or tasks are currently working with this replicate
char	memory_node	
enum _starpu_cache_state	state: 2	describes the state of the local data in term of coherency
unsigned	relaxed_coherency:2	A buffer that is used for SCRATCH or reduction cannnot be used with filters.
unsigned	initialized:1	We may need to initialize the replicate with some value before using it.
unsigned	allocated:1	is the data locally allocated ?
unsigned	automatically_allocated:1	was it automatically allocated ? (else it's the application-provided buffer, don't ever try to free it!) perhaps the allocation was perform higher in the hiearchy for now this is just translated into !automatically_allocated
struct _starpu_data_request ∗	request[STARPU_MAXNODES]	To help the scheduling policies to make some decision, we may keep a track of the tasks that are likely to request this data on the current node. It is the responsability of the scheduling _policy_ to set that flag when it assigns a task to a queue, policies which do not use this hint can simply ignore it.
struct _starpu_mem_chunk ∗	mc	Pointer to memchunk for LRU strategy

6.5.1.2 struct _starpu_jobid_list

Data Fields

unsigned long	id	
struct _starpu_jobid_list ∗	next	

6.5.1.3 struct _starpu_task_wrapper_list

This structure describes a simply-linked list of task

Data Fields

struct starpu_task ∗	task	
struct _starpu_task_wrapper_list ∗	next	

6.5.1.4 struct _starpu_task_wrapper_dlist

This structure describes a doubly-linked list of task

Data Fields

struct starpu_task ∗	task	
struct _starpu_task_wrapper_dlist ∗	next	
struct _starpu_task_wrapper_dlist ∗	prev	

6.5.1.5 struct _starpu_data_state

This is initialized in both _starpu_register_new_data and _starpu_data_partition

Data Fields

int	magic	
struct _starpu_data_requester_prio_list	req_list	
unsigned	refcnt	the number of requests currently in the scheduling engine (not in the req_list anymore), i.e. the number of holders of the current_mode rwlock
unsigned	unlocking_reqs	whether we are already unlocking data requests
enum starpu_data_access_mode	current_mode	Current access mode. Is always either STARPU_R, STARPU_W, STARPU_SCRATCH or STARPU_REDUX, but never a combination such as STARPU_RW.
struct _starpu_spinlock	header_lock	protect meta data
unsigned	busy_count	Condition to make application wait for all transfers before freeing handle busy_count is the number of handle->refcnt, handle->per_node[∗]->refcnt, number of starpu_data_requesters, and number of tasks that have released it but are still registered on the implicit data dependency lists. Core code which releases busy_count has to call _starpu_data_check_not_busy to let starpu_data_unregister proceed
unsigned	busy_waiting	Is starpu_data_unregister waiting for busy_count?
starpu_pthread_mutex_t	busy_mutex	
starpu_pthread_cond_t	busy_cond	
struct _starpu_data_state ∗	root_handle	In case we user filters, the handle may describe a sub-data
struct _starpu_data_state ∗	father_handle	root of the tree
starpu_data_handle_t ∗	active_children	father of the node, NULL if the current node is the root
starpu_data_handle_t ∗∗	active_readonly_children	The currently active set of read-write children

Data Fields

unsigned	nactive_readonly_children	The currently active set of read-only children
unsigned	nsiblings	Size of active_readonly_children array Our siblings in the father partitioning
starpu_data_handle_t *	siblings	How many siblings
unsigned	sibling_index	
unsigned	depth	indicate which child this node is from the father's perpsective (if any)
starpu_data_handle_t	children	what's the depth of the tree ? Synchronous partitioning
unsigned	nchildren	
unsigned	nplans	How many partition plans this handle has
struct starpu_codelet *	switch_cl	Switch codelet for asynchronous partitioning
unsigned	switch_cl_nparts	size of dyn_nodes recorded in switch_cl
unsigned	partitioned	Whether a partition plan is currently submitted and the corresponding unpartition has not been yet Or the number of partition plans currently submitted in readonly mode.
unsigned	readonly:1	Whether a partition plan is currently submitted in readonly mode
unsigned	active:1	Whether our father is currently partitioned into ourself
unsigned	active_ro:1	
struct _starpu_data_replicate	per_node[STARPU_MAXNODES]	describe the state of the data in term of coherency
struct _starpu_data_replicate *	per_worker	
struct starpu_data_interface_ops *	ops	
uint32_t	footprint	Footprint which identifies data layout
int	home_node	where is the data home, i.e. which node it was registered from ? -1 if none yet
uint32_t	wt_mask	what is the default write-through mask for that data ?
unsigned	is_not_important	in some case, the application may explicitly tell StarPU that a piece of data is not likely to be used soon again
unsigned	sequential_consistency	Does StarPU have to enforce some implicit data-dependencies ?
unsigned	initialized	Is the data initialized, or a task is already submitted to initialize it

Data Fields

unsigned	ooc	Can the data be pushed to the disk?
starpu_pthread_mutex_t	sequential_consistency_mutex	This lock should protect any operation to enforce sequential_consistency
enum starpu_data_access_mode	last_submitted_mode	The last submitted task (or application data request) that declared it would modify the piece of data ? Any task accessing the data in a read-only mode should depend on that task implicitly if the sequential_consistency flag is enabled.
struct starpu_task ∗	last_sync_task	
struct _starpu_task_wrapper_dlist	last_submitted_accessors	
unsigned	last_submitted_ghost_sync_id_is_valid	If FxT is enabled, we keep track of "ghost dependencies": that is to say the dependencies that are not needed anymore, but that should appear in the post-mortem DAG. For instance if we have the sequence f(Aw) g(Aw), and that g is submitted after the termination of f, we want to have f->g appear in the DAG even if StarPU does not need to enforce this dependency anymore.
unsigned long	last_submitted_ghost_sync_id	
struct _starpu_jobid_list ∗	last_submitted_ghost_accessors_id	
struct _starpu_task_wrapper_list ∗	post_sync_tasks	protected by sequential_consistency_mutex
unsigned	post_sync_tasks_cnt	
struct starpu_codelet ∗	redux_cl	During reduction we need some specific methods: redux_func performs the reduction of an interface into another one (eg. "+="), and init_func initializes the data interface to a default value that is stable by reduction (eg. 0 for +=).
struct starpu_codelet ∗	init_cl	
unsigned	reduction_refcnt	Are we currently performing a reduction on that handle ? If so the reduction_refcnt should be non null until there are pending tasks that are performing the reduction.
struct _starpu_data_requester_prio_list	reduction_req_list	List of requesters that are specific to the pending reduction. This list is used when the requests in the req_list list are frozen until the end of the reduction.
starpu_data_handle_t ∗	reduction_tmp_handles	
struct _starpu_data_request ∗	write_invalidation_req	Final request for write invalidation

Data Fields

unsigned	lazy_unregister	
unsigned	removed_from_context_hash	
void *	mpi_data	Used for MPI
_starpu_memory_stats_t	memory_stats	
unsigned int	mf_node	
_starpu_data_handle_unregister_hook	unregister_hook	hook to be called when unregistering the data
struct starpu_arbiter *	arbiter	
struct _starpu_data_requester_prio_list	arbitered_req_list	This is protected by the arbiter mutex
int	last_locality	Data maintained by schedulers themselves Last worker that took this data in locality mode, or -1 if nobody took it yet
unsigned	dimensions	Application-provided coordinates. The maximum dimension (5) is relatively arbitrary.
int	coordinates[5]	
void *	user_data	A generic pointer to data in the user land (could be anything and this is not manage by StarPU)

6.5.2 Function Documentation

6.5.2.1 _starpu_fetch_data_on_node()

```
int _starpu_fetch_data_on_node (
        starpu_data_handle_t handle,
        int node,
        struct _starpu_data_replicate * replicate,
        enum starpu_data_access_mode mode,
        unsigned detached,
        enum _starpu_is_prefetch is_prefetch,
        unsigned async,
        void(*)(void *) callback_func,
        void * callback_arg,
        int prio,
        const char * origin )
```

This does not take a reference on the handle, the caller has to do it, e.g. through _starpu_attempt_to_submit↵
_data_request_from_apps() detached means that the core is allowed to drop the request. The caller should thus
not take a reference since it can not know whether the request will complete async means that _starpu_fetch_↵
data_on_node will wait for completion of the request

6.5.2.2 _starpu_release_data_on_node()

```
void _starpu_release_data_on_node (
        struct _starpu_data_state * state,
        uint32_t default_wt_mask,
        struct _starpu_data_replicate * replicate )
```

This releases a reference on the handle

6.5.2.3 _starpu_push_task_output()

```
void _starpu_push_task_output (
            struct _starpu_job * j )
```
Version with driver trace

6.5.2.4 _starpu_create_request_to_fetch_data()

```
struct _starpu_data_request* _starpu_create_request_to_fetch_data (
            starpu_data_handle_t handle,
            struct _starpu_data_replicate * dst_replicate,
            enum starpu_data_access_mode mode,
            enum _starpu_is_prefetch is_prefetch,
            unsigned async,
            void(*)(void *) callback_func,
            void * callback_arg,
            int prio,
            const char * origin )
```
is_prefetch is whether the DSM may drop the request (when there is not enough memory for instance async is whether the caller wants a reference on the last request, to be able to wait for it (which will release that reference).

6.6 combined_workers.h File Reference

```
#include <starpu.h>
#include <common/config.h>
```

6.7 config.h File Reference

Macros

- #define **CONFIG_FUT**
- #define **HAVE_AIO_H**
- #define **HAVE_AYUDAME_H**
- #define **HAVE_CBLAS_SGEMV**
- #define **HAVE_CLENQUEUEMARKERWITHWAITLIST**
- #define **HAVE_CLGETEXTENSIONFUNCTIONADDRESSFORPLATFORM**
- #define **HAVE_CLOCK_GETTIME**
- #define **HAVE_CL_CL_EXT_H**
- #define **HAVE_COPY_FILE_RANGE**
- #define **HAVE_CUDA_GL_INTEROP_H**
- #define **HAVE_CXX11**
- #define **HAVE_DECL_CUSPARSESETSTREAM**
- #define **HAVE_DECL_ENABLE_FUT_FLUSH**
- #define **HAVE_DECL_FUT_SETUP_FLUSH_CALLBACK**
- #define **HAVE_DECL_FUT_SET_FILENAME**
- #define **HAVE_DECL_HWLOC_CUDA_GET_DEVICE_OSDEV_BY_INDEX**
- #define **HAVE_DECL_NVMLDEVICEGETTOTALENERGYCONSUMPTION**
- #define **HAVE_DECL_SMPI_PROCESS_SET_USER_DATA**
- #define **HAVE_DLB_H**
- #define **HAVE_DLFCN_H**
- #define **HAVE_ENABLE_FUT_FLUSH**
- #define **HAVE_FUT_SETUP_FLUSH_CALLBACK**
- #define **HAVE_FUT_SET_FILENAME**
- #define **HAVE_FXT_CLOSE**
- #define **HAVE_GETRLIMIT**

- #define **HAVE_GLPK_H**
- #define **HAVE_HDF5_H**
- #define **HAVE_HWLOC_CPUKINDS_GET_NR**
- #define **HAVE_HWLOC_GLIBC_SCHED_H**
- #define **HAVE_HWLOC_TOPOLOGY_DUP**
- #define **HAVE_HWLOC_TOPOLOGY_SET_COMPONENTS**
- #define **HAVE_INTTYPES_H**
- #define **HAVE_LEVELDB_DB_H**
- #define **HAVE_LIBATLAS**
- #define **HAVE_LIBBLAS_OPENBLAS**
- #define **HAVE_LIBCBLAS**
- #define **HAVE_LIBCUSPARSE**
- #define **HAVE_LIBDLB**
- #define **HAVE_LIBGFORTRAN**
- #define **HAVE_LIBGL**
- #define **HAVE_LIBGLPK**
- #define **HAVE_LIBGLU**
- #define **HAVE_LIBGLUT**
- #define **HAVE_LIBGOTO**
- #define **HAVE_LIBGOTO2**
- #define **HAVE_LIBHDF5**
- #define **HAVE_LIBIFCORE**
- #define **HAVE_LIBLEVELDB**
- #define **HAVE_LIBOPENBLAS**
- #define **HAVE_LIBRT**
- #define **HAVE_LIBSIMGRID**
- #define **HAVE_LIBWS2_32**
- #define **HAVE_MALLOC_H**
- #define **HAVE_MEMALIGN**
- #define **HAVE_MEMORY_H**
- #define **HAVE_MKDTEMP**
- #define **HAVE_MKOSTEMP**
- #define **HAVE_MPI_COMM_F2C**
- #define **HAVE_MSG_ENVIRONMENT_GET_ROUTING_ROOT**
- #define **HAVE_MSG_GET_AS_BY_NAME**
- #define **HAVE_MSG_HOST_GET_SPEED**
- #define **HAVE_MSG_MSG_H**
- #define **HAVE_MSG_PROCESS_ATTACH**
- #define **HAVE_MSG_PROCESS_SELF_NAME**
- #define **HAVE_MSG_PROCESS_USERDATA_INIT**
- #define **HAVE_MSG_ZONE_GET_BY_NAME**
- #define **HAVE_MSG_ZONE_GET_HOSTS**
- #define **HAVE_NVMLDEVICEGETTOTALENERGYCONSUMPTION**
- #define **HAVE_PIOM_LTASK_SET_BOUND_THREAD_OS_INDEXES**
- #define **HAVE_POSIX_MEMALIGN**
- #define **HAVE_POTI_INIT_CUSTOM**
- #define **HAVE_POTI_USER_NEWEVENT**
- #define **HAVE_PREAD**
- #define **HAVE_PTHREAD_SETAFFINITY_NP**
- #define **HAVE_PTHREAD_SPIN_LOCK**
- #define **HAVE_PWRITE**
- #define **HAVE_SCANDIR**
- #define **HAVE_SG_ACTOR_ATTACH**
- #define **HAVE_SG_ACTOR_DATA**
- #define **HAVE_SG_ACTOR_EXECUTE**

- #define **HAVE_SG_ACTOR_GET_DATA**
- #define **HAVE_SG_ACTOR_INIT**
- #define **HAVE_SG_ACTOR_ON_EXIT**
- #define **HAVE_SG_ACTOR_REF**
- #define **HAVE_SG_ACTOR_SELF**
- #define **HAVE_SG_ACTOR_SELF_EXECUTE**
- #define **HAVE_SG_ACTOR_SET_DATA**
- #define **HAVE_SG_ACTOR_SLEEP_FOR**
- #define **HAVE_SG_CFG_SET_INT**
- #define **HAVE_SG_CONFIG_CONTINUE_AFTER_HELP**
- #define **HAVE_SG_HOST_GET_PROPERTIES**
- #define **HAVE_SG_HOST_GET_ROUTE**
- #define **HAVE_SG_HOST_GET_SPEED**
- #define **HAVE_SG_HOST_LIST**
- #define **HAVE_SG_HOST_ROUTE**
- #define **HAVE_SG_HOST_SELF**
- #define **HAVE_SG_HOST_SENDTO**
- #define **HAVE_SG_HOST_SEND_TO**
- #define **HAVE_SG_HOST_SPEED**
- #define **HAVE_SG_LINK_BANDWIDTH_SET**
- #define **HAVE_SG_LINK_GET_NAME**
- #define **HAVE_SG_LINK_NAME**
- #define **HAVE_SG_LINK_SET_BANDWIDTH**
- #define **HAVE_SG_ZONE_GET_BY_NAME**
- #define **HAVE_SG_ZONE_GET_HOSTS**
- #define **HAVE_SIMCALL_PROCESS_CREATE**
- #define **HAVE_SIMGRID_ACTOR_H**
- #define **HAVE_SIMGRID_BARRIER_H**
- #define **HAVE_SIMGRID_COND_H**
- #define **HAVE_SIMGRID_ENGINE_H**
- #define **HAVE_SIMGRID_GET_CLOCK**
- #define **HAVE_SIMGRID_HOST_H**
- #define **HAVE_SIMGRID_INIT**
- #define **HAVE_SIMGRID_MSG_H**
- #define **HAVE_SIMGRID_MUTEX_H**
- #define **HAVE_SIMGRID_SEMAPHORE_H**
- #define **HAVE_SIMGRID_SET_MAESTRO**
- #define **HAVE_SIMGRID_SIMDAG_H**
- #define **HAVE_SIMGRID_VERSION_H**
- #define **HAVE_SIMGRID_ZONE_H**
- #define **HAVE_SMPI_PROCESS_SET_USER_DATA**
- #define **HAVE_SMPI_THREAD_CREATE**
- #define **HAVE_SMX_ACTOR_T**
- #define **HAVE_STDINT_H**
- #define **HAVE_STDLIB_H**
- #define **HAVE_STRINGS_H**
- #define **HAVE_STRING_H**
- #define **HAVE_SYSCONF**
- #define **HAVE_SYS_STAT_H**
- #define **HAVE_SYS_TYPES_H**
- #define **HAVE_UNISTD_H**
- #define **HAVE_VALGRIND_HELGRIND_H**
- #define **HAVE_VALGRIND_MEMCHECK_H**
- #define **HAVE_VALGRIND_VALGRIND_H**
- #define **HAVE_XBT_BARRIER_INIT**

- #define **HAVE_XBT_BASE_H**
- #define **HAVE_XBT_CONFIG_H**
- #define **HAVE_XBT_MUTEX_TRY_ACQUIRE**
- #define **HAVE_XBT_SYNCHRO_H**
- #define **LT_OBJDIR**
- #define **PACKAGE**
- #define **PACKAGE_BUGREPORT**
- #define **PACKAGE_NAME**
- #define **PACKAGE_STRING**
- #define **PACKAGE_TARNAME**
- #define **PACKAGE_URL**
- #define **PACKAGE_VERSION**
- #define **SIZEOF_VOID_P**
- #define **STARPURM_DLB_VERBOSE**
- #define **STARPURM_HAVE_DLB**
- #define **STARPURM_HAVE_DLB_CALLBACK_ARG**
- #define **STARPURM_STARPU_HAVE_WORKER_CALLBACKS**
- #define **STARPURM_VERBOSE**
- #define **STARPU_ARMPL**
- #define **STARPU_ATLAS**
- #define **STARPU_BUILD_DIR**
- #define **STARPU_BUILT_IN_MIN_DGELS**
- #define **STARPU_CLUSTER**
- #define **STARPU_DEBUG**
- #define **STARPU_DEVEL**
- #define **STARPU_DISABLE_ASYNCHRONOUS_COPY**
- #define **STARPU_DISABLE_ASYNCHRONOUS_CUDA_COPY**
- #define **STARPU_DISABLE_ASYNCHRONOUS_MIC_COPY**
- #define **STARPU_DISABLE_ASYNCHRONOUS_MPI_MS_COPY**
- #define **STARPU_DISABLE_ASYNCHRONOUS_OPENCL_COPY**
- #define **STARPU_EXTRA_VERBOSE**
- #define **STARPU_FXT_LOCK_TRACES**
- #define **STARPU_GDB_PATH**
- #define **STARPU_GOTO**
- #define **STARPU_HAVE_ATOMIC_COMPARE_EXCHANGE_N**
- #define **STARPU_HAVE_ATOMIC_EXCHANGE_N**
- #define **STARPU_HAVE_ATOMIC_FETCH_ADD**
- #define **STARPU_HAVE_ATOMIC_FETCH_OR**
- #define **STARPU_HAVE_ATOMIC_TEST_AND_SET**
- #define **STARPU_HAVE_BLAS**
- #define **STARPU_HAVE_BUSID**
- #define **STARPU_HAVE_CBLAS_H**
- #define **STARPU_HAVE_CUDA_MEMCPY_PEER**
- #define **STARPU_HAVE_CUFFTDOUBLECOMPLEX**
- #define **STARPU_HAVE_CURAND**
- #define **STARPU_HAVE_CXX11**
- #define **STARPU_HAVE_DARWIN**
- #define **STARPU_HAVE_DOMAINID**
- #define **STARPU_HAVE_F77_H**
- #define **STARPU_HAVE_FC**
- #define **STARPU_HAVE_FFTW**
- #define **STARPU_HAVE_FFTWF**
- #define **STARPU_HAVE_FFTWL**
- #define **STARPU_HAVE_GLPK_H**
- #define **STARPU_HAVE_HDF5**

- #define **STARPU_HAVE_HELGRIND_H**
- #define **STARPU_HAVE_HWLOC**
- #define **STARPU_HAVE_ICC**
- #define **STARPU_HAVE_LEVELDB**
- #define **STARPU_HAVE_LIBNUMA**
- #define **STARPU_HAVE_LIBNVIDIA_ML**
- #define **STARPU_HAVE_MAGMA**
- #define **STARPU_HAVE_MALLOC_H**
- #define **STARPU_HAVE_MEMALIGN**
- #define **STARPU_HAVE_MEMCHECK_H**
- #define **STARPU_HAVE_MSG_MSG_H**
- #define **STARPU_HAVE_NEARBYINTF**
- #define **STARPU_HAVE_POSIX_MEMALIGN**
- #define **STARPU_HAVE_POTI**
- #define **STARPU_HAVE_PTHREAD_BARRIER**
- #define **STARPU_HAVE_PTHREAD_SETNAME_NP**
- #define **STARPU_HAVE_PTHREAD_SPIN_LOCK**
- #define **STARPU_HAVE_RINTF**
- #define **STARPU_HAVE_SCHED_YIELD**
- #define **STARPU_HAVE_SETENV**
- #define **STARPU_HAVE_SIMGRID_ACTOR_H**
- #define **STARPU_HAVE_SIMGRID_BARRIER_H**
- #define **STARPU_HAVE_SIMGRID_COND_H**
- #define **STARPU_HAVE_SIMGRID_ENGINE_H**
- #define **STARPU_HAVE_SIMGRID_HOST_H**
- #define **STARPU_HAVE_SIMGRID_MSG_H**
- #define **STARPU_HAVE_SIMGRID_MUTEX_H**
- #define **STARPU_HAVE_SIMGRID_SEMAPHORE_H**
- #define **STARPU_HAVE_SIMGRID_SIMDAG_H**
- #define **STARPU_HAVE_SIMGRID_VERSION_H**
- #define **STARPU_HAVE_SIMGRID_ZONE_H**
- #define **STARPU_HAVE_SMX_ACTOR_T**
- #define **STARPU_HAVE_STATEMENT_EXPRESSIONS**
- #define **STARPU_HAVE_STRERROR_R**
- #define **STARPU_HAVE_STRUCT_TIMESPEC**
- #define **STARPU_HAVE_SYNC_BOOL_COMPARE_AND_SWAP**
- #define **STARPU_HAVE_SYNC_FETCH_AND_ADD**
- #define **STARPU_HAVE_SYNC_FETCH_AND_OR**
- #define **STARPU_HAVE_SYNC_LOCK_TEST_AND_SET**
- #define **STARPU_HAVE_SYNC_SYNCHRONIZE**
- #define **STARPU_HAVE_SYNC_VAL_COMPARE_AND_SWAP**
- #define **STARPU_HAVE_UNISTD_H**
- #define **STARPU_HAVE_UNSETENV**
- #define **STARPU_HAVE_VALGRIND_H**
- #define **STARPU_HAVE_WINDOWS**
- #define **STARPU_HAVE_X11**
- #define **STARPU_HAVE_XBT_BASE_H**
- #define **STARPU_HAVE_XBT_CONFIG_H**
- #define **STARPU_HAVE_XBT_SYNCHRO_H**
- #define **STARPU_HISTORYMAXERROR**
- #define **STARPU_LINUX_SYS**
- #define **STARPU_LONG_CHECK**
- #define **STARPU_MAJOR_VERSION**
- #define **STARPU_MAXCPUS**
- #define **STARPU_MAXCUDADEVS**

- #define **STARPU_MAXIMPLEMENTATIONS**
- #define **STARPU_MAXMICCORES**
- #define **STARPU_MAXMICDEVS**
- #define **STARPU_MAXMPIDEVS**
- #define **STARPU_MAXMPKERNELS**
- #define **STARPU_MAXNODES**
- #define **STARPU_MAXNUMANODES**
- #define **STARPU_MAXOPENCLDEVS**
- #define **STARPU_MEMORY_STATS**
- #define **STARPU_MIC_USE_RMA**
- #define **STARPU_MINOR_VERSION**
- #define **STARPU_MKL**
- #define **STARPU_MLR_MODEL**
- #define **STARPU_MODEL_DEBUG**
- #define **STARPU_MPI_EXTRA_VERBOSE**
- #define **STARPU_MPI_MASTER_SLAVE_MULTIPLE_THREAD**
- #define **STARPU_MPI_PEDANTIC_ISEND**
- #define **STARPU_MPI_VERBOSE**
- #define **STARPU_NATIVE_WINTHREADS**
- #define **STARPU_NEW_CHECK**
- #define **STARPU_NMAXBUFS**
- #define **STARPU_NMAXWORKERS**
- #define **STARPU_NMAX_COMBINEDWORKERS**
- #define **STARPU_NMAX_SCHED_CTXS**
- #define **STARPU_NON_BLOCKING_DRIVERS**
- #define **STARPU_NO_ASSERT**
- #define **STARPU_OPENBLAS**
- #define **STARPU_OPENBSD_SYS**
- #define **STARPU_OPENCL_SIMULATOR**
- #define **STARPU_OPENGL_RENDER**
- #define **STARPU_OPENMP**
- #define **STARPU_PERF_DEBUG**
- #define **STARPU_PERF_MODEL_DIR**
- #define **STARPU_QUICK_CHECK**
- #define **STARPU_RELEASE_VERSION**
- #define **STARPU_SC_HYPERVISOR_DEBUG**
- #define **STARPU_SIMGRID**
- #define **STARPU_SIMGRID_HAVE_SIMGRID_INIT**
- #define **STARPU_SIMGRID_HAVE_XBT_BARRIER_INIT**
- #define **STARPU_SIMGRID_MC**
- #define **STARPU_SPINLOCK_CHECK**
- #define **STARPU_SRC_DIR**
- #define **STARPU_STATIC_ONLY**
- #define **STARPU_SYSTEM_BLAS**
- #define **STARPU_USE_ALLOCATION_CACHE**
- #define **STARPU_USE_AYUDAME1**
- #define **STARPU_USE_AYUDAME2**
- #define **STARPU_USE_CPU**
- #define **STARPU_USE_CUDA**
- #define **STARPU_USE_DRAND48**
- #define **STARPU_USE_ERAND48_R**
- #define **STARPU_USE_FXT**
- #define **STARPU_USE_MIC**
- #define **STARPU_USE_MP**
- #define **STARPU_USE_MPI**

- #define **STARPU_USE_MPI_MASTER_SLAVE**
- #define **STARPU_USE_MPI_MPI**
- #define **STARPU_USE_MPI_NMAD**
- #define **STARPU_USE_OPENCL**
- #define **STARPU_USE_SC_HYPERVISOR**
- #define **STARPU_VALGRIND_FULL**
- #define **STARPU_VERBOSE**
- #define **STARPU_WORKER_CALLBACKS**
- #define **STDC_HEADERS**
- #define **VERSION**
- #define **X_DISPLAY_MISSING**
- #define **restrict**

6.8 copy_driver.h File Reference

```
#include <common/config.h>
#include <common/list.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include <starpu_opencl.h>
```

Data Structures

- struct _starpu_mic_async_event
- struct _starpu_disk_backend_event
- struct _starpu_disk_async_event
- union _starpu_async_channel_event
- struct _starpu_async_channel
- struct _starpu_async_channel_event.__unnamed__

Enumerations

- enum _starpu_is_prefetch { STARPU_FETCH, STARPU_PREFETCH, STARPU_IDLEFETCH, **STARPU_↩ NFETCH** }

Functions

- void **_starpu_wake_all_blocked_workers_on_node** (unsigned nodeid)
- int **_starpu_driver_copy_data_1_to_1** (starpu_data_handle_t handle, struct _starpu_data_replicate ∗src↩ _replicate, struct _starpu_data_replicate ∗dst_replicate, unsigned donotread, struct _starpu_data_request ∗req, unsigned may_alloc, enum _starpu_is_prefetch prefetch)
- unsigned **_starpu_driver_test_request_completion** (struct _starpu_async_channel ∗async_channel)
- void **_starpu_driver_wait_request_completion** (struct _starpu_async_channel ∗async_channel)

6.8.1 Data Structure Documentation

6.8.1.1 struct _starpu_mic_async_event

MIC needs memory_node to know which MIC is concerned. mark is used to wait asynchronous request. signal is used to test asynchronous request.

Data Fields

unsigned	memory_node	
int	mark	
uint64_t ∗	signal	

6.8.1.2 struct _starpu_disk_backend_event

Data Fields

void ∗	backend_event	

6.8.1.3 struct _starpu_disk_async_event

Data Fields

unsigned	memory_node	
struct _starpu_disk_backend_event_list ∗	requests	
void ∗	ptr	
unsigned	node	
size_t	size	
starpu_data_handle_t	handle	

6.8.1.4 union _starpu_async_channel_event

this is a structure that can be queried to see whether an asynchronous transfer has terminated or not

Data Fields

struct _starpu_async_channel_event	__unnamed__	
cudaEvent_t	cuda_event	
cl_event	opencl_event	
struct _starpu_mic_async_event	mic_event	
struct _starpu_disk_async_event	disk_event	

6.8.1.5 struct _starpu_async_channel

Data Fields

union _starpu_async_channel_event	event	
struct _starpu_node_ops ∗	node_ops	
struct _starpu_mp_node ∗	polling_node_sender	Which node to polling when needing ACK msg
struct _starpu_mp_node ∗	polling_node_receiver	
volatile int	starpu_mp_common_finished_sender	Used to know if the acknowlegdment msg is arrived from sinks
volatile int	starpu_mp_common_finished_receiver	

6.8.1.6 struct _starpu_async_channel_event.__unnamed__

Data Fields

unsigned	finished	
starpu_pthread_queue_t ∗	queue	

6.8.2 Enumeration Type Documentation

6.8.2.1 _starpu_is_prefetch

enum [_starpu_is_prefetch](#)

Enumerator

STARPU_FETCH	A task really needs it now!
STARPU_PREFETCH	It is a good idea to have it asap
STARPU_IDLEFETCH	Get this here when you have time to

6.9 data_concurrency.h File Reference

```
#include <core/jobs.h>
```

Functions

- void **_starpu_job_set_ordered_buffers** ([struct _starpu_job](#) ∗j)
- unsigned **_starpu_concurrent_data_access** ([struct _starpu_job](#) ∗j)
- void **_starpu_submit_job_enforce_arbitered_deps** ([struct _starpu_job](#) ∗j, unsigned buf, unsigned nbuffers)
- void **_starpu_enforce_data_deps_notify_job_ready_soon** ([struct _starpu_job](#) ∗j, _starpu_notify_job_↩ start_data ∗data)
- int **_starpu_notify_data_dependencies** (starpu_data_handle_t handle)
- void **_starpu_notify_arbitered_dependencies** (starpu_data_handle_t handle)
- unsigned **_starpu_attempt_to_submit_data_request_from_apps** (starpu_data_handle_t handle, enum starpu_data_access_mode mode, void(∗callback)(void ∗), void ∗argcb)
- unsigned **_starpu_attempt_to_submit_arbitered_data_request** (unsigned request_from_codelet, starpu_data_handle_t handle, enum starpu_data_access_mode mode, void(∗callback)(void ∗), void ∗argcb, [struct _starpu_job](#) ∗j, unsigned buffer_index)

6.10 data_interface.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/uthash.h>
#include <util/openmp_runtime_support.h>
```

Data Structures

- union [_starpu_interface](#)

Macros

- #define **_starpu_data_check_not_busy**(handle)
- #define **_starpu_data_is_multiformat_handle**(handle)

Functions

- void **_starpu_data_free_interfaces** (starpu_data_handle_t handle) STARPU_ATTRIBUTE_INTERNAL
- int **_starpu_data_handle_init** (starpu_data_handle_t handle, struct starpu_data_interface_ops ∗interface↩_ops, unsigned int mf_node)
- void **_starpu_data_initialize_per_worker** (starpu_data_handle_t handle)
- void **_starpu_data_interface_init** (void) STARPU_ATTRIBUTE_INTERNAL
- int **__starpu_data_check_not_busy** (starpu_data_handle_t handle) STARPU_ATTRIBUTE_INTERNAL STARPU_ATTRIBUTE_WARN_UNUSED_RESULT
- void **_starpu_data_interface_shutdown** (void) STARPU_ATTRIBUTE_INTERNAL
- void **_starpu_omp_unregister_region_handles** (struct starpu_omp_region ∗region)
- void **_starpu_omp_unregister_task_handles** (struct starpu_omp_task ∗task)
- struct starpu_data_interface_ops ∗ **_starpu_data_interface_get_ops** (unsigned interface_id)
- void **_starpu_data_register_ram_pointer** (starpu_data_handle_t handle, void ∗ptr) STARPU_ATTRIBU↩TE_INTERNAL
- void **_starpu_data_unregister_ram_pointer** (starpu_data_handle_t handle, unsigned node) STARPU_A↩TTRIBUTE_INTERNAL
- void **_starpu_data_invalidate_submit_noplan** (starpu_data_handle_t handle)

Variables

- struct starpu_data_interface_ops starpu_interface_matrix_ops
- struct starpu_data_interface_ops **starpu_interface_block_ops**
- struct starpu_data_interface_ops **starpu_interface_vector_ops**
- struct starpu_data_interface_ops **starpu_interface_csr_ops**
- struct starpu_data_interface_ops **starpu_interface_bcsr_ops**
- struct starpu_data_interface_ops **starpu_interface_variable_ops**
- struct starpu_data_interface_ops **starpu_interface_void_ops**
- struct starpu_data_interface_ops **starpu_interface_multiformat_ops**
- struct starpu_arbiter ∗ **_starpu_global_arbiter**

6.10.1 Data Structure Documentation

6.10.1.1 union _starpu_interface

Generic type representing an interface, for now it's only used before execution on message-passing devices but it can be useful in other cases.

Data Fields

struct starpu_variable_interface	variable	
struct starpu_vector_interface	vector	
struct starpu_matrix_interface	matrix	
struct starpu_block_interface	block	
struct starpu_csr_interface	csr	
struct starpu_bcsr_interface	bcsr	
struct starpu_coo_interface	coo	

6.10.2 Variable Documentation

6.10.2.1 starpu_interface_matrix_ops

```
struct starpu_data_interface_ops starpu_interface_matrix_ops
```

Some data interfaces or filters use this interface internally

6.11 data_request.h File Reference

```
#include <datawizard/coherency.h>
#include <semaphore.h>
#include <datawizard/copy_driver.h>
#include <common/list.h>
#include <common/prio_list.h>
#include <common/starpu_spinlock.h>
```

Data Structures

- struct _starpu_callback_list

Macros

- #define **MAX_PENDING_REQUESTS_PER_NODE**
- #define **MAX_PENDING_PREFETCH_REQUESTS_PER_NODE**
- #define **MAX_PENDING_IDLE_REQUESTS_PER_NODE**
- #define MAX_PUSH_TIME

6.11.1 Macro Definition Documentation

6.11.1.1 MAX_PUSH_TIME

```
#define MAX_PUSH_TIME
```
Maximum time in us that we can afford pushing requests before going back to the driver loop, e.g. for checking GPU task termination

6.12 datastats.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <stdint.h>
#include <stdlib.h>
```

Macros

- #define **_starpu_msi_cache_hit**(node)
- #define **_starpu_msi_cache_miss**(node)
- #define **_starpu_allocation_cache_hit**(node)
- #define **_starpu_data_allocation_inc_stats**(node)

Functions

- void **_starpu_datastats_init** ()
- static int **starpu_enable_stats** (void)
- void **__starpu_msi_cache_hit** (unsigned node)
- void **__starpu_msi_cache_miss** (unsigned node)
- void **_starpu_display_msi_stats** (FILE ∗stream)
- void **__starpu_allocation_cache_hit** (unsigned node STARPU_ATTRIBUTE_UNUSED)

- void **__starpu_data_allocation_inc_stats** (unsigned node STARPU_ATTRIBUTE_UNUSED)
- void **_starpu_display_alloc_cache_stats** (FILE ∗stream)

Variables

- int **_starpu_enable_stats**

6.13 datawizard.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/utils.h>
#include <datawizard/coherency.h>
#include <datawizard/filters.h>
#include <datawizard/copy_driver.h>
#include <datawizard/footprint.h>
#include <datawizard/data_request.h>
#include <datawizard/interfaces/data_interface.h>
#include <core/dependencies/implicit_data_deps.h>
```

Functions

- int **___starpu_datawizard_progress** (unsigned memory_node, unsigned may_alloc, unsigned push_↩ requests)
- int **__starpu_datawizard_progress** (unsigned may_alloc, unsigned push_requests)
- void **_starpu_datawizard_progress** (unsigned may_alloc)

6.14 debug.h File Reference

```
#include <stdio.h>
#include <stdlib.h>
#include <stdarg.h>
#include <common/config.h>
#include <core/workers.h>
```

Macros

- #define **STARPU_AYU_EVENT**
- #define **STARPU_AYU_PREINIT**()
- #define **STARPU_AYU_INIT**()
- #define **STARPU_AYU_FINISH**()
- #define **STARPU_AYU_ADDDEPENDENCY**(previous, handle, next_job)
- #define **STARPU_AYU_REMOVETASK**(job_id)
- #define **STARPU_AYU_ADDTASK**(job_id, task)
- #define **STARPU_AYU_PRERUNTASK**(job_id, workerid)
- #define **STARPU_AYU_RUNTASK**(job_id)
- #define **STARPU_AYU_POSTRUNTASK**(job_id)
- #define **STARPU_AYU_ADDTOTASKQUEUE**(job_id, worker_id)
- #define **STARPU_AYU_BARRIER**()

Functions

- void [_starpu_open_debug_logfile](void)
- void [_starpu_close_debug_logfile](void)
- void [_starpu_print_to_logfile](const char *format,...) STARPU_ATTRIBUTE_FORMAT(printf
- void **_starpu_watchdog_init** (void)
- void **_starpu_watchdog_shutdown** (void)

Variables

- void int [_starpu_use_fxt](void)

6.14.1   Function Documentation

6.14.1.1   _starpu_open_debug_logfile()

```
void _starpu_open_debug_logfile (
            void  )
```
Create a file that will contain StarPU's log

6.14.1.2   _starpu_close_debug_logfile()

```
void _starpu_close_debug_logfile (
            void  )
```
Close StarPU's log file

6.14.1.3   _starpu_print_to_logfile()

```
void _starpu_print_to_logfile (
            const char * format,
            ... )
```
Write into StarPU's log file

6.14.2   Variable Documentation

6.14.2.1   _starpu_use_fxt

```
void int _starpu_use_fxt
```
Tell gdb whether FXT is compiled in or not

6.15   detect_combined_workers.h File Reference

```
#include <starpu.h>
```

Functions

- void [_starpu_sched_find_worker_combinations](int *workerids, int nworkers)

Variables

- int **_starpu_initialized_combined_workers**

6.15.1 Function Documentation

6.15.1.1 _starpu_sched_find_worker_combinations()

```
void _starpu_sched_find_worker_combinations (
            int * workerids,
            int nworkers )
```
Initialize combined workers

6.16 disk.h File Reference

```
#include <datawizard/copy_driver.h>
#include <datawizard/malloc.h>
```

Macros

- #define **STARPU_DISK_ALL**
- #define **STARPU_DISK_NO_RECLAIM**

Functions

- void ∗ _starpu_disk_alloc (unsigned node, size_t size) STARPU_ATTRIBUTE_MALLOC
- void **_starpu_disk_free** (unsigned node, void ∗obj, size_t size)
- int _starpu_disk_read (unsigned src_node, unsigned dst_node, void ∗obj, void ∗buf, off_t offset, size_t size, struct _starpu_async_channel ∗async_channel)
- int _starpu_disk_write (unsigned src_node, unsigned dst_node, void ∗obj, void ∗buf, off_t offset, size_t size, struct _starpu_async_channel ∗async_channel)
- int **_starpu_disk_full_read** (unsigned src_node, unsigned dst_node, void ∗obj, void ∗∗ptr, size_t ∗size, struct _starpu_async_channel ∗async_channel)
- int **_starpu_disk_full_write** (unsigned src_node, unsigned dst_node, void ∗obj, void ∗ptr, size_t size, struct _starpu_async_channel ∗async_channel)
- int **_starpu_disk_copy** (unsigned node_src, void ∗obj_src, off_t offset_src, unsigned node_dst, void ∗obj↩_dst, off_t offset_dst, size_t size, struct _starpu_async_channel ∗async_channel)
- void starpu_disk_wait_request (struct _starpu_async_channel ∗async_channel)
- int starpu_disk_test_request (struct _starpu_async_channel ∗async_channel)
- void **starpu_disk_free_request** (struct _starpu_async_channel ∗async_channel)
- int _starpu_disk_can_copy (unsigned node1, unsigned node2)
- void _starpu_set_disk_flag (unsigned node, int flag)
- int **_starpu_get_disk_flag** (unsigned node)
- void _starpu_disk_unregister (void)
- void **_starpu_swap_init** (void)

6.16.1 Function Documentation

6.16.1.1 _starpu_disk_alloc()

```
void* _starpu_disk_alloc (
            unsigned node,
            size_t size )
```
interface to manipulate memory disk

6.16.1.2 _starpu_disk_read()

```
int _starpu_disk_read (
            unsigned src_node,
            unsigned dst_node,
            void * obj,
            void * buf,
            off_t offset,
            size_t size,
            struct _starpu_async_channel * async_channel )
```
src_node is a disk node, dst_node is for the moment the STARPU_MAIN_RAM

6.16.1.3 _starpu_disk_write()

```
int _starpu_disk_write (
            unsigned src_node,
            unsigned dst_node,
            void * obj,
            void * buf,
            off_t offset,
            size_t size,
            struct _starpu_async_channel * async_channel )
```
src_node is for the moment the STARU_MAIN_RAM, dst_node is a disk node

6.16.1.4 starpu_disk_wait_request()

```
void starpu_disk_wait_request (
            struct _starpu_async_channel * async_channel )
```
force the request to compute

6.16.1.5 starpu_disk_test_request()

```
int starpu_disk_test_request (
            struct _starpu_async_channel * async_channel )
```
return 1 if the request is finished, 0 if not finished

6.16.1.6 _starpu_disk_can_copy()

```
int _starpu_disk_can_copy (
            unsigned node1,
            unsigned node2 )
```
interface to compare memory disk

6.16.1.7 _starpu_set_disk_flag()

```
void _starpu_set_disk_flag (
            unsigned node,
            int flag )
```
change disk flag

6.16.1.8 _starpu_disk_unregister()

```
void _starpu_disk_unregister (
            void  )
```
unregister disk

6.17   disk_unistd_global.h File Reference

```
#include <fcntl.h>
```

Data Structures

- struct starpu_unistd_global_obj

Macros

- #define **O_BINARY**
- #define **STARPU_UNISTD_USE_COPY**

Functions

- void ∗ **starpu_unistd_global_alloc** (struct starpu_unistd_global_obj ∗obj, void ∗base, size_t size)
- void **starpu_unistd_global_free** (void ∗base, void ∗obj, size_t size)
- void ∗ **starpu_unistd_global_open** (struct starpu_unistd_global_obj ∗obj, void ∗base, void ∗pos, size_t size)
- void **starpu_unistd_global_close** (void ∗base, void ∗obj, size_t size)
- int **starpu_unistd_global_read** (void ∗base, void ∗obj, void ∗buf, off_t offset, size_t size)
- int **starpu_unistd_global_write** (void ∗base, void ∗obj, const void ∗buf, off_t offset, size_t size)
- void ∗ **starpu_unistd_global_plug** (void ∗parameter, starpu_ssize_t size)
- void **starpu_unistd_global_unplug** (void ∗base)
- int **_starpu_get_unistd_global_bandwidth_between_disk_and_main_ram** (unsigned node, void ∗base)
- void ∗ **starpu_unistd_global_async_read** (void ∗base, void ∗obj, void ∗buf, off_t offset, size_t size)
- void ∗ **starpu_unistd_global_async_write** (void ∗base, void ∗obj, void ∗buf, off_t offset, size_t size)
- void ∗ **starpu_unistd_global_async_full_write** (void ∗base, void ∗obj, void ∗ptr, size_t size)
- void ∗ **starpu_unistd_global_async_full_read** (void ∗base, void ∗obj, void ∗∗ptr, size_t ∗size, unsigned dst_node)
- void **starpu_unistd_global_wait_request** (void ∗async_channel)
- int **starpu_unistd_global_test_request** (void ∗async_channel)
- void **starpu_unistd_global_free_request** (void ∗async_channel)
- int **starpu_unistd_global_full_read** (void ∗base, void ∗obj, void ∗∗ptr, size_t ∗size, unsigned dst_node)
- int **starpu_unistd_global_full_write** (void ∗base, void ∗obj, void ∗ptr, size_t size)

6.17.1   Data Structure Documentation

6.17.1.1   struct starpu_unistd_global_obj

Data Fields

int	descriptor	
char ∗	path	
size_t	size	
int	flags	
starpu_pthread_mutex_t	mutex	

6.18   driver_common.h File Reference

```
#include <starpu.h>
#include <starpu_util.h>
#include <core/jobs.h>
#include <common/utils.h>
```

Functions

- void **_starpu_driver_start_job** (struct _starpu_worker ∗args, struct _starpu_job ∗j, struct starpu_↩
  perfmodel_arch ∗perf_arch, int rank, int profiling)
- void **_starpu_driver_end_job** (struct _starpu_worker ∗args, struct _starpu_job ∗j, struct starpu_perfmodel↩
  _arch ∗perf_arch, int rank, int profiling)
- void **_starpu_driver_update_job_feedback** (struct _starpu_job ∗j, struct _starpu_worker ∗worker_args,
  struct starpu_perfmodel_arch ∗perf_arch, int profiling)
- struct starpu_task ∗ **_starpu_get_worker_task** (struct _starpu_worker ∗args, int workerid, unsigned
  memnode)
- int **_starpu_get_multi_worker_task** (struct _starpu_worker ∗workers, struct starpu_task ∗∗tasks, int
  nworker, unsigned memnode)

6.19 driver_cpu.h File Reference

```
#include <common/config.h>
#include <datawizard/node_ops.h>
```

Functions

- void ∗ **_starpu_cpu_worker** (void ∗)
- int **_starpu_cpu_copy_interface** (starpu_data_handle_t handle, void ∗src_interface, unsigned src_node,
  void ∗dst_interface, unsigned dst_node, struct _starpu_data_request ∗req)
- int **_starpu_cpu_copy_data** (uintptr_t src_ptr, size_t src_offset, unsigned src_node, uintptr_t dst_ptr, size↩
  _t dst_offset, unsigned dst_node, size_t ssize, struct _starpu_async_channel ∗async_channel)
- int **_starpu_cpu_is_direct_access_supported** (unsigned node, unsigned handling_node)
- uintptr_t **_starpu_cpu_malloc_on_node** (unsigned dst_node, size_t size, int flags)
- void **_starpu_cpu_free_on_node** (unsigned dst_node, uintptr_t addr, size_t size, int flags)

Variables

- struct _starpu_driver_ops **_starpu_driver_cpu_ops**
- struct _starpu_node_ops **_starpu_driver_cpu_node_ops**

6.20 driver_cuda.h File Reference

```
#include <common/config.h>
#include <cuda.h>
#include <cuda_runtime_api.h>
#include <cublas.h>
#include <starpu.h>
#include <core/workers.h>
#include <datawizard/node_ops.h>
```

Functions

- void **_starpu_cuda_init** (void)
- unsigned **_starpu_get_cuda_device_count** (void)
- void **_starpu_cuda_discover_devices** (struct _starpu_machine_config ∗)
- void **_starpu_init_cuda** (void)
- void ∗ **_starpu_cuda_worker** (void ∗)
- cudaStream_t **starpu_cuda_get_local_in_transfer_stream** (void)

- cudaStream_t **starpu_cuda_get_in_transfer_stream** (unsigned dst_node)
- cudaStream_t **starpu_cuda_get_local_out_transfer_stream** (void)
- cudaStream_t **starpu_cuda_get_out_transfer_stream** (unsigned src_node)
- cudaStream_t **starpu_cuda_get_peer_transfer_stream** (unsigned src_node, unsigned dst_node)
- unsigned **_starpu_cuda_test_request_completion** (struct _starpu_async_channel *async_channel)
- void **_starpu_cuda_wait_request_completion** (struct _starpu_async_channel *async_channel)
- int **_starpu_cuda_copy_interface_from_cpu_to_cuda** (starpu_data_handle_t handle, void *src_interface, unsigned src_node, void *dst_interface, unsigned dst_node, struct _starpu_data_request *req)
- int **_starpu_cuda_copy_interface_from_cuda_to_cuda** (starpu_data_handle_t handle, void *src_↩ interface, unsigned src_node, void *dst_interface, unsigned dst_node, struct _starpu_data_request *req)
- int **_starpu_cuda_copy_interface_from_cuda_to_cpu** (starpu_data_handle_t handle, void *src_interface, unsigned src_node, void *dst_interface, unsigned dst_node, struct _starpu_data_request *req)
- int **_starpu_cuda_copy_data_from_cuda_to_cuda** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel *async_↩ channel)
- int **_starpu_cuda_copy_data_from_cuda_to_cpu** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel *async_↩ channel)
- int **_starpu_cuda_copy_data_from_cpu_to_cuda** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel *async_↩ channel)
- int **_starpu_cuda_copy2d_data_from_cuda_to_cuda** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks, size_t ld_src, size_t ld_dst, struct _starpu_async_channel *async_channel)
- int **_starpu_cuda_copy2d_data_from_cuda_to_cpu** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks, size_t ld_src, size_t ld_dst, struct _starpu_async_channel *async_channel)
- int **_starpu_cuda_copy2d_data_from_cpu_to_cuda** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks, size_t ld_src, size_t ld_dst, struct _starpu_async_channel *async_channel)
- int **_starpu_cuda_copy3d_data_from_cuda_to_cuda** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks_1, size_t ld1_src, size_t ld1_dst, size_t numblocks_2, size_t ld2_src, size_t ld2_dst, struct _starpu_async_channel *async_↩ channel)
- int **_starpu_cuda_copy3d_data_from_cuda_to_cpu** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks_1, size_t ld1_src, size_t ld1_dst, size_t numblocks_2, size_t ld2_src, size_t ld2_dst, struct _starpu_async_channel *async_↩ channel)
- int **_starpu_cuda_copy3d_data_from_cpu_to_cuda** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks_1, size_t ld1_src, size_t ld1_dst, size_t numblocks_2, size_t ld2_src, size_t ld2_dst, struct _starpu_async_channel *async_↩ channel)
- int **_starpu_cuda_is_direct_access_supported** (unsigned node, unsigned handling_node)
- uintptr_t **_starpu_cuda_malloc_on_node** (unsigned dst_node, size_t size, int flags)
- void **_starpu_cuda_free_on_node** (unsigned dst_node, uintptr_t addr, size_t size, int flags)

Variables

- struct _starpu_driver_ops **_starpu_driver_cuda_ops**
- struct _starpu_node_ops **_starpu_driver_cuda_node_ops**
- int **_starpu_cuda_bus_ids** [STARPU_MAXCUDADEVS+STARPU_MAXNUMANODES][STARPU_MAX↩ CUDADEVS+STARPU_MAXNUMANODES]

6.21 driver_disk.h File Reference

```
#include <datawizard/node_ops.h>
```

Functions

- int **_starpu_disk_copy_src_to_disk** (void ∗src, unsigned src_node, void ∗dst, size_t dst_offset, unsigned dst_node, size_t size, void ∗async_channel)
- int **_starpu_disk_copy_disk_to_src** (void ∗src, size_t src_offset, unsigned src_node, void ∗dst, unsigned dst_node, size_t size, void ∗async_channel)
- int **_starpu_disk_copy_disk_to_disk** (void ∗src, size_t src_offset, unsigned src_node, void ∗dst, size_t dst_offset, unsigned dst_node, size_t size, void ∗async_channel)
- unsigned **_starpu_disk_test_request_completion** ([struct _starpu_async_channel](#) ∗async_channel)
- void **_starpu_disk_wait_request_completion** ([struct _starpu_async_channel](#) ∗async_channel)
- int **_starpu_disk_copy_interface_from_disk_to_cpu** (starpu_data_handle_t handle, void ∗src_interface, unsigned src_node, void ∗dst_interface, unsigned dst_node, [struct](#) _starpu_data_request ∗req)
- int **_starpu_disk_copy_interface_from_disk_to_disk** (starpu_data_handle_t handle, void ∗src_interface, unsigned src_node, void ∗dst_interface, unsigned dst_node, [struct](#) _starpu_data_request ∗req)
- int **_starpu_disk_copy_interface_from_cpu_to_disk** (starpu_data_handle_t handle, void ∗src_interface, unsigned src_node, void ∗dst_interface, unsigned dst_node, [struct](#) _starpu_data_request ∗req)
- int **_starpu_disk_copy_data_from_disk_to_cpu** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, [struct _starpu_async_channel](#) ∗async_↩ channel)
- int **_starpu_disk_copy_data_from_disk_to_disk** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, [struct _starpu_async_channel](#) ∗async_↩ channel)
- int **_starpu_disk_copy_data_from_cpu_to_disk** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, [struct _starpu_async_channel](#) ∗async_↩ channel)
- int **_starpu_disk_is_direct_access_supported** (unsigned node, unsigned handling_node)
- uintptr_t **_starpu_disk_malloc_on_node** (unsigned dst_node, size_t size, int flags)
- void **_starpu_disk_free_on_node** (unsigned dst_node, uintptr_t addr, size_t size, int flags)

Variables

- [struct _starpu_node_ops](#) **_starpu_driver_disk_node_ops**

6.22 driver_mic_common.h File Reference

```
#include <common/config.h>
#include <source/COIProcess_source.h>
```

Data Structures

- struct [_starpu_mic_free_command](#)

Macros

- #define **STARPU_TO_MIC_ID**(id)
- #define **STARPU_MIC_PORTS_BEGIN**
- #define **STARPU_MIC_SOURCE_PORT_NUMBER**
- #define **STARPU_MIC_SINK_PORT_NUMBER**(id)
- #define **STARPU_MIC_SOURCE_DT_PORT_NUMBER**
- #define **STARPU_MIC_SINK_DT_PORT_NUMBER**(id)
- #define **STARPU_MIC_SINK_SINK_DT_PORT_NUMBER**(me, peer_id)
- #define **STARPU_MIC_PAGE_SIZE**
- #define **STARPU_MIC_GET_PAGE_SIZE_MULTIPLE**(size)
- #define **STARPU_MIC_COMMON_REPORT_SCIF_ERROR**(status)

Functions

- void **_starpu_mic_common_report_scif_error** (const char *func, const char *file, int line, const int status)
- int **_starpu_mic_common_recv_is_ready** (const struct _starpu_mp_node *mp_node)
- void **_starpu_mic_common_send** (const struct _starpu_mp_node *node, void *msg, int len)
- void **_starpu_mic_common_recv** (const struct _starpu_mp_node *node, void *msg, int len)
- void **_starpu_mic_common_dt_send** (const struct _starpu_mp_node *node, void *msg, int len, void *event)
- void **_starpu_mic_common_dt_recv** (const struct _starpu_mp_node *node, void *msg, int len, void *event)
- void **_starpu_mic_common_connect** (scif_epd_t *endpoint, uint16_t remote_node, COIPROCESS process, uint16_t local_port_number, uint16_t remote_port_number)
- void **_starpu_mic_common_accept** (scif_epd_t *endpoint, uint16_t port_number)

6.22.1 Data Structure Documentation

6.22.1.1 struct _starpu_mic_free_command

Data Fields

void *	addr	
size_t	size	

6.23 driver_mic_sink.h File Reference

```
#include <common/config.h>
#include <scif.h>
#include <drivers/mp_common/mp_common.h>
#include <drivers/mp_common/sink_common.h>
```

Macros

- #define **STARPU_MIC_SINK_REPORT_ERROR**(status)

Functions

- void **_starpu_mic_sink_report_error** (const char *func, const char *file, const int line, const int status)
- void **_starpu_mic_sink_init** (struct _starpu_mp_node *node)
- void **_starpu_mic_sink_launch_workers** (struct _starpu_mp_node *node)
- void **_starpu_mic_sink_deinit** (struct _starpu_mp_node *node)
- void **_starpu_mic_sink_allocate** (const struct _starpu_mp_node *mp_node, void *arg, int arg_size)
- void **_starpu_mic_sink_free** (const struct _starpu_mp_node *mp_node STARPU_ATTRIBUTE_UNUSED, void *arg, int arg_size)
- void **_starpu_mic_sink_bind_thread** (const struct _starpu_mp_node *mp_node STARPU_ATTRIBUTE↩_UNUSED, int coreid, int *core_table, int nb_core)

Variables

- void(*)(void) **_starpu_mic_sink_lookup** (const struct _starpu_mp_node *node STARPU_ATTRIBUTE_↩UNUSED, char *func_name)

6.24 driver_mic_source.h File Reference

```
#include <starpu_mic.h>
#include <common/config.h>
#include <source/COIProcess_source.h>
```

```
#include <source/COIEngine_source.h>
#include <core/workers.h>
#include <drivers/mp_common/mp_common.h>
#include <datawizard/node_ops.h>
```

Macros

- #define **STARPU_MIC_REQUEST_COMPLETE**
- #define **STARPU_MIC_SRC_REPORT_COI_ERROR**(status)
- #define **STARPU_MIC_SRC_REPORT_SCIF_ERROR**(status)

Functions

- struct _starpu_mp_node ∗ **_starpu_mic_src_get_actual_thread_mp_node** ()
- struct _starpu_mp_node ∗ **_starpu_mic_src_get_mp_node_from_memory_node** (int memory_node)
- int **_starpu_mic_src_register_kernel** (starpu_mic_func_symbol_t ∗symbol, const char ∗func_name)
- starpu_mic_kernel_t **_starpu_mic_src_get_kernel** (starpu_mic_func_symbol_t symbol)
- void **_starpu_mic_src_report_coi_error** (const char ∗func, const char ∗file, int line, const COIRESULT status)
- void **_starpu_mic_src_report_scif_error** (const char ∗func, const char ∗file, int line, const int status)
- unsigned **_starpu_mic_src_get_device_count** (void)
- starpu_mic_kernel_t **_starpu_mic_src_get_kernel_from_codelet** (struct starpu_codelet ∗cl, unsigned nimpl)
- void **_starpu_mic_src_init** (struct _starpu_mp_node ∗node)
- void **_starpu_mic_clear_kernels** (void)
- void **_starpu_mic_src_deinit** (struct _starpu_mp_node ∗node)
- size_t **_starpu_mic_get_global_mem_size** (int devid)
- size_t **_starpu_mic_get_free_mem_size** (int devid)
- int **_starpu_mic_allocate_memory** (void ∗∗addr, size_t size, unsigned memory_node)
- void **_starpu_mic_free_memory** (void ∗addr, size_t size, unsigned memory_node)
- int **_starpu_mic_copy_ram_to_mic** (void ∗src, unsigned src_node STARPU_ATTRIBUTE_UNUSED, void ∗dst, unsigned dst_node, size_t size)
- int **_starpu_mic_copy_mic_to_ram** (void ∗src, unsigned src_node, void ∗dst, unsigned dst_node STAR↩PU_ATTRIBUTE_UNUSED, size_t size)
- int **_starpu_mic_copy_ram_to_mic_async** (void ∗src, unsigned src_node STARPU_ATTRIBUTE_UNU↩SED, void ∗dst, unsigned dst_node, size_t size)
- int **_starpu_mic_copy_mic_to_ram_async** (void ∗src, unsigned src_node, void ∗dst, unsigned dst_node STARPU_ATTRIBUTE_UNUSED, size_t size)
- int **_starpu_mic_init_event** (struct _starpu_mic_async_event ∗event, unsigned memory_node)
- void ∗ **_starpu_mic_src_worker** (void ∗arg)
- unsigned **_starpu_mic_test_request_completion** (struct _starpu_async_channel ∗async_channel)
- void **_starpu_mic_wait_request_completion** (struct _starpu_async_channel ∗async_channel)
- int **_starpu_mic_copy_data_from_mic_to_cpu** (starpu_data_handle_t handle, void ∗src_interface, un-signed src_node, void ∗dst_interface, unsigned dst_node, struct _starpu_data_request ∗req)
- int **_starpu_mic_copy_data_from_cpu_to_mic** (starpu_data_handle_t handle, void ∗src_interface, un-signed src_node, void ∗dst_interface, unsigned dst_node, struct _starpu_data_request ∗req)
- int **_starpu_mic_copy_interface_from_mic_to_cpu** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel ∗async_↩channel)
- int **_starpu_mic_copy_interface_from_cpu_to_mic** (uintptr_t src, size_t src_offset, unsigned src_node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel ∗async_↩channel)
- int **_starpu_mic_is_direct_access_supported** (unsigned node, unsigned handling_node)
- uintptr_t **_starpu_mic_malloc_on_node** (unsigned dst_node, size_t size, int flags)
- void **_starpu_mic_free_on_node** (unsigned dst_node, uintptr_t addr, size_t size, int flags)

Variables

- struct _starpu_node_ops **_starpu_driver_mic_node_ops**
- struct _starpu_mp_node * _starpu_mic_nodes [STARPU_MAXMICDEVS]
- struct _starpu_mic_async_event * **event**
- void(∗)(void) **_starpu_mic_src_get_kernel_from_job** (const struct _starpu_mp_node ∗node STARPU_↩
  ATTRIBUTE_UNUSED, struct _starpu_job ∗j)

6.24.1 Variable Documentation

6.24.1.1 _starpu_mic_nodes

```
struct _starpu_mp_node* _starpu_mic_nodes[STARPU_MAXMICDEVS]
```
Array of structures containing all the informations useful to send and receive informations with devices

6.25 driver_mpi_common.h File Reference

```
#include <drivers/mp_common/mp_common.h>
#include <drivers/mpi/driver_mpi_source.h>
```

6.26 driver_mpi_sink.h File Reference

```
#include <drivers/mp_common/sink_common.h>
```

6.27 driver_mpi_source.h File Reference

```
#include <drivers/mp_common/mp_common.h>
#include <starpu_mpi_ms.h>
#include <datawizard/node_ops.h>
```

6.28 driver_opencl.h File Reference

```
#include <CL/cl.h>
#include <core/workers.h>
#include <datawizard/node_ops.h>
```

Macros

- #define **_GNU_SOURCE**
- #define **CL_TARGET_OPENCL_VERSION**

Functions

- void **_starpu_opencl_discover_devices** (struct _starpu_machine_config ∗config)
- unsigned **_starpu_opencl_get_device_count** (void)
- void **_starpu_opencl_init** (void)
- void ∗ **_starpu_opencl_worker** (void ∗)
- int **_starpu_run_opencl** (struct _starpu_worker ∗)
- int **_starpu_opencl_driver_init** (struct _starpu_worker ∗)

- int **_starpu_opencl_driver_run_once** (struct _starpu_worker ∗)
- int **_starpu_opencl_driver_deinit** (struct _starpu_worker ∗)
- int **_starpu_opencl_init_context** (int devid)
- int **_starpu_opencl_deinit_context** (int devid)
- cl_device_type **_starpu_opencl_get_device_type** (int devid)
- unsigned **_starpu_opencl_test_request_completion** (struct _starpu_async_channel ∗async_channel)
- void **_starpu_opencl_wait_request_completion** (struct _starpu_async_channel ∗async_channel)
- int **_starpu_opencl_copy_interface_from_opencl_to_opencl** (starpu_data_handle_t handle, void ∗src↩
  _interface, unsigned src_node, void ∗dst_interface, unsigned dst_node, struct _starpu_data_request ∗req)
- int **_starpu_opencl_copy_interface_from_opencl_to_cpu** (starpu_data_handle_t handle, void ∗src↩
  interface, unsigned src_node, void ∗dst_interface, unsigned dst_node, struct _starpu_data_request ∗req)
- int **_starpu_opencl_copy_interface_from_cpu_to_opencl** (starpu_data_handle_t handle, void ∗src↩
  interface, unsigned src_node, void ∗dst_interface, unsigned dst_node, struct _starpu_data_request ∗req)
- int **_starpu_opencl_copy_data_from_opencl_to_cpu** (uintptr_t src, size_t src_offset, unsigned src_node,
  uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel ∗async↩
  _channel)
- int **_starpu_opencl_copy_data_from_opencl_to_opencl** (uintptr_t src, size_t src_offset, unsigned src↩
  _node, uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel ∗async↩
  _channel)
- int **_starpu_opencl_copy_data_from_cpu_to_opencl** (uintptr_t src, size_t src_offset, unsigned src_node,
  uintptr_t dst, size_t dst_offset, unsigned dst_node, size_t size, struct _starpu_async_channel ∗async↩
  channel)
- int **_starpu_opencl_is_direct_access_supported** (unsigned node, unsigned handling_node)
- uintptr_t **_starpu_opencl_malloc_on_node** (unsigned dst_node, size_t size, int flags)
- void **_starpu_opencl_free_on_node** (unsigned dst_node, uintptr_t addr, size_t size, int flags)

Variables

- struct _starpu_node_ops **_starpu_driver_opencl_node_ops**
- struct _starpu_driver_ops **_starpu_driver_opencl_ops**
- char ∗ **_starpu_opencl_program_dir**

6.29 driver_opencl_utils.h File Reference

Macros

- #define **_STARPU_OPENCL_PLATFORM_MAX**

Functions

- char ∗ **_starpu_opencl_get_device_type_as_string** (int id)

6.30 drivers.h File Reference

Data Structures

- struct _starpu_driver_ops

6.31 errorcheck.h File Reference

```
#include <starpu.h>
```

Enumerations

- enum _starpu_worker_status {
  STATUS_INVALID, STATUS_UNKNOWN, STATUS_INITIALIZING, STATUS_EXECUTING,
  STATUS_CALLBACK, STATUS_SCHEDULING, STATUS_WAITING, STATUS_SLEEPING_SCHEDULI↵
  NG,
  STATUS_SLEEPING }

Functions

- void _starpu_set_worker_status (struct _starpu_worker ∗worker, enum _starpu_worker_status st)
- void **_starpu_set_local_worker_status** (enum _starpu_worker_status st)
- enum _starpu_worker_status _starpu_get_local_worker_status (void)
- unsigned _starpu_worker_may_perform_blocking_calls (void)

6.31.1 Enumeration Type Documentation

6.31.1.1 _starpu_worker_status

enum _starpu_worker_status
This type describes in which state a worker may be.

Enumerator

STATUS_INVALID	invalid status (for instance if we request the status of some thread that is not controlled by StarPU
STATUS_UNKNOWN	everything that does not fit the other status
STATUS_INITIALIZING	during the initialization
STATUS_EXECUTING	during the execution of a codelet
STATUS_CALLBACK	during the execution of the callback
STATUS_SCHEDULING	while executing the scheduler code
STATUS_WAITING	while waiting for a data transfer
STATUS_SLEEPING_SCHEDULING	while sleeping because there is nothing to do, but looking for tasks to do
STATUS_SLEEPING	while sleeping because there is nothing to do, and not even scheduling

6.31.2 Function Documentation

6.31.2.1 _starpu_set_worker_status()

void _starpu_set_worker_status (
            struct _starpu_worker ∗ *worker,*
            enum _starpu_worker_status *st*)
Specify what the local worker is currently doing (eg. executing a callback). This permits to detect if this is legal to do a blocking call for instance.

6.31.2.2 _starpu_get_local_worker_status()

enum _starpu_worker_status _starpu_get_local_worker_status (
            void)
Indicate what type of operation the worker is currently doing.

6.31.2.3 _starpu_worker_may_perform_blocking_calls()

```
unsigned _starpu_worker_may_perform_blocking_calls (
            void )
```

It is forbidden to do blocking calls during some operations such as callback or during the execution of a task. This function indicates whether it is legal to call a blocking operation in the current context.

6.32 fifo_queues.h File Reference

```
#include <starpu.h>
#include <core/task.h>
```

Data Structures

- struct _starpu_fifo_taskq

Functions

- struct _starpu_fifo_taskq ∗ **_starpu_create_fifo** (void) STARPU_ATTRIBUTE_MALLOC
- void **_starpu_destroy_fifo** (struct _starpu_fifo_taskq ∗fifo)
- int **_starpu_fifo_empty** (struct _starpu_fifo_taskq ∗fifo)
- double **_starpu_fifo_get_exp_len_prev_task_list** (struct _starpu_fifo_taskq ∗fifo_queue, struct starpu_↩
task ∗task, int workerid, int nimpl, int ∗fifo_ntasks)
- int **_starpu_fifo_push_sorted_task** (struct _starpu_fifo_taskq ∗fifo_queue, struct starpu_task ∗task)
- int **_starpu_fifo_push_task** (struct _starpu_fifo_taskq ∗fifo, struct starpu_task ∗task)
- int **_starpu_fifo_push_back_task** (struct _starpu_fifo_taskq ∗fifo_queue, struct starpu_task ∗task)
- int **_starpu_fifo_pop_this_task** (struct _starpu_fifo_taskq ∗fifo_queue, int workerid, struct starpu_task
∗task)
- struct starpu_task ∗ **_starpu_fifo_pop_task** (struct _starpu_fifo_taskq ∗fifo, int workerid)
- struct starpu_task ∗ **_starpu_fifo_pop_local_task** (struct _starpu_fifo_taskq ∗fifo)
- struct starpu_task ∗ **_starpu_fifo_pop_every_task** (struct _starpu_fifo_taskq ∗fifo, int workerid)
- int **_starpu_normalize_prio** (int priority, int num_priorities, unsigned sched_ctx_id)
- int **_starpu_count_non_ready_buffers** (struct starpu_task ∗task, unsigned worker)
- size_t **_starpu_size_non_ready_buffers** (struct starpu_task ∗task, unsigned worker)
- struct starpu_task ∗ **_starpu_fifo_pop_first_ready_task** (struct _starpu_fifo_taskq ∗fifo_queue, unsigned
workerid, int num_priorities)

6.32.1 Data Structure Documentation

6.32.1.1 struct _starpu_fifo_taskq

Data Fields

struct starpu_task_list	taskq	the actual list
unsigned	ntasks	the number of tasks currently in the queue
unsigned ∗	ntasks_per_priority	the number of tasks currently in the queue corresponding to each priority
unsigned	nprocessed	the number of tasks that were processed
double	exp_start	only meaningful if the queue is only used by a single worker
double	exp_end	Expected start date of next item to do in the queue (i.e. not started yet). This is thus updated when we start it.
double	exp_len	Expected end date of last task in the queue
double ∗	exp_len_per_priority	Expected duration of the set of tasks in the queue
double	pipeline_len	Expected duration of the set of tasks in the queue corresponding to each priority

6.33 filters.h File Reference

```
#include <stdarg.h>
#include <datawizard/coherency.h>
#include <datawizard/memalloc.h>
#include <starpu.h>
#include <common/config.h>
```

Functions

- void _starpu_data_partition_access_submit (starpu_data_handle_t target, int write)

6.33.1 Function Documentation

6.33.1.1 _starpu_data_partition_access_submit()

```
void _starpu_data_partition_access_submit (
            starpu_data_handle_t target,
            int write )
```
submit asynchronous unpartitioning / partitioning to make target active read-only or read-write

6.34 footprint.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <core/jobs.h>
```

Functions

- uint32_t _starpu_compute_buffers_footprint (struct starpu_perfmodel ∗model, struct starpu_perfmodel_arch ∗arch, unsigned nimpl, struct _starpu_job ∗j)
- uint32_t _starpu_compute_data_footprint (starpu_data_handle_t handle)
- uint32_t _starpu_compute_data_alloc_footprint (starpu_data_handle_t handle)

6.34.1 Function Documentation

6.34.1.1 _starpu_compute_buffers_footprint()

```
uint32_t _starpu_compute_buffers_footprint (
            struct starpu_perfmodel * model,
            struct starpu_perfmodel_arch * arch,
            unsigned nimpl,
            struct _starpu_job * j )
```
Compute the footprint that characterizes the job and cache it into the job structure.

6.34.1.2 _starpu_compute_data_footprint()

```
uint32_t _starpu_compute_data_footprint (
            starpu_data_handle_t handle )
```
Compute the footprint that characterizes the layout of the data handle.

6.34.1.3 _starpu_compute_data_alloc_footprint()

```
uint32_t _starpu_compute_data_alloc_footprint (
            starpu_data_handle_t handle )
```
Compute the footprint that characterizes the allocation of the data handle.

6.35 fxt.h File Reference

```
#include <string.h>
#include <sys/types.h>
#include <stdlib.h>
#include <common/config.h>
#include <common/utils.h>
#include <starpu.h>
```

Macros

- #define **_GNU_SOURCE**
- #define **_STARPU_FUT_APPS_KEY**
- #define **_STARPU_FUT_CPU_KEY**
- #define **_STARPU_FUT_CUDA_KEY**
- #define **_STARPU_FUT_OPENCL_KEY**
- #define **_STARPU_FUT_MIC_KEY**
- #define **_STARPU_FUT_MPI_KEY**
- #define **_STARPU_FUT_WORKER_INIT_START**
- #define **_STARPU_FUT_WORKER_INIT_END**
- #define **_STARPU_FUT_START_CODELET_BODY**
- #define **_STARPU_FUT_END_CODELET_BODY**
- #define **_STARPU_FUT_JOB_PUSH**
- #define **_STARPU_FUT_JOB_POP**
- #define **_STARPU_FUT_UPDATE_TASK_CNT**
- #define **_STARPU_FUT_START_FETCH_INPUT_ON_TID**
- #define **_STARPU_FUT_END_FETCH_INPUT_ON_TID**
- #define **_STARPU_FUT_START_PUSH_OUTPUT_ON_TID**
- #define **_STARPU_FUT_END_PUSH_OUTPUT_ON_TID**
- #define **_STARPU_FUT_TAG**
- #define **_STARPU_FUT_TAG_DEPS**
- #define **_STARPU_FUT_TASK_DEPS**
- #define **_STARPU_FUT_DATA_COPY**
- #define **_STARPU_FUT_WORK_STEALING**
- #define **_STARPU_FUT_WORKER_DEINIT_START**
- #define **_STARPU_FUT_WORKER_DEINIT_END**
- #define **_STARPU_FUT_WORKER_SLEEP_START**
- #define **_STARPU_FUT_WORKER_SLEEP_END**
- #define **_STARPU_FUT_TASK_SUBMIT**
- #define **_STARPU_FUT_CODELET_DATA_HANDLE**
- #define **_STARPU_FUT_MODEL_NAME**
- #define **_STARPU_FUT_DATA_NAME**
- #define **_STARPU_FUT_DATA_COORDINATES**
- #define **_STARPU_FUT_HANDLE_DATA_UNREGISTER**
- #define **_STARPU_FUT_USER_DEFINED_START**
- #define **_STARPU_FUT_USER_DEFINED_END**
- #define **_STARPU_FUT_NEW_MEM_NODE**
- #define **_STARPU_FUT_START_CALLBACK**

- #define **_STARPU_FUT_END_CALLBACK**
- #define **_STARPU_FUT_TASK_DONE**
- #define **_STARPU_FUT_TAG_DONE**
- #define **_STARPU_FUT_START_ALLOC**
- #define **_STARPU_FUT_END_ALLOC**
- #define **_STARPU_FUT_START_ALLOC_REUSE**
- #define **_STARPU_FUT_END_ALLOC_REUSE**
- #define **_STARPU_FUT_USED_MEM**
- #define **_STARPU_FUT_TASK_NAME**
- #define **_STARPU_FUT_DATA_WONT_USE**
- #define **_STARPU_FUT_TASK_COLOR**
- #define **_STARPU_FUT_DATA_DOING_WONT_USE**
- #define **_STARPU_FUT_START_MEMRECLAIM**
- #define **_STARPU_FUT_END_MEMRECLAIM**
- #define **_STARPU_FUT_START_DRIVER_COPY**
- #define **_STARPU_FUT_END_DRIVER_COPY**
- #define **_STARPU_FUT_START_DRIVER_COPY_ASYNC**
- #define **_STARPU_FUT_END_DRIVER_COPY_ASYNC**
- #define **_STARPU_FUT_START_PROGRESS_ON_TID**
- #define **_STARPU_FUT_END_PROGRESS_ON_TID**
- #define **_STARPU_FUT_USER_EVENT**
- #define **_STARPU_FUT_SET_PROFILING**
- #define **_STARPU_FUT_TASK_WAIT_FOR_ALL**
- #define **_STARPU_FUT_EVENT**
- #define **_STARPU_FUT_THREAD_EVENT**
- #define **_STARPU_FUT_CODELET_DETAILS**
- #define **_STARPU_FUT_CODELET_DATA**
- #define **_STARPU_FUT_LOCKING_MUTEX**
- #define **_STARPU_FUT_MUTEX_LOCKED**
- #define **_STARPU_FUT_UNLOCKING_MUTEX**
- #define **_STARPU_FUT_MUTEX_UNLOCKED**
- #define **_STARPU_FUT_TRYLOCK_MUTEX**
- #define **_STARPU_FUT_RDLOCKING_RWLOCK**
- #define **_STARPU_FUT_RWLOCK_RDLOCKED**
- #define **_STARPU_FUT_WRLOCKING_RWLOCK**
- #define **_STARPU_FUT_RWLOCK_WRLOCKED**
- #define **_STARPU_FUT_UNLOCKING_RWLOCK**
- #define **_STARPU_FUT_RWLOCK_UNLOCKED**
- #define **_STARPU_FUT_LOCKING_SPINLOCK**
- #define **_STARPU_FUT_SPINLOCK_LOCKED**
- #define **_STARPU_FUT_UNLOCKING_SPINLOCK**
- #define **_STARPU_FUT_SPINLOCK_UNLOCKED**
- #define **_STARPU_FUT_TRYLOCK_SPINLOCK**
- #define **_STARPU_FUT_COND_WAIT_BEGIN**
- #define **_STARPU_FUT_COND_WAIT_END**
- #define **_STARPU_FUT_MEMORY_FULL**
- #define **_STARPU_FUT_DATA_LOAD**
- #define **_STARPU_FUT_START_UNPARTITION_ON_TID**
- #define **_STARPU_FUT_END_UNPARTITION_ON_TID**
- #define **_STARPU_FUT_START_FREE**
- #define **_STARPU_FUT_END_FREE**
- #define **_STARPU_FUT_START_WRITEBACK**
- #define **_STARPU_FUT_END_WRITEBACK**
- #define **_STARPU_FUT_SCHED_COMPONENT_PUSH_PRIO**
- #define **_STARPU_FUT_SCHED_COMPONENT_POP_PRIO**

- #define **_STARPU_FUT_START_WRITEBACK_ASYNC**
- #define **_STARPU_FUT_END_WRITEBACK_ASYNC**
- #define **_STARPU_FUT_HYPERVISOR_BEGIN**
- #define **_STARPU_FUT_HYPERVISOR_END**
- #define **_STARPU_FUT_BARRIER_WAIT_BEGIN**
- #define **_STARPU_FUT_BARRIER_WAIT_END**
- #define **_STARPU_FUT_WORKER_SCHEDULING_START**
- #define **_STARPU_FUT_WORKER_SCHEDULING_END**
- #define **_STARPU_FUT_WORKER_SCHEDULING_PUSH**
- #define **_STARPU_FUT_WORKER_SCHEDULING_POP**
- #define **_STARPU_FUT_START_EXECUTING**
- #define **_STARPU_FUT_END_EXECUTING**
- #define **_STARPU_FUT_SCHED_COMPONENT_NEW**
- #define **_STARPU_FUT_SCHED_COMPONENT_CONNECT**
- #define **_STARPU_FUT_SCHED_COMPONENT_PUSH**
- #define **_STARPU_FUT_SCHED_COMPONENT_PULL**
- #define **_STARPU_FUT_TASK_SUBMIT_START**
- #define **_STARPU_FUT_TASK_SUBMIT_END**
- #define **_STARPU_FUT_TASK_BUILD_START**
- #define **_STARPU_FUT_TASK_BUILD_END**
- #define **_STARPU_FUT_TASK_MPI_DECODE_START**
- #define **_STARPU_FUT_TASK_MPI_DECODE_END**
- #define **_STARPU_FUT_TASK_MPI_PRE_START**
- #define **_STARPU_FUT_TASK_MPI_PRE_END**
- #define **_STARPU_FUT_TASK_MPI_POST_START**
- #define **_STARPU_FUT_TASK_MPI_POST_END**
- #define **_STARPU_FUT_TASK_WAIT_START**
- #define **_STARPU_FUT_TASK_WAIT_END**
- #define **_STARPU_FUT_TASK_WAIT_FOR_ALL_START**
- #define **_STARPU_FUT_TASK_WAIT_FOR_ALL_END**
- #define **_STARPU_FUT_HANDLE_DATA_REGISTER**
- #define **_STARPU_FUT_START_FETCH_INPUT**
- #define **_STARPU_FUT_END_FETCH_INPUT**
- #define **_STARPU_FUT_TASK_THROTTLE_START**
- #define **_STARPU_FUT_TASK_THROTTLE_END**
- #define **_STARPU_FUT_DATA_STATE_INVALID**
- #define **_STARPU_FUT_DATA_STATE_OWNER**
- #define **_STARPU_FUT_DATA_STATE_SHARED**
- #define **_STARPU_FUT_DATA_REQUEST_CREATED**
- #define **_STARPU_FUT_TASK_EXCLUDE_FROM_DAG**
- #define **_STARPU_TRACE_NEW_MEM_NODE**(nodeid)
- #define **_STARPU_TRACE_WORKER_INIT_START**(a, b, c, d, e, f)
- #define **_STARPU_TRACE_WORKER_INIT_END**(workerid)
- #define **_STARPU_TRACE_START_CODELET_BODY**(job, nimpl, perf_arch, workerid)
- #define **_STARPU_TRACE_END_CODELET_BODY**(job, nimpl, perf_arch, workerid)
- #define **_STARPU_TRACE_START_EXECUTING**()
- #define **_STARPU_TRACE_END_EXECUTING**()
- #define **_STARPU_TRACE_START_CALLBACK**(job)
- #define **_STARPU_TRACE_END_CALLBACK**(job)
- #define **_STARPU_TRACE_JOB_PUSH**(task, prio)
- #define **_STARPU_TRACE_JOB_POP**(task, prio)
- #define **_STARPU_TRACE_UPDATE_TASK_CNT**(counter)
- #define **_STARPU_TRACE_START_FETCH_INPUT**(job)
- #define **_STARPU_TRACE_END_FETCH_INPUT**(job)
- #define **_STARPU_TRACE_START_PUSH_OUTPUT**(job)

- #define **_STARPU_TRACE_END_PUSH_OUTPUT**(job)
- #define **_STARPU_TRACE_TAG**(tag, job)
- #define **_STARPU_TRACE_TAG_DEPS**(a, b)
- #define **_STARPU_TRACE_TASK_DEPS**(a, b)
- #define **_STARPU_TRACE_GHOST_TASK_DEPS**(a, b)
- #define **_STARPU_TRACE_TASK_EXCLUDE_FROM_DAG**(a)
- #define **_STARPU_TRACE_TASK_NAME**(a)
- #define **_STARPU_TRACE_TASK_COLOR**(a)
- #define **_STARPU_TRACE_TASK_DONE**(a)
- #define **_STARPU_TRACE_TAG_DONE**(a)
- #define **_STARPU_TRACE_DATA_NAME**(a, b)
- #define **_STARPU_TRACE_DATA_COORDINATES**(a, b, c)
- #define **_STARPU_TRACE_DATA_COPY**(a, b, c)
- #define **_STARPU_TRACE_DATA_WONT_USE**(a)
- #define **_STARPU_TRACE_DATA_DOING_WONT_USE**(a)
- #define **_STARPU_TRACE_START_DRIVER_COPY**(a, b, c, d, e, f)
- #define **_STARPU_TRACE_END_DRIVER_COPY**(a, b, c, d, e)
- #define **_STARPU_TRACE_START_DRIVER_COPY_ASYNC**(a, b)
- #define **_STARPU_TRACE_END_DRIVER_COPY_ASYNC**(a, b)
- #define **_STARPU_TRACE_WORK_STEALING**(a, b)
- #define **_STARPU_TRACE_WORKER_DEINIT_START**
- #define **_STARPU_TRACE_WORKER_DEINIT_END**(a)
- #define **_STARPU_TRACE_WORKER_SCHEDULING_START**
- #define **_STARPU_TRACE_WORKER_SCHEDULING_END**
- #define **_STARPU_TRACE_WORKER_SCHEDULING_PUSH**
- #define **_STARPU_TRACE_WORKER_SCHEDULING_POP**
- #define **_STARPU_TRACE_WORKER_SLEEP_START**
- #define **_STARPU_TRACE_WORKER_SLEEP_END**
- #define **_STARPU_TRACE_TASK_SUBMIT**(job, a, b)
- #define **_STARPU_TRACE_TASK_SUBMIT_START**()
- #define **_STARPU_TRACE_TASK_SUBMIT_END**()
- #define **_STARPU_TRACE_TASK_THROTTLE_START**()
- #define **_STARPU_TRACE_TASK_THROTTLE_END**()
- #define **_STARPU_TRACE_TASK_BUILD_START**()
- #define **_STARPU_TRACE_TASK_BUILD_END**()
- #define **_STARPU_TRACE_TASK_MPI_DECODE_START**()
- #define **_STARPU_TRACE_TASK_MPI_DECODE_END**()
- #define **_STARPU_TRACE_TASK_MPI_PRE_START**()
- #define **_STARPU_TRACE_TASK_MPI_PRE_END**()
- #define **_STARPU_TRACE_TASK_MPI_POST_START**()
- #define **_STARPU_TRACE_TASK_MPI_POST_END**()
- #define **_STARPU_TRACE_TASK_WAIT_START**(job)
- #define **_STARPU_TRACE_TASK_WAIT_END**()
- #define **_STARPU_TRACE_TASK_WAIT_FOR_ALL_START**()
- #define **_STARPU_TRACE_TASK_WAIT_FOR_ALL_END**()
- #define **_STARPU_TRACE_USER_DEFINED_START**()
- #define **_STARPU_TRACE_USER_DEFINED_END**()
- #define **_STARPU_TRACE_START_ALLOC**(memnode, size, handle, is_prefetch)
- #define **_STARPU_TRACE_END_ALLOC**(memnode, handle, r)
- #define **_STARPU_TRACE_START_ALLOC_REUSE**(a, size, handle, is_prefetch)
- #define **_STARPU_TRACE_END_ALLOC_REUSE**(a, handle, r)
- #define **_STARPU_TRACE_START_FREE**(memnode, size, handle)
- #define **_STARPU_TRACE_END_FREE**(memnode, handle)
- #define **_STARPU_TRACE_START_WRITEBACK**(memnode, handle)
- #define **_STARPU_TRACE_END_WRITEBACK**(memnode, handle)

- #define **_STARPU_TRACE_USED_MEM**(memnode, used)
- #define **_STARPU_TRACE_START_MEMRECLAIM**(memnode, is_prefetch)
- #define **_STARPU_TRACE_END_MEMRECLAIM**(memnode, is_prefetch)
- #define **_STARPU_TRACE_START_WRITEBACK_ASYNC**(memnode)
- #define **_STARPU_TRACE_END_WRITEBACK_ASYNC**(memnode)
- #define **_STARPU_TRACE_START_PROGRESS**(memnode)
- #define **_STARPU_TRACE_END_PROGRESS**(memnode)
- #define **_STARPU_TRACE_USER_EVENT**(code)
- #define **_STARPU_TRACE_SET_PROFILING**(status)
- #define **_STARPU_TRACE_TASK_WAIT_FOR_ALL**()
- #define **_STARPU_TRACE_EVENT**(S)
- #define **_STARPU_TRACE_THREAD_EVENT**(S)
- #define **_STARPU_TRACE_LOCKING_MUTEX**()
- #define **_STARPU_TRACE_MUTEX_LOCKED**()
- #define **_STARPU_TRACE_UNLOCKING_MUTEX**()
- #define **_STARPU_TRACE_MUTEX_UNLOCKED**()
- #define **_STARPU_TRACE_TRYLOCK_MUTEX**()
- #define **_STARPU_TRACE_RDLOCKING_RWLOCK**()
- #define **_STARPU_TRACE_RWLOCK_RDLOCKED**()
- #define **_STARPU_TRACE_WRLOCKING_RWLOCK**()
- #define **_STARPU_TRACE_RWLOCK_WRLOCKED**()
- #define **_STARPU_TRACE_UNLOCKING_RWLOCK**()
- #define **_STARPU_TRACE_RWLOCK_UNLOCKED**()
- #define **_STARPU_TRACE_LOCKING_SPINLOCK**(file, line)
- #define **_STARPU_TRACE_SPINLOCK_LOCKED**(file, line)
- #define **_STARPU_TRACE_UNLOCKING_SPINLOCK**(file, line)
- #define **_STARPU_TRACE_SPINLOCK_UNLOCKED**(file, line)
- #define **_STARPU_TRACE_TRYLOCK_SPINLOCK**(file, line)
- #define **_STARPU_TRACE_COND_WAIT_BEGIN**()
- #define **_STARPU_TRACE_COND_WAIT_END**()
- #define **_STARPU_TRACE_BARRIER_WAIT_BEGIN**()
- #define **_STARPU_TRACE_BARRIER_WAIT_END**()
- #define **_STARPU_TRACE_MEMORY_FULL**(size)
- #define **_STARPU_TRACE_DATA_LOAD**(workerid, size)
- #define **_STARPU_TRACE_START_UNPARTITION**(handle, memnode)
- #define **_STARPU_TRACE_END_UNPARTITION**(handle, memnode)
- #define **_STARPU_TRACE_SCHED_COMPONENT_PUSH_PRIO**(workerid, ntasks, exp_len)
- #define **_STARPU_TRACE_SCHED_COMPONENT_POP_PRIO**(workerid, ntasks, exp_len)
- #define **_STARPU_TRACE_HYPERVISOR_BEGIN**()
- #define **_STARPU_TRACE_HYPERVISOR_END**()
- #define **_STARPU_TRACE_SCHED_COMPONENT_NEW**(component)
- #define **_STARPU_TRACE_SCHED_COMPONENT_CONNECT**(parent, child)
- #define **_STARPU_TRACE_SCHED_COMPONENT_PUSH**(from, to, task, prio)
- #define **_STARPU_TRACE_SCHED_COMPONENT_PULL**(from, to, task)
- #define **_STARPU_TRACE_HANDLE_DATA_REGISTER**(handle)
- #define **_STARPU_TRACE_HANDLE_DATA_UNREGISTER**(handle)
- #define **_STARPU_TRACE_WORKER_START_FETCH_INPUT**(job, id)
- #define **_STARPU_TRACE_WORKER_END_FETCH_INPUT**(job, id)
- #define **_STARPU_TRACE_DATA_STATE_INVALID**(handle, node)
- #define **_STARPU_TRACE_DATA_STATE_OWNER**(handle, node)
- #define **_STARPU_TRACE_DATA_STATE_SHARED**(handle, node)
- #define **_STARPU_TRACE_DATA_REQUEST_CREATED**(handle, orig, dest, prio, is_pre)

Functions

- static unsigned long **_starpu_fxt_get_job_id** (void)

Variables

- unsigned long **_starpu_job_cnt**

6.36 graph.h File Reference

```
#include <common/list.h>
```

Data Structures

- struct _starpu_graph_node

Functions

- void **_starpu_graph_init** (void)
- void **_starpu_graph_wrlock** (void)
- void **_starpu_graph_rdlock** (void)
- void **_starpu_graph_wrunlock** (void)
- void **_starpu_graph_rdunlock** (void)
- void _starpu_graph_add_job (struct _starpu_job ∗job)
- void _starpu_graph_add_job_dep (struct _starpu_job ∗job, struct _starpu_job ∗prev_job)
- void _starpu_graph_drop_job (struct _starpu_job ∗job)
- void _starpu_graph_drop_dropped_nodes (void)
- void _starpu_graph_compute_depths (void)
- void _starpu_graph_compute_descendants (void)
- void _starpu_graph_foreach (void(∗func)(void ∗data, struct _starpu_graph_node ∗node), void ∗data)

Variables

- int **_starpu_graph_record**

6.36.1 Data Structure Documentation

6.36.1.1 struct _starpu_graph_node

Data Fields

starpu_pthread_mutex_t	mutex	protects access to the job
struct _starpu_job ∗	job	pointer to the job, if it is still alive, NULL otherwise
struct _starpu_graph_node_multilist_top	top	Fields for graph analysis for scheduling heuristicsMember of list of all jobs without incoming dependency
struct _starpu_graph_node_multilist_bottom	bottom	Member of list of all jobs without outgoing dependency
struct _starpu_graph_node_multilist_all	all	Member of list of all jobs
struct _starpu_graph_node_multilist_dropped	dropped	Member of list of dropped jobs

Data Fields

[struct _starpu_graph_node](#) **	incoming	set of incoming dependencies May contain NULLs for terminated jobs
unsigned *	incoming_slot	Index within corresponding outgoing array
unsigned	n_incoming	Number of slots used
unsigned	alloc_incoming	Size of incoming
[struct _starpu_graph_node](#) **	outgoing	set of outgoing dependencies
unsigned *	outgoing_slot	Index within corresponding incoming array
unsigned	n_outgoing	Number of slots used
unsigned	alloc_outgoing	Size of outgoing
unsigned	depth	Rank from bottom, in number of jobs Only available if _starpu_graph_compute_depths was called
unsigned	descendants	Number of children, grand-children, etc. Only available if _starpu_graph_compute_descendants was called
int	graph_n	Variable available for graph flow

6.36.2 Function Documentation

6.36.2.1 _starpu_graph_add_job()

```
void _starpu_graph_add_job (
            struct _starpu_job * job )
```
Add a job to the graph, called before any _starpu_graph_add_job_dep call

6.36.2.2 _starpu_graph_add_job_dep()

```
void _starpu_graph_add_job_dep (
            struct _starpu_job * job,
            struct _starpu_job * prev_job )
```
Add a dependency between jobs

6.36.2.3 _starpu_graph_drop_job()

```
void _starpu_graph_drop_job (
            struct _starpu_job * job )
```
Remove a job from the graph

6.36.2.4 _starpu_graph_drop_dropped_nodes()

```
void _starpu_graph_drop_dropped_nodes (
            void  )
```
Really drop the nodes from the graph now

6.36.2.5 _starpu_graph_compute_depths()

```
void _starpu_graph_compute_depths (
            void  )
```
This make StarPU compute for each task the depth, i.e. the length of the longest path to a task without outgoing dependencies. This does not take job duration into account, just the number

6.36.2.6 _starpu_graph_compute_descendants()

```
void _starpu_graph_compute_descendants (
            void  )
```
Compute the descendants of jobs in the graph

6.36.2.7 _starpu_graph_foreach()

```
void _starpu_graph_foreach (
            void(*)(void *data, struct _starpu_graph_node *node) func,
            void * data )
```
This calls *func* for each node of the task graph, passing also *data* as it Apply func on each job of the graph

6.37 helper_mct.h File Reference

Data Structures

- struct _starpu_mct_data

Functions

- struct _starpu_mct_data ∗ **starpu_mct_init_parameters** (struct starpu_sched_component_mct_data ∗params)
- unsigned **starpu_mct_compute_execution_times** (struct starpu_sched_component ∗component, struct starpu_task ∗task, double ∗estimated_lengths, double ∗estimated_transfer_length, unsigned ∗suitable_↩ components)
- void **starpu_mct_compute_expected_times** (struct starpu_sched_component ∗component, struct starpu_task ∗task, double ∗estimated_lengths, double ∗estimated_transfer_length, double ∗estimated↩ _ends_with_task, double ∗min_exp_end_with_task, double ∗max_exp_end_with_task, unsigned ∗suitable↩ _components, unsigned nsuitable_components)
- double **starpu_mct_compute_fitness** (struct _starpu_mct_data ∗d, double exp_end, double min_exp_end, double max_exp_end, double transfer_len, double local_energy)
- int **starpu_mct_get_best_component** (struct _starpu_mct_data ∗d, struct starpu_task ∗task, double ∗estimated_lengths, double ∗estimated_transfer_length, double ∗estimated_ends_with_task, double min_↩ exp_end_with_task, double max_exp_end_with_task, unsigned ∗suitable_components, unsigned nsuitable↩ _components)

6.37.1 Data Structure Documentation

6.37.1.1 struct _starpu_mct_data

Data Fields

double	alpha	
double	beta	
double	_gamma	
double	idle_power	
starpu_pthread_mutex_t	scheduling_mutex	

6.38 idle_hook.h File Reference

Functions

- void **_starpu_init_idle_hooks** (void)
- unsigned **_starpu_execute_registered_idle_hooks** (void)

6.39   implicit_data_deps.h File Reference

```
#include <starpu.h>
#include <common/config.h>
```

Functions

- struct starpu_task ∗ **_starpu_detect_implicit_data_deps_with_handle** (struct starpu_task ∗pre_↩
  sync_task, struct starpu_task ∗post_sync_task, struct _starpu_task_wrapper_dlist ∗post_sync_task_↩
  dependency_slot, starpu_data_handle_t handle, enum starpu_data_access_mode mode, unsigned task_↩
  handle_sequential_consistency)
- int **_starpu_test_implicit_data_deps_with_handle** (starpu_data_handle_t handle, enum starpu_data_↩
  access_mode mode)
- void **_starpu_detect_implicit_data_deps** (struct starpu_task ∗task)
- void **_starpu_release_data_enforce_sequential_consistency** (struct starpu_task ∗task, struct _starpu↩
  _task_wrapper_dlist ∗task_dependency_slot, starpu_data_handle_t handle)
- void **_starpu_release_task_enforce_sequential_consistency** (struct _starpu_job ∗j)
- void **_starpu_add_post_sync_tasks** (struct starpu_task ∗post_sync_task, starpu_data_handle_t handle)
- void **_starpu_unlock_post_sync_tasks** (starpu_data_handle_t handle)
- void _starpu_implicit_data_deps_write_hook (void(∗func)(starpu_data_handle_t))
- int _starpu_data_wait_until_available (starpu_data_handle_t handle, enum starpu_data_access_mode
  mode, const char ∗sync_name)
- void **_starpu_data_clear_implicit** (starpu_data_handle_t handle)

6.39.1   Function Documentation

6.39.1.1   _starpu_implicit_data_deps_write_hook()

```
void _starpu_implicit_data_deps_write_hook (
            void(*)(starpu_data_handle_t) func )
```
Register a hook to be called when a write is submitted

6.39.1.2   _starpu_data_wait_until_available()

```
int _starpu_data_wait_until_available (
            starpu_data_handle_t handle,
            enum starpu_data_access_mode mode,
            const char * sync_name )
```
This function blocks until the handle is available in the requested mode

6.40   jobs.h File Reference

```
#include <starpu.h>
#include <semaphore.h>
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <string.h>
#include <stdarg.h>
#include <common/config.h>
#include <common/timing.h>
#include <common/list.h>
#include <common/fxt.h>
#include <core/dependencies/tags.h>
```

```
#include <datawizard/datawizard.h>
#include <core/perfmodel/perfmodel.h>
#include <core/errorcheck.h>
#include <common/barrier.h>
#include <common/utils.h>
#include <cuda.h>
```

Data Structures

- struct _starpu_data_descr
- struct _starpu_job

Macros

- #define **_STARPU_CPU_MAY_PERFORM**(j)
- #define **_STARPU_CUDA_MAY_PERFORM**(j)
- #define **_STARPU_OPENCL_MAY_PERFORM**(j)
- #define **_STARPU_MIC_MAY_PERFORM**(j)
- #define **_STARPU_JOB_GET_ORDERED_BUFFER_INDEX**(job, i)
- #define **_STARPU_JOB_GET_ORDERED_BUFFER_HANDLE**(job, i)
- #define **_STARPU_JOB_GET_ORDERED_BUFFER_MODE**(job, i)
- #define **_STARPU_JOB_GET_ORDERED_BUFFER_NODE**(job, i)
- #define **_STARPU_JOB_SET_ORDERED_BUFFER_INDEX**(job, __index, i)
- #define **_STARPU_JOB_SET_ORDERED_BUFFER_HANDLE**(job, __handle, i)
- #define **_STARPU_JOB_SET_ORDERED_BUFFER_MODE**(job, __mode, i)
- #define **_STARPU_JOB_SET_ORDERED_BUFFER_NODE**(job, __node, i)
- #define **_STARPU_JOB_SET_ORDERED_BUFFER**(job, buffer, i)
- #define **_STARPU_JOB_GET_ORDERED_BUFFERS**(job)
- #define **_STARPU_JOB_GET_DEP_SLOTS**(job)

Typedefs

- typedef void(∗ _starpu_cl_func_t) (void ∗∗, void ∗)

Functions

- void **_starpu_job_init** (void)
- void **_starpu_job_fini** (void)
- struct _starpu_job ∗ _starpu_job_create (struct starpu_task ∗task) STARPU_ATTRIBUTE_MALLOC
- void _starpu_job_destroy (struct _starpu_job ∗j)
- int _starpu_job_finished (struct _starpu_job ∗j)
- void _starpu_wait_job (struct _starpu_job ∗j)
- int _starpu_test_job_termination (struct _starpu_job ∗j)
- void _starpu_job_prepare_for_continuation_ext (struct _starpu_job ∗j, unsigned continuation_resubmit, void(∗continuation_callback_on_sleep)(void ∗arg), void ∗continuation_callback_on_sleep_arg)
- void **_starpu_job_prepare_for_continuation** (struct _starpu_job ∗j)
- void **_starpu_job_set_omp_cleanup_callback** (struct _starpu_job ∗j, void(∗omp_cleanup_callback)(void ∗arg), void ∗omp_cleanup_callback_arg)
- void _starpu_exclude_task_from_dag (struct starpu_task ∗task)
- unsigned _starpu_enforce_deps_and_schedule (struct _starpu_job ∗j)
- unsigned **_starpu_enforce_deps_starting_from_task** (struct _starpu_job ∗j)
- unsigned _starpu_reenforce_task_deps_and_schedule (struct _starpu_job ∗j)
- void **_starpu_enforce_deps_notify_job_ready_soon** (struct _starpu_job ∗j, _starpu_notify_job_start_data ∗data, int tag)
- void _starpu_handle_job_submission (struct _starpu_job ∗j)

- void _starpu_handle_job_termination (struct _starpu_job ∗j)
- size_t _starpu_job_get_data_size (struct starpu_perfmodel ∗model, struct starpu_perfmodel_arch ∗arch, unsigned nimpl, struct _starpu_job ∗j)
- struct starpu_task ∗ _starpu_pop_local_task (struct _starpu_worker ∗worker)
- int _starpu_push_local_task (struct _starpu_worker ∗worker, struct starpu_task ∗task, int prio)

6.40.1 Data Structure Documentation

6.40.1.1 struct _starpu_data_descr

Data Fields

starpu_data_handle_t	handle	
enum starpu_data_access_mode	mode	
int	node	
int	index	This is the value actually chosen, only set by _starpu_fetch_task_input for coherency with __starpu_push_task_output
int	orderedindex	

6.40.2 Typedef Documentation

6.40.2.1 _starpu_cl_func_t

```
typedef void(* _starpu_cl_func_t) (void **, void *)
```
codelet function

6.40.3 Function Documentation

6.40.3.1 _starpu_job_create()

```
struct _starpu_job* _starpu_job_create (
            struct starpu_task * task )
```
Create an internal struct _starpu_job ∗structure to encapsulate the task.

6.40.3.2 _starpu_job_destroy()

```
void _starpu_job_destroy (
            struct _starpu_job * j )
```
Destroy the data structure associated to the job structure

6.40.3.3 _starpu_job_finished()

```
int _starpu_job_finished (
            struct _starpu_job * j )
```
Test for the termination of the job

6.40.3.4 _starpu_wait_job()

```
void _starpu_wait_job (
            struct _starpu_job * j )
```
Wait for the termination of the job

6.40.3.5 _starpu_test_job_termination()

```
int _starpu_test_job_termination (
            struct _starpu_job * j )
```
Test for the termination of the job

6.40.3.6 _starpu_job_prepare_for_continuation_ext()

```
void _starpu_job_prepare_for_continuation_ext (
            struct _starpu_job * j,
            unsigned continuation_resubmit,
            void(*)(void *arg) continuation_callback_on_sleep,
            void * continuation_callback_on_sleep_arg )
```
Prepare the job for accepting new dependencies before becoming a continuation.

6.40.3.7 _starpu_exclude_task_from_dag()

```
void _starpu_exclude_task_from_dag (
            struct starpu_task * task )
```
Specify that the task should not appear in the DAG generated by debug tools.

6.40.3.8 _starpu_enforce_deps_and_schedule()

```
unsigned _starpu_enforce_deps_and_schedule (
            struct _starpu_job * j )
```
try to submit job j, enqueue it if it's not schedulable yet. The job's sync mutex is supposed to be held already

6.40.3.9 _starpu_reenforce_task_deps_and_schedule()

```
unsigned _starpu_reenforce_task_deps_and_schedule (
            struct _starpu_job * j )
```
When waking up a continuation, we only enforce new task dependencies

6.40.3.10 _starpu_handle_job_submission()

```
void _starpu_handle_job_submission (
            struct _starpu_job * j )
```
Called at the submission of the job

6.40.3.11 _starpu_handle_job_termination()

```
void _starpu_handle_job_termination (
            struct _starpu_job * j )
```
This function must be called after the execution of a job, this triggers all job's dependencies and perform the callback function if any.

6.40.3.12 _starpu_job_get_data_size()

```
size_t _starpu_job_get_data_size (
            struct starpu_perfmodel * model,
            struct starpu_perfmodel_arch * arch,
            unsigned nimpl,
            struct _starpu_job * j )
```
Get the sum of the size of the data accessed by the job.

6.40.3.13 _starpu_pop_local_task()

```
struct starpu_task* _starpu_pop_local_task (
            struct _starpu_worker * worker )
```
Get a task from the local pool of tasks that were explicitly attributed to that worker.

6.40.3.14 _starpu_push_local_task()

```
int _starpu_push_local_task (
            struct _starpu_worker * worker,
            struct starpu_task * task,
            int prio )
```

Put a task into the pool of tasks that are explicitly attributed to the specified worker. If "back" is set, the task is put at the back of the list. Considering the tasks are popped from the back, this value should be 0 to enforce a FIFO ordering.

6.41 malloc.h File Reference

Functions

- void **_starpu_malloc_init** (unsigned dst_node)
- void **_starpu_malloc_shutdown** (unsigned dst_node)
- void **_starpu_free_on_node** (unsigned dst_node, uintptr_t addr, size_t size)
- int **_starpu_malloc_flags_on_node** (unsigned dst_node, void ∗∗A, size_t dim, int flags)
- int **_starpu_free_flags_on_node** (unsigned dst_node, void ∗A, size_t dim, int flags)
- int _starpu_malloc_willpin_on_node (unsigned dst_node)

6.41.1 Function Documentation

6.41.1.1 _starpu_malloc_willpin_on_node()

```
int _starpu_malloc_willpin_on_node (
            unsigned dst_node )
```

Returns whether when allocating data on `dst_node`, we will do pinning, i.e. the allocation will be very expensive, and should thus be moved out from the critical path

6.42 memalloc.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/list.h>
#include <datawizard/interfaces/data_interface.h>
#include <datawizard/coherency.h>
#include <datawizard/copy_driver.h>
#include <datawizard/data_request.h>
```

6.43 memory_manager.h File Reference

```
#include <starpu.h>
```

Functions

- int _starpu_memory_manager_init ()
- void _starpu_memory_manager_set_global_memory_size (unsigned node, size_t size)
- size_t _starpu_memory_manager_get_global_memory_size (unsigned node)
- int **_starpu_memory_manager_test_allocate_size** (unsigned node, size_t size)

6.43.1 Function Documentation

6.43.1.1 _starpu_memory_manager_init()

```
int _starpu_memory_manager_init ( )
```
Initialises the memory manager

6.43.1.2 _starpu_memory_manager_set_global_memory_size()

```
void _starpu_memory_manager_set_global_memory_size (
            unsigned node,
            size_t size )
```
Initialises the global memory size for the given node

6.43.1.3 _starpu_memory_manager_get_global_memory_size()

```
size_t _starpu_memory_manager_get_global_memory_size (
            unsigned node )
```
Gets the global memory size for the given node

6.44 memory_nodes.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <datawizard/coherency.h>
#include <datawizard/memalloc.h>
#include <datawizard/node_ops.h>
#include <common/utils.h>
#include <core/workers.h>
#include <core/simgrid.h>
```

Data Structures

- struct _starpu_cond_and_worker
- struct _starpu_memory_node_descr

Macros

- #define **starpu_node_get_kind**
- #define **starpu_memory_nodes_get_count**
- #define **starpu_worker_get_memory_node**
- #define **starpu_worker_get_local_memory_node**

Functions

- void **_starpu_memory_nodes_init** (void)
- void **_starpu_memory_nodes_deinit** (void)
- static void **_starpu_memory_node_add_nworkers** (unsigned node)
- void _starpu_worker_drives_memory_node (struct _starpu_worker ∗worker, unsigned memnode)
- static struct _starpu_node_ops ∗ **_starpu_memory_node_get_node_ops** (unsigned node)
- static unsigned **_starpu_memory_node_get_nworkers** (unsigned node)
- static void **_starpu_simgrid_memory_node_set_host** (unsigned node, starpu_sg_host_t host)
- static starpu_sg_host_t **_starpu_simgrid_memory_node_get_host** (unsigned node)
- unsigned **_starpu_memory_node_register** (enum starpu_node_kind kind, int devid, struct _starpu_node↩
  _ops ∗node_ops)

- void **_starpu_memory_node_register_condition** (struct _starpu_worker ∗worker, starpu_pthread_cond↩
_t ∗cond, unsigned nodeid)
- static struct _starpu_memory_node_descr ∗ **_starpu_memory_node_get_description** (void)
- static enum starpu_node_kind **_starpu_node_get_kind** (unsigned node)
- static unsigned **_starpu_memory_nodes_get_count** (void)
- static unsigned _starpu_worker_get_memory_node (unsigned workerid)
- static unsigned **_starpu_worker_get_local_memory_node** (void)

Variables

- char **_starpu_worker_drives_memory** [STARPU_NMAXWORKERS][STARPU_MAXNODES]
- struct _starpu_memory_node_descr **_starpu_descr**

6.44.1 Data Structure Documentation

6.44.1.1 struct _starpu_cond_and_worker

Data Fields

starpu_pthread_cond_t ∗	cond	
struct _starpu_worker ∗	worker	

6.44.1.2 struct _starpu_memory_node_descr

Data Fields

unsigned	nnodes	
enum starpu_node_kind	nodes[STARPU_MAXNODES]	
struct _starpu_node_ops ∗	node_ops[STARPU_MAXNODES]	
int	devid[STARPU_MAXNODES]	Get the device id associated to this node, or -1 if not applicable
unsigned	nworkers[STARPU_MAXNODES]	
starpu_sg_host_t	host[STARPU_MAXNODES]	
starpu_pthread_rwlock_t	conditions_rwlock	Every worker is associated to a condition variable on which the worker waits when there is task available. It is possible that multiple worker share the same condition variable, so we maintain a list of all these condition variables so that we can wake up all worker attached to a memory node that are waiting on a task.
struct _starpu_cond_and_worker	conditions_attached_to_node[STARPU_MAXNODES][STARPU_NMAXWORKERS]	
struct _starpu_cond_and_worker	conditions_all[STARPU_MAXNODES ∗STARPU_NMAXWORKERS]	
unsigned	total_condition_count	the number of queues attached to each node
unsigned	condition_count[STARPU_MAXNODES]	

6.44.2 Function Documentation

6.44.2.1 _starpu_worker_drives_memory_node()

```
void _starpu_worker_drives_memory_node (
            struct _starpu_worker * worker,
            unsigned memnode )
```
same utility as _starpu_memory_node_add_nworkers

6.44.2.2 _starpu_worker_get_memory_node()

```
static unsigned _starpu_worker_get_memory_node (
            unsigned workerid ) [inline], [static]
```
This workerid may either be a basic worker or a combined worker
We have a combined worker

6.45 memstats.h File Reference

```
#include <starpu.h>
#include <common/config.h>
```

Typedefs

- typedef void ∗ **_starpu_memory_stats_t**

Functions

- void **_starpu_memory_stats_init** (starpu_data_handle_t handle)
- void **_starpu_memory_stats_init_per_node** (starpu_data_handle_t handle, unsigned node)
- void **_starpu_memory_stats_free** (starpu_data_handle_t handle)
- void **_starpu_memory_display_handle_stats** (FILE ∗stream, starpu_data_handle_t handle)
- void **_starpu_memory_handle_stats_cache_hit** (starpu_data_handle_t handle, unsigned node)
- void **_starpu_memory_handle_stats_loaded_shared** (starpu_data_handle_t handle, unsigned node)
- void **_starpu_memory_handle_stats_loaded_owner** (starpu_data_handle_t handle, unsigned node)
- void **_starpu_memory_handle_stats_shared_to_owner** (starpu_data_handle_t handle, unsigned node)
- void **_starpu_memory_handle_stats_invalidated** (starpu_data_handle_t handle, unsigned node)

6.46 mp_common.h File Reference

```
#include <semaphore.h>
#include <starpu.h>
#include <common/config.h>
#include <common/list.h>
#include <common/barrier.h>
#include <common/thread.h>
#include <datawizard/interfaces/data_interface.h>
#include <datawizard/copy_driver.h>
```

6.47 multiple_regression.h File Reference

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include <core/perfmodel/perfmodel.h>
#include <starpu.h>
```

Functions

- int **_starpu_multiple_regression** ([struct](#) starpu_perfmodel_history_list ∗ptr, double ∗coeff, unsigned ncoeff, unsigned nparameters, const char ∗∗parameters_names, unsigned ∗∗combinations, const char ∗codelet_↩ name)

6.48 node_ops.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <datawizard/copy_driver.h>
```

Data Structures

- struct [_starpu_node_ops](#)

Typedefs

- typedef int(∗ **copy_interface_func_t**) (starpu_data_handle_t handle, void ∗src_interface, unsigned src_↩ node, void ∗dst_interface, unsigned dst_node, [struct](#) _starpu_data_request ∗req)
- typedef int(∗ **copy_data_t**) (uintptr_t src_ptr, size_t src_offset, unsigned src_node, uintptr_t dst_ptr, size_t dst_offset, unsigned dst_node, size_t ssize, [struct _starpu_async_channel](#) ∗async_channel)
- typedef int(∗ **copy2d_data_t**) (uintptr_t src_ptr, size_t src_offset, unsigned src_node, uintptr_t dst_ptr, size↩ _t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks, size_t ld_src, size_t ld_dst, [struct _↩ starpu_async_channel](#) ∗async_channel)
- typedef int(∗ **copy3d_data_t**) (uintptr_t src_ptr, size_t src_offset, unsigned src_node, uintptr_t dst_ptr, size↩ _t dst_offset, unsigned dst_node, size_t blocksize, size_t numblocks_1, size_t ld1_src, size_t ld1_dst, size_t numblocks_2, size_t ld2_src, size_t ld2_dst, [struct _starpu_async_channel](#) ∗async_channel)

Functions

- const char ∗ **_starpu_node_get_prefix** (enum starpu_node_kind kind)

6.49 openmp_runtime_support.h File Reference

```
#include <starpu.h>
#include <common/list.h>
#include <common/starpu_spinlock.h>
#include <common/uthash.h>
#include <ucontext.h>
```

Data Structures

- struct [starpu_omp_numeric_place](#)
- struct [starpu_omp_place](#)
- struct [starpu_omp_data_environment_icvs](#)
- struct [starpu_omp_device_icvs](#)
- struct [starpu_omp_implicit_task_icvs](#)
- struct [starpu_omp_global_icvs](#)
- struct [starpu_omp_initial_icv_values](#)
- struct [starpu_omp_task_group](#)
- struct [starpu_omp_task_link](#)
- struct [starpu_omp_condition](#)
- struct [starpu_omp_critical](#)

Macros

- #define _XOPEN_SOURCE
- #define STARPU_OMP_MAX_ACTIVE_LEVELS

Enumerations

- enum starpu_omp_place_name {
  starpu_omp_place_undefined, **starpu_omp_place_threads**, **starpu_omp_place_cores**, **starpu_omp↩
  _place_sockets**,
  starpu_omp_place_numerical }
- enum starpu_omp_task_state {
  starpu_omp_task_state_clear, **starpu_omp_task_state_preempted**, **starpu_omp_task_state_↩
  terminated**, **starpu_omp_task_state_zombie**,
  starpu_omp_task_state_target }
- enum **starpu_omp_task_wait_on** {
  starpu_omp_task_wait_on_task_childs, **starpu_omp_task_wait_on_region_tasks**, **starpu_omp_↩
  task_wait_on_barrier**, **starpu_omp_task_wait_on_group**,
  starpu_omp_task_wait_on_critical, **starpu_omp_task_wait_on_ordered**, **starpu_omp_task_wait_on↩
  _lock**, **starpu_omp_task_wait_on_nest_lock** }
- enum **starpu_omp_task_flags** { **STARPU_OMP_TASK_FLAGS_IMPLICIT**, **STARPU_OMP_TASK_FL↩
  AGS_UNDEFERRED**, **STARPU_OMP_TASK_FLAGS_FINAL**, **STARPU_OMP_TASK_FLAGS_UNTIED** }

Variables

- starpu_pthread_key_t **omp_thread_key**
- starpu_pthread_key_t **omp_task_key**

6.49.1 Data Structure Documentation

6.49.1.1 struct starpu_omp_numeric_place

Data Fields

int	excluded_place	
int ∗	included_numeric_items	
int	nb_included_numeric_items	
int ∗	excluded_numeric_items	
int	nb_excluded_numeric_items	

6.49.1.2 struct starpu_omp_place

OpenMP place for thread afinity, defined by the OpenMP spec

Data Fields

int	abstract_name	
int	abstract_excluded	
int	abstract_length	
struct starpu_omp_numeric_place ∗	numeric_places	
int	nb_numeric_places	

6.49.1.3 struct starpu_omp_data_environment_icvs

Internal Control Variables (ICVs) declared following OpenMP 4.0.0 spec section 2.3.1

Data Fields

int	dyn_var	parallel region icvs
int	nest_var	
int ∗	nthreads_var	
int	thread_limit_var	nthreads_var ICV is a list
int	active_levels_var	
int	levels_var	
int ∗	bind_var	
int	run_sched_var	bind_var ICV is a list loop region icvs
unsigned long long	run_sched_chunk_var	
int	default_device_var	program execution icvs
int	max_task_priority_var	

6.49.1.4 struct starpu_omp_device_icvs

Data Fields

int	max_active_levels_var	parallel region icvs
int	def_sched_var	loop region icvs
unsigned long long	def_sched_chunk_var	
int	stacksize_var	program execution icvs
int	wait_policy_var	

6.49.1.5 struct starpu_omp_implicit_task_icvs

Data Fields

int	place_partition_var	parallel region icvs

6.49.1.6 struct starpu_omp_global_icvs

Data Fields

int	cancel_var	program execution icvs

6.49.1.7 struct starpu_omp_initial_icv_values

Data Fields

int	dyn_var	
int	nest_var	
int ∗	nthreads_var	
int	run_sched_var	
unsigned long long	run_sched_chunk_var	
int	def_sched_var	
unsigned long long	def_sched_chunk_var	
int ∗	bind_var	
int	stacksize_var	
int	wait_policy_var	

Data Fields

int	thread_limit_var	
int	max_active_levels_var	
int	active_levels_var	
int	levels_var	
int	place_partition_var	
int	cancel_var	
int	default_device_var	
int	max_task_priority_var	
struct starpu_omp_place	places	not a real ICV, but needed to store the contents of OMP_PLACES

6.49.1.8   struct starpu_omp_task_group

Data Fields

int	descendent_task_count	
struct starpu_omp_task ∗	leader_task	
struct starpu_omp_task_group ∗	p_previous_task_group	

6.49.1.9   struct starpu_omp_task_link

Data Fields

struct starpu_omp_task ∗	task	
struct starpu_omp_task_link ∗	next	

6.49.1.10   struct starpu_omp_condition

Data Fields

struct starpu_omp_task_link ∗	contention_list_head	

6.49.1.11   struct starpu_omp_critical

Data Fields

UT_hash_handle	hh	
struct _starpu_spinlock	lock	
unsigned	state	
struct starpu_omp_task_link ∗	contention_list_head	
const char ∗	name	

6.49.2   Macro Definition Documentation

6.49.2.1   _XOPEN_SOURCE

`#define _XOPEN_SOURCE`
ucontexts have been deprecated as of POSIX 1-2004 _XOPEN_SOURCE required at least on OS/X

TODO: add detection in configure.ac

6.49.2.2 STARPU_OMP_MAX_ACTIVE_LEVELS

```
#define STARPU_OMP_MAX_ACTIVE_LEVELS
```
Arbitrary limit on the number of nested parallel sections

6.49.3 Enumeration Type Documentation

6.49.3.1 starpu_omp_place_name

```
enum starpu_omp_place_name
```
Possible abstract names for OpenMP places

6.49.3.2 starpu_omp_task_state

```
enum starpu_omp_task_state
```

Enumerator

starpu_omp_task_state_target	target tasks are non-preemptible tasks, without dedicated stack and OpenMP Runtime Support context

6.50 perfmodel.h File Reference

```
#include <common/config.h>
#include <starpu.h>
#include <core/task_bundle.h>
#include <stdio.h>
```

Data Structures

- struct _starpu_perfmodel_state

Macros

- #define _STARPU_PERFMODEL_VERSION

Functions

- char ∗ **_starpu_get_perf_model_dir_codelet** ()
- char ∗ **_starpu_get_perf_model_dir_bus** ()
- char ∗ **_starpu_get_perf_model_dir_debug** ()
- double **_starpu_history_based_job_expected_perf** (struct starpu_perfmodel ∗model, struct starpu_←perfmodel_arch ∗arch, struct _starpu_job ∗j, unsigned nimpl)
- void **_starpu_load_history_based_model** (struct starpu_perfmodel ∗model, unsigned scan_history)
- void **_starpu_init_and_load_perfmodel** (struct starpu_perfmodel ∗model)
- void **_starpu_initialize_registered_performance_models** (void)
- void **_starpu_deinitialize_registered_performance_models** (void)
- void **_starpu_deinitialize_performance_model** (struct starpu_perfmodel ∗model)
- double **_starpu_regression_based_job_expected_perf** (struct starpu_perfmodel ∗model, struct starpu←_perfmodel_arch ∗arch, struct _starpu_job ∗j, unsigned nimpl)

- double **_starpu_non_linear_regression_based_job_expected_perf** (struct starpu_perfmodel *model, struct starpu_perfmodel_arch *arch, struct _starpu_job *j, unsigned nimpl)
- double **_starpu_multiple_regression_based_job_expected_perf** (struct starpu_perfmodel *model, struct starpu_perfmodel_arch *arch, struct _starpu_job *j, unsigned nimpl)
- void **_starpu_update_perfmodel_history** (struct _starpu_job *j, struct starpu_perfmodel *model, struct starpu_perfmodel_arch *arch, unsigned cpuid, double measured, unsigned nimpl)
- int **_starpu_perfmodel_create_comb_if_needed** (struct starpu_perfmodel_arch *arch)
- void **_starpu_create_sampling_directory_if_needed** (void)
- void **_starpu_load_bus_performance_files** (void)
- void **_starpu_set_calibrate_flag** (unsigned val)
- unsigned **_starpu_get_calibrate_flag** (void)
- unsigned * **_starpu_get_cuda_affinity_vector** (unsigned gpuid)
- unsigned * **_starpu_get_opencl_affinity_vector** (unsigned gpuid)
- void **_starpu_save_bandwidth_and_latency_disk** (double bandwidth_write, double bandwidth_read, double latency_write, double latency_read, unsigned node, const char *name)
- void **_starpu_write_double** (FILE *f, const char *format, double val)
- int **_starpu_read_double** (FILE *f, char *format, double *val)
- void **_starpu_simgrid_get_platform_path** (int version, char *path, size_t maxlen)
- void **_starpu_perfmodel_realloc** (struct starpu_perfmodel *model, int nb)
- void **_starpu_free_arch_combs** (void)
- hwloc_topology_t **_starpu_perfmodel_get_hwtopology** ()

Variables

- unsigned **_starpu_calibration_minimum**

6.50.1 Data Structure Documentation

6.50.1.1 struct _starpu_perfmodel_state

Data Fields

struct starpu_perfmodel_per_arch **	per_arch	
int **	per_arch_is_set	
starpu_pthread_rwlock_t	model_rwlock	
int *	nimpls	
int *	nimpls_set	
int	ncombs	The number of combinations currently used by the model
int	ncombs_set	The number of combinations allocated in the array nimpls and ncombs
int *	combs	

6.50.2 Macro Definition Documentation

6.50.2.1 _STARPU_PERFMODEL_VERSION

```
#define _STARPU_PERFMODEL_VERSION
```
Performance models files are stored in a directory whose name include the version of the performance model format. The version number is also written in the file itself. When updating the format, the variable _STARPU_PE←
RFMODEL_VERSION should be updated. It is then possible to switch easily between differents versions of StarPU having different performance model formats.

6.51 prio_deque.h File Reference

```
#include <starpu.h>
#include <starpu_scheduler.h>
#include <core/task.h>
```

Data Structures

- struct _starpu_prio_deque

Functions

- static void **_starpu_prio_deque_init** (struct _starpu_prio_deque ∗pdeque)
- static void **_starpu_prio_deque_destroy** (struct _starpu_prio_deque ∗pdeque)
- static int _starpu_prio_deque_is_empty (struct _starpu_prio_deque ∗pdeque)
- static void **_starpu_prio_deque_erase** (struct _starpu_prio_deque ∗pdeque, struct starpu_task ∗task)
- static int _starpu_prio_deque_push_front_task (struct _starpu_prio_deque ∗pdeque, struct starpu_task ∗task)
- static int **_starpu_prio_deque_push_back_task** (struct _starpu_prio_deque ∗pdeque, struct starpu_task ∗task)
- static struct starpu_task ∗ **_starpu_prio_deque_highest_task** (struct _starpu_prio_deque ∗pdeque)
- static struct starpu_task ∗ _starpu_prio_deque_pop_task (struct _starpu_prio_deque ∗pdeque)
- static struct starpu_task ∗ **_starpu_prio_deque_pop_back_task** (struct _starpu_prio_deque ∗pdeque)
- static int **_starpu_prio_deque_pop_this_task** (struct _starpu_prio_deque ∗pdeque, int workerid, struct starpu_task ∗task)
- struct starpu_task ∗ _starpu_prio_deque_pop_task_for_worker (struct _starpu_prio_deque ∗, int workerid, int ∗skipped)
- struct starpu_task ∗ _starpu_prio_deque_deque_task_for_worker (struct _starpu_prio_deque ∗, int workerid, int ∗skipped)
- struct starpu_task ∗ **_starpu_prio_deque_deque_first_ready_task** (struct _starpu_prio_deque ∗, unsigned workerid)

6.51.1 Data Structure Documentation

6.51.1.1 struct _starpu_prio_deque

Data Fields

struct starpu_task_prio_list	list	
unsigned	ntasks	
unsigned	nprocessed	
double	exp_start	
double	exp_end	
double	exp_len	

6.51.2 Function Documentation

6.51.2.1 _starpu_prio_deque_is_empty()

```
static int _starpu_prio_deque_is_empty (
            struct _starpu_prio_deque * pdeque ) [inline], [static]
```

return 0 iff the struct _starpu_prio_deque is not empty

6.51.2.2 _starpu_prio_deque_push_front_task()

```
static int _starpu_prio_deque_push_front_task (
            struct _starpu_prio_deque * pdeque,
            struct starpu_task * task ) [inline], [static]
```
push a task in O(lg(nb priorities))

6.51.2.3 _starpu_prio_deque_pop_task()

```
static struct starpu_task* _starpu_prio_deque_pop_task (
            struct _starpu_prio_deque * pdeque ) [static]
```
all _starpu_prio_deque_pop/deque_task function return a task or a NULL pointer if none are available in O(lg(nb priorities))

6.51.2.4 _starpu_prio_deque_pop_task_for_worker()

```
struct starpu_task* _starpu_prio_deque_pop_task_for_worker (
            struct _starpu_prio_deque * ,
            int workerid,
            int * skipped )
```
return a task that can be executed by workerid

6.51.2.5 _starpu_prio_deque_deque_task_for_worker()

```
struct starpu_task* _starpu_prio_deque_deque_task_for_worker (
            struct _starpu_prio_deque * ,
            int workerid,
            int * skipped )
```
return a task that can be executed by workerid

6.52 prio_list.h File Reference

```
#include <common/rbtree.h>
```

Macros

- #define **PRIO_LIST_INLINE**
- #define **PRIO_struct**
- #define **PRIO_LIST_CREATE_TYPE**(ENAME, PRIOFIELD)

6.53 profiling.h File Reference

```
#include <starpu.h>
#include <starpu_profiling.h>
#include <starpu_util.h>
#include <common/config.h>
```

Functions

- struct starpu_profiling_task_info ∗ _starpu_allocate_profiling_info_if_needed (struct starpu_task ∗task)
- void _starpu_worker_update_profiling_info_executing (int workerid, struct timespec ∗executing_time, int executed_tasks, uint64_t used_cycles, uint64_t stall_cycles, double consumed_energy, double flops)
- void _starpu_worker_restart_sleeping (int workerid)
- void _starpu_worker_stop_sleeping (int workerid)
- void _starpu_worker_register_executing_start_date (int workerid, struct timespec ∗executing_start)

- void [_starpu_worker_register_executing_end](link) (int workerid)
- void [_starpu_initialize_busid_matrix](link) (void)
- int [_starpu_register_bus](link) (int src_node, int dst_node)
- void [_starpu_bus_update_profiling_info](link) (int src_node, int dst_node, size_t size)
- void **_starpu_profiling_set_task_push_start_time** ([struct](link) starpu_task *task)
- void **_starpu_profiling_set_task_push_end_time** ([struct](link) starpu_task *task)
- void [_starpu_profiling_init](link) (void)
- void [_starpu_profiling_start](link) (void)
- void **_starpu_profiling_terminate** (void)

6.53.1 Function Documentation

6.53.1.1 _starpu_allocate_profiling_info_if_needed()

```
struct starpu_profiling_task_info* _starpu_allocate_profiling_info_if_needed (
            struct starpu_task * task )
```
Create a task profiling info structure (with the proper time stamps) in case profiling is enabled.

6.53.1.2 _starpu_worker_update_profiling_info_executing()

```
void _starpu_worker_update_profiling_info_executing (
            int workerid,
            struct timespec * executing_time,
            int executed_tasks,
            uint64_t used_cycles,
            uint64_t stall_cycles,
            double consumed_energy,
            double flops )
```
Update the per-worker profiling info after a task (or more) was executed. This tells StarPU how much time was spent doing computation.

6.53.1.3 _starpu_worker_restart_sleeping()

```
void _starpu_worker_restart_sleeping (
            int workerid )
```
Record the date when the worker started to sleep. This permits to measure how much time was spent sleeping.

6.53.1.4 _starpu_worker_stop_sleeping()

```
void _starpu_worker_stop_sleeping (
            int workerid )
```
Record the date when the worker stopped sleeping. This permits to measure how much time was spent sleeping.

6.53.1.5 _starpu_worker_register_executing_start_date()

```
void _starpu_worker_register_executing_start_date (
            int workerid,
            struct timespec * executing_start )
```
Record the date when the worker started to execute a piece of code. This permits to measure how much time was really spent doing computation at the end of the codelet.

6.53.1.6 _starpu_worker_register_executing_end()

```
void _starpu_worker_register_executing_end (
            int workerid )
```
Record that the worker is not executing any more.

6.53.1.7 _starpu_initialize_busid_matrix()

```
void _starpu_initialize_busid_matrix (
              void  )
```
When StarPU is initialized, a matrix describing all the bus between memory nodes is created: it indicates whether there is a physical link between two memory nodes or not. This matrix should contain the identifier of the bus between two nodes or -1 in case there is no link.

6.53.1.8 _starpu_register_bus()

```
int _starpu_register_bus (
              int src_node,
              int dst_node )
```
Tell StarPU that there exists a link between the two memory nodes. This function returns the identifier associated to the bus which can be used to retrieve profiling information about the bus activity later on.

6.53.1.9 _starpu_bus_update_profiling_info()

```
void _starpu_bus_update_profiling_info (
              int src_node,
              int dst_node,
              size_t size )
```
Tell StarPU that "size" bytes were transferred between the two specified memory nodes.

6.53.1.10 _starpu_profiling_init()

```
void _starpu_profiling_init (
              void  )
```
This function needs to be called before other starpu_profile_* functions

6.53.1.11 _starpu_profiling_start()

```
void _starpu_profiling_start (
              void  )
```
This function starts profiling if the STARPU_PROFILING environment variable was set

6.54 progress_hook.h File Reference

Functions

- void **_starpu_init_progression_hooks** (void)
- unsigned **_starpu_execute_registered_progression_hooks** (void)

6.55 rbtree.h File Reference

```
#include <stddef.h>
#include <assert.h>
#include <stdint.h>
#include <sys/types.h>
#include "rbtree_i.h"
```

Macros

- #define **MACRO_BEGIN**
- #define **MACRO_END**
- #define **STARPU_RBTREE_LEFT**
- #define **STARPU_RBTREE_RIGHT**

- #define STARPU_RBTREE_INITIALIZER
- #define starpu_rbtree_entry(node, type, member)
- #define starpu_rbtree_lookup(tree, key, cmp_fn)
- #define starpu_rbtree_lookup_nearest(tree, key, cmp_fn, dir)
- #define starpu_rbtree_insert(tree, node, cmp_fn)
- #define starpu_rbtree_lookup_slot(tree, key, cmp_fn, slot)
- #define starpu_rbtree_first(tree)
- #define starpu_rbtree_last(tree)
- #define starpu_rbtree_prev(node)
- #define starpu_rbtree_next(node)
- #define starpu_rbtree_for_each_remove(tree, node, tmp)

Functions

- static void starpu_rbtree_init (struct starpu_rbtree ∗tree)
- static void starpu_rbtree_node_init (struct starpu_rbtree_node ∗node)
- static int **starpu_rbtree_node_unlinked** (const struct starpu_rbtree_node ∗node)
- static int starpu_rbtree_empty (const struct starpu_rbtree ∗tree)
- static void starpu_rbtree_insert_slot (struct starpu_rbtree ∗tree, uintptr_t slot, struct starpu_rbtree_node ∗node)
- void starpu_rbtree_remove (struct starpu_rbtree ∗tree, struct starpu_rbtree_node ∗node)

6.55.1 Macro Definition Documentation

6.55.1.1 STARPU_RBTREE_INITIALIZER

```
#define STARPU_RBTREE_INITIALIZER
```
Static tree initializer.

6.55.1.2 starpu_rbtree_entry

```
#define starpu_rbtree_entry(
            node,
            type,
            member )
```
Macro that evaluates to the address of the structure containing the given node based on the given type and member.

6.55.1.3 starpu_rbtree_lookup

```
#define starpu_rbtree_lookup(
            tree,
            key,
            cmp_fn )
```
Look up a node in a tree.
Note that implementing the lookup algorithm as a macro gives two benefits: First, it avoids the overhead of a callback function. Next, the type of the cmp_fn parameter isn't rigid. The only guarantee offered by this implementation is that the key parameter is the first parameter given to cmp_fn. This way, users can pass only the value they need for comparison instead of e.g. allocating a full structure on the stack.
See starpu_rbtree_insert().

6.55.1.4 starpu_rbtree_lookup_nearest

```
#define starpu_rbtree_lookup_nearest(
            tree,
            key,
            cmp_fn,
            dir )
```
Look up a node or one of its nearest nodes in a tree.

This macro essentially acts as starpu_rbtree_lookup() but if no entry matched the key, an additional step is performed to obtain the next or previous node, depending on the direction (left or right).

The constraints that apply to the key parameter are the same as for starpu_rbtree_lookup().

6.55.1.5 starpu_rbtree_insert

```
#define starpu_rbtree_insert(
            tree,
            node,
            cmp_fn )
```
Insert a node in a tree.

This macro performs a standard lookup to obtain the insertion point of the given node in the tree (it is assumed that the inserted node never compares equal to any other entry in the tree) and links the node. It then checks red-black rules violations, and rebalances the tree if necessary.

Unlike starpu_rbtree_lookup(), the cmp_fn parameter must compare two complete entries, so it is suggested to use two different comparison inline functions, such as myobj_cmp_lookup() and myobj_cmp_insert(). There is no guarantee about the order of the nodes given to the comparison function.

See starpu_rbtree_lookup().

6.55.1.6 starpu_rbtree_lookup_slot

```
#define starpu_rbtree_lookup_slot(
            tree,
            key,
            cmp_fn,
            slot )
```
Look up a node/slot pair in a tree.

This macro essentially acts as starpu_rbtree_lookup() but in addition to a node, it also returns a slot, which identifies an insertion point in the tree. If the returned node is null, the slot can be used by starpu_rbtree_insert_slot() to insert without the overhead of an additional lookup. The slot is a simple uintptr_t integer.

The constraints that apply to the key parameter are the same as for starpu_rbtree_lookup().

6.55.1.7 starpu_rbtree_first

```
#define starpu_rbtree_first(
            tree )
```
Return the first node of a tree.

6.55.1.8 starpu_rbtree_last

```
#define starpu_rbtree_last(
            tree )
```
Return the last node of a tree.

6.55.1.9 starpu_rbtree_prev

```
#define starpu_rbtree_prev(
            node )
```
Return the node previous to the given node.

6.55.1.10   starpu_rbtree_next

```
#define starpu_rbtree_next(
               node )
```
Return the node next to the given node.

6.55.1.11   starpu_rbtree_for_each_remove

```
#define starpu_rbtree_for_each_remove(
               tree,
               node,
               tmp )
```
Forge a loop to process all nodes of a tree, removing them when visited.

This macro can only be used to destroy a tree, so that the resources used by the entries can be released by the user. It basically removes all nodes without doing any color checking.

After completion, all nodes and the tree root member are stale.

6.55.2   Function Documentation

6.55.2.1   starpu_rbtree_init()

```
static void starpu_rbtree_init (
               struct starpu_rbtree * tree )  [inline], [static]
```
Initialize a tree.

6.55.2.2   starpu_rbtree_node_init()

```
static void starpu_rbtree_node_init (
               struct starpu_rbtree_node * node )  [inline], [static]
```
Initialize a node.

A node is in no tree when its parent points to itself.

6.55.2.3   starpu_rbtree_empty()

```
static int starpu_rbtree_empty (
               const struct starpu_rbtree * tree )  [inline], [static]
```
Return true if tree is empty.

6.55.2.4   starpu_rbtree_insert_slot()

```
static void starpu_rbtree_insert_slot (
               struct starpu_rbtree * tree,
               uintptr_t slot,
               struct starpu_rbtree_node * node )  [inline], [static]
```
Insert a node at an insertion point in a tree.

This macro essentially acts as starpu_rbtree_insert() except that it doesn't obtain the insertion point with a standard lookup. The insertion point is obtained by calling starpu_rbtree_lookup_slot(). In addition, the new node must not compare equal to an existing node in the tree (i.e. the slot must denote a null node).

6.55.2.5   starpu_rbtree_remove()

```
void starpu_rbtree_remove (
               struct starpu_rbtree * tree,
               struct starpu_rbtree_node * node )
```
Remove a node from a tree.

After completion, the node is stale.

6.56 rbtree_i.h File Reference

```
#include <assert.h>
```

Data Structures

- struct starpu_rbtree_node
- struct starpu_rbtree

Macros

- #define STARPU_RBTREE_COLOR_MASK
- #define **STARPU_RBTREE_PARENT_MASK**
- #define STARPU_RBTREE_COLOR_RED
- #define **STARPU_RBTREE_COLOR_BLACK**
- #define STARPU_RBTREE_SLOT_INDEX_MASK
- #define **STARPU_RBTREE_SLOT_PARENT_MASK**

Functions

- static int starpu_rbtree_check_alignment (const struct starpu_rbtree_node ∗node)
- static int starpu_rbtree_check_index (int index)
- static int starpu_rbtree_d2i (int diff)
- static struct starpu_rbtree_node ∗ starpu_rbtree_parent (const struct starpu_rbtree_node ∗node)
- static uintptr_t starpu_rbtree_slot (struct starpu_rbtree_node ∗parent, int index)
- static struct starpu_rbtree_node ∗ starpu_rbtree_slot_parent (uintptr_t slot)
- static int starpu_rbtree_slot_index (uintptr_t slot)
- void starpu_rbtree_insert_rebalance (struct starpu_rbtree ∗tree, struct starpu_rbtree_node ∗parent, int index, struct starpu_rbtree_node ∗node)
- struct starpu_rbtree_node ∗ starpu_rbtree_nearest (struct starpu_rbtree_node ∗parent, int index, int direction)
- struct starpu_rbtree_node ∗ starpu_rbtree_firstlast (const struct starpu_rbtree ∗tree, int direction)
- struct starpu_rbtree_node ∗ starpu_rbtree_walk (struct starpu_rbtree_node ∗node, int direction)
- struct starpu_rbtree_node ∗ starpu_rbtree_postwalk_deepest (const struct starpu_rbtree ∗tree)
- struct starpu_rbtree_node ∗ starpu_rbtree_postwalk_unlink (struct starpu_rbtree_node ∗node)

6.56.1 Data Structure Documentation

6.56.1.1 struct starpu_rbtree_node

Red-black node structure.

To reduce the number of branches and the instruction cache footprint, the left and right child pointers are stored in an array, and the symmetry of most tree operations is exploited by using left/right variables when referring to children.

In addition, this implementation assumes that all nodes are 4-byte aligned, so that the least significant bit of the parent member can be used to store the color of the node. This is true for all modern 32 and 64 bits architectures, as long as the nodes aren't embedded in structures with special alignment constraints such as member packing.

Data Fields

uintptr_t	parent	
struct starpu_rbtree_node ∗	children[2]	

6.56.1.2 struct starpu_rbtree

Red-black tree structure.

Data Fields

struct starpu_rbtree_node ∗	root	

6.56.2 Macro Definition Documentation

6.56.2.1 STARPU_RBTREE_COLOR_MASK

```
#define STARPU_RBTREE_COLOR_MASK
```
Masks applied on the parent member of a node to obtain either the color or the parent address.

6.56.2.2 STARPU_RBTREE_COLOR_RED

```
#define STARPU_RBTREE_COLOR_RED
```
Node colors.

6.56.2.3 STARPU_RBTREE_SLOT_INDEX_MASK

```
#define STARPU_RBTREE_SLOT_INDEX_MASK
```
Masks applied on slots to obtain either the child index or the parent address.

6.56.3 Function Documentation

6.56.3.1 starpu_rbtree_check_alignment()

```
static int starpu_rbtree_check_alignment (
            const struct starpu_rbtree_node ∗ node ) [inline], [static]
```
Return true if the given pointer is suitably aligned.

6.56.3.2 starpu_rbtree_check_index()

```
static int starpu_rbtree_check_index (
            int index ) [inline], [static]
```
Return true if the given index is a valid child index.

6.56.3.3 starpu_rbtree_d2i()

```
static int starpu_rbtree_d2i (
            int diff ) [inline], [static]
```
Convert the result of a comparison into an index in the children array (0 or 1).
This function is mostly used when looking up a node.

6.56.3.4 starpu_rbtree_parent()

```
static struct starpu_rbtree_node∗ starpu_rbtree_parent (
            const struct starpu_rbtree_node ∗ node ) [static]
```
Return the parent of a node.

6.56.3.5 starpu_rbtree_slot()

```
static uintptr_t starpu_rbtree_slot (
            struct starpu_rbtree_node * parent,
            int index )  [inline], [static]
```
Translate an insertion point into a slot.

6.56.3.6 starpu_rbtree_slot_parent()

```
static struct starpu_rbtree_node* starpu_rbtree_slot_parent (
            uintptr_t slot )  [static]
```
Extract the parent address from a slot.

6.56.3.7 starpu_rbtree_slot_index()

```
static int starpu_rbtree_slot_index (
            uintptr_t slot )  [inline], [static]
```
Extract the index from a slot.

6.56.3.8 starpu_rbtree_insert_rebalance()

```
void starpu_rbtree_insert_rebalance (
            struct starpu_rbtree * tree,
            struct starpu_rbtree_node * parent,
            int index,
            struct starpu_rbtree_node * node )
```
Insert a node in a tree, rebalancing it if necessary.

The index parameter is the index in the children array of the parent where the new node is to be inserted. It is ignored if the parent is null.

This function is intended to be used by the starpu_rbtree_insert() macro only.

6.56.3.9 starpu_rbtree_nearest()

```
struct starpu_rbtree_node* starpu_rbtree_nearest (
            struct starpu_rbtree_node * parent,
            int index,
            int direction )
```
Return the previous or next node relative to a location in a tree.

The parent and index parameters define the location, which can be empty. The direction parameter is either S←
TARPU_RBTREE_LEFT (to obtain the previous node) or STARPU_RBTREE_RIGHT (to obtain the next one).

6.56.3.10 starpu_rbtree_firstlast()

```
struct starpu_rbtree_node* starpu_rbtree_firstlast (
            const struct starpu_rbtree * tree,
            int direction )
```
Return the first or last node of a tree.

The direction parameter is either STARPU_RBTREE_LEFT (to obtain the first node) or STARPU_RBTREE_RIGHT (to obtain the last one).

6.56.3.11 starpu_rbtree_walk()

```
struct starpu_rbtree_node* starpu_rbtree_walk (
            struct starpu_rbtree_node * node,
            int direction )
```
Return the node next to, or previous to the given node.

The direction parameter is either STARPU_RBTREE_LEFT (to obtain the previous node) or STARPU_RBTREE←
_RIGHT (to obtain the next one).

6.56.3.12 starpu_rbtree_postwalk_deepest()

struct starpu_rbtree_node∗ starpu_rbtree_postwalk_deepest (
            const struct starpu_rbtree ∗ *tree*)

Return the left-most deepest node of a tree, which is the starting point of the postorder traversal performed by
starpu_rbtree_for_each_remove().

6.56.3.13 starpu_rbtree_postwalk_unlink()

struct starpu_rbtree_node∗ starpu_rbtree_postwalk_unlink (
            struct starpu_rbtree_node ∗ *node*)

Unlink a node from its tree and return the next (right) node in postorder.

6.57 regression.h File Reference

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>
#include <core/perfmodel/perfmodel.h>
#include <starpu.h>
```

Functions

- int **_starpu_regression_non_linear_power** (struct starpu_perfmodel_history_list ∗ptr, double ∗a, double ∗b, double ∗c)

6.58 rwlock.h File Reference

```
#include <stdint.h>
#include <starpu.h>
```

Data Structures

- struct _starpu_rw_lock

Functions

- void _starpu_init_rw_lock (struct _starpu_rw_lock ∗lock)
- void _starpu_take_rw_lock_write (struct _starpu_rw_lock ∗lock)
- void _starpu_take_rw_lock_read (struct _starpu_rw_lock ∗lock)
- int _starpu_take_rw_lock_write_try (struct _starpu_rw_lock ∗lock)
- int _starpu_take_rw_lock_read_try (struct _starpu_rw_lock ∗lock)
- void _starpu_release_rw_lock (struct _starpu_rw_lock ∗lock)

6.58.1 Data Structure Documentation

6.58.1.1 struct _starpu_rw_lock

Dummy implementation of a RW-lock using a spinlock.

Data Fields

uint32_t	busy	
uint8_t	writer	
uint16_t	readercnt	

6.58.2 Function Documentation

6.58.2.1 _starpu_init_rw_lock()

```
void _starpu_init_rw_lock (
             struct _starpu_rw_lock * lock )
```
Initialize the RW-lock

6.58.2.2 _starpu_take_rw_lock_write()

```
void _starpu_take_rw_lock_write (
             struct _starpu_rw_lock * lock )
```
Grab the RW-lock in a write mode

6.58.2.3 _starpu_take_rw_lock_read()

```
void _starpu_take_rw_lock_read (
             struct _starpu_rw_lock * lock )
```
Grab the RW-lock in a read mode

6.58.2.4 _starpu_take_rw_lock_write_try()

```
int _starpu_take_rw_lock_write_try (
             struct _starpu_rw_lock * lock )
```
Try to grab the RW-lock in a write mode. Returns 0 in case of success, -1 otherwise.

6.58.2.5 _starpu_take_rw_lock_read_try()

```
int _starpu_take_rw_lock_read_try (
             struct _starpu_rw_lock * lock )
```
Try to grab the RW-lock in a read mode. Returns 0 in case of success, -1 otherwise.

6.58.2.6 _starpu_release_rw_lock()

```
void _starpu_release_rw_lock (
             struct _starpu_rw_lock * lock )
```
Unlock the RW-lock.

6.59 sched_component.h File Reference

```
#include <starpu_sched_component.h>
```

Functions

- void _starpu_sched_component_lock_all_workers (void)
- void **_starpu_sched_component_unlock_all_workers** (void)
- void **_starpu_sched_component_workers_destroy** (void)
- struct _starpu_worker ∗ **_starpu_sched_component_worker_get_worker** (struct starpu_sched_↩
  component ∗)
- struct starpu_bitmap ∗ **_starpu_get_worker_mask** (unsigned sched_ctx_id)

6.59.1 Function Documentation

6.59.1.1 _starpu_sched_component_lock_all_workers()

```
void _starpu_sched_component_lock_all_workers (
            void  )
```

lock and unlock drivers for modifying schedulers

6.60 sched_ctx.h File Reference

```
#include <starpu.h>
#include <starpu_sched_ctx.h>
#include <starpu_sched_ctx_hypervisor.h>
#include <starpu_scheduler.h>
#include <common/config.h>
#include <common/barrier_counter.h>
#include <common/utils.h>
#include <profiling/profiling.h>
#include <semaphore.h>
#include <core/task.h>
#include "sched_ctx_list.h"
#include <hwloc.h>
```

Data Structures

- struct _starpu_sched_ctx
- struct _starpu_ctx_change

Macros

- #define **NO_RESIZE**
- #define **REQ_RESIZE**
- #define **DO_RESIZE**
- #define **STARPU_GLOBAL_SCHED_CTX**
- #define **STARPU_NMAXSMS**
- #define **_starpu_sched_ctx_get_sched_ctx_for_worker_and_job**(w, j)
- #define **STARPU_SCHED_CTX_CHECK_LOCK**(sched_ctx_id)

Functions

- void _starpu_init_all_sched_ctxs (struct _starpu_machine_config ∗config)
- struct _starpu_sched_ctx ∗ _starpu_create_sched_ctx (struct starpu_sched_policy ∗policy, int ∗workerid, int nworkerids, unsigned is_init_sched, const char ∗sched_name, int min_prio_set, int min_prio, int max_←prio_set, int max_prio, unsigned awake_workers, void(∗sched_policy_init)(unsigned), void ∗user_data, int nsub_ctxs, int ∗sub_ctxs, int nsms)
- void _starpu_delete_all_sched_ctxs ()
- int _starpu_wait_for_all_tasks_of_sched_ctx (unsigned sched_ctx_id)
- int _starpu_wait_for_n_submitted_tasks_of_sched_ctx (unsigned sched_ctx_id, unsigned n)
- void _starpu_decrement_nsubmitted_tasks_of_sched_ctx (unsigned sched_ctx_id)
- void **_starpu_increment_nsubmitted_tasks_of_sched_ctx** (unsigned sched_ctx_id)
- int **_starpu_get_nsubmitted_tasks_of_sched_ctx** (unsigned sched_ctx_id)
- int **_starpu_check_nsubmitted_tasks_of_sched_ctx** (unsigned sched_ctx_id)
- void **_starpu_decrement_nready_tasks_of_sched_ctx** (unsigned sched_ctx_id, double ready_flops)
- unsigned **_starpu_increment_nready_tasks_of_sched_ctx** (unsigned sched_ctx_id, double ready_flops, struct starpu_task ∗task)
- int **_starpu_wait_for_no_ready_of_sched_ctx** (unsigned sched_ctx_id)
- int _starpu_get_index_in_ctx_of_workerid (unsigned sched_ctx, unsigned workerid)

- starpu_pthread_mutex_t ∗ _starpu_get_sched_mutex (struct _starpu_sched_ctx ∗sched_ctx, int worker)
- int _starpu_get_workers_of_sched_ctx (unsigned sched_ctx_id, int ∗pus, enum starpu_worker_archtype arch)
- void _starpu_worker_gets_out_of_ctx (unsigned sched_ctx_id, struct _starpu_worker ∗worker)
- unsigned _starpu_worker_belongs_to_a_sched_ctx (int workerid, unsigned sched_ctx_id)
- unsigned _starpu_sched_ctx_last_worker_awake (struct _starpu_worker ∗worker)
- unsigned _starpu_sched_ctx_get_current_context ()
- int _starpu_workers_able_to_execute_task (struct starpu_task ∗task, struct _starpu_sched_ctx ∗sched_ctx)
- void **_starpu_fetch_tasks_from_empty_ctx_list** (struct _starpu_sched_ctx ∗sched_ctx)
- unsigned **_starpu_sched_ctx_allow_hypervisor** (unsigned sched_ctx_id)
- struct starpu_perfmodel_arch ∗ **_starpu_sched_ctx_get_perf_archtype** (unsigned sched_ctx)
- void _starpu_sched_ctx_post_exec_task_cb (int workerid, struct starpu_task ∗task, size_t data_size, uint32_t footprint)
- void **starpu_sched_ctx_add_combined_workers** (int ∗combined_workers_to_add, unsigned n_↩ combined_workers_to_add, unsigned sched_ctx_id)
- struct _starpu_sched_ctx ∗ __starpu_sched_ctx_get_sched_ctx_for_worker_and_job (struct _starpu_worker ∗worker, struct _starpu_job ∗j)
- static struct _starpu_sched_ctx ∗ **_starpu_get_sched_ctx_struct** (unsigned id)
- static int **_starpu_sched_ctx_check_write_locked** (unsigned sched_ctx_id)
- static void **_starpu_sched_ctx_lock_write** (unsigned sched_ctx_id)
- static void **_starpu_sched_ctx_unlock_write** (unsigned sched_ctx_id)
- static void **_starpu_sched_ctx_lock_read** (unsigned sched_ctx_id)
- static void **_starpu_sched_ctx_unlock_read** (unsigned sched_ctx_id)
- static unsigned **_starpu_sched_ctx_worker_is_master_for_child_ctx** (unsigned sched_ctx_id, unsigned workerid, struct starpu_task ∗task)
- void _starpu_worker_apply_deferred_ctx_changes (void)

6.60.1 Data Structure Documentation

6.60.1.1 struct _starpu_ctx_change

per-worker list of deferred ctx_change ops

Data Fields

int	sched_ctx_id	
int	op	
int	nworkers_to_notify	
int ∗	workerids_to_notify	
int	nworkers_to_change	
int ∗	workerids_to_change	

6.60.2 Function Documentation

6.60.2.1 _starpu_init_all_sched_ctxs()

```
void _starpu_init_all_sched_ctxs (
            struct _starpu_machine_config * config )
```
init sched_ctx_id of all contextes

6.60.2.2 _starpu_create_sched_ctx()

```
struct _starpu_sched_ctx* _starpu_create_sched_ctx (
```

```
                struct starpu_sched_policy * policy,
                int * workerid,
                int nworkerids,
                unsigned is_init_sched,
                const char * sched_name,
                int min_prio_set,
                int min_prio,
                int max_prio_set,
                int max_prio,
                unsigned awake_workers,
                void(*)(unsigned) sched_policy_init,
                void * user_data,
                int nsub_ctxs,
                int * sub_ctxs,
                int nsms )
```
allocate all structures belonging to a context

6.60.2.3 _starpu_delete_all_sched_ctxs()

```
void _starpu_delete_all_sched_ctxs ( )
```
delete all sched_ctx

6.60.2.4 _starpu_wait_for_all_tasks_of_sched_ctx()

```
int _starpu_wait_for_all_tasks_of_sched_ctx (
                unsigned sched_ctx_id )
```
This function waits until all the tasks that were already submitted to a specific context have been executed.

6.60.2.5 _starpu_wait_for_n_submitted_tasks_of_sched_ctx()

```
int _starpu_wait_for_n_submitted_tasks_of_sched_ctx (
                unsigned sched_ctx_id,
                unsigned n )
```
This function waits until at most n tasks are still submitted.

6.60.2.6 _starpu_decrement_nsubmitted_tasks_of_sched_ctx()

```
void _starpu_decrement_nsubmitted_tasks_of_sched_ctx (
                unsigned sched_ctx_id )
```
In order to implement starpu_wait_for_all_tasks_of_ctx, we keep track of the number of task currently submitted to the context

6.60.2.7 _starpu_get_index_in_ctx_of_workerid()

```
int _starpu_get_index_in_ctx_of_workerid (
                unsigned sched_ctx,
                unsigned workerid )
```
Return the corresponding index of the workerid in the ctx table

6.60.2.8 _starpu_get_sched_mutex()

```
starpu_pthread_mutex_t* _starpu_get_sched_mutex (
                struct _starpu_sched_ctx * sched_ctx,
                int worker )
```
Get the mutex corresponding to the global workerid

6.60.2.9 _starpu_get_workers_of_sched_ctx()

```
int _starpu_get_workers_of_sched_ctx (
                unsigned sched_ctx_id,
```

```
            int ∗ pus,
            enum starpu_worker_archtype arch )
```
Get workers belonging to a certain context, it returns the number of workers take care: no mutex taken, the list of workers might not be updated

6.60.2.10 _starpu_worker_gets_out_of_ctx()

```
void _starpu_worker_gets_out_of_ctx (
            unsigned sched_ctx_id,
            struct _starpu_worker ∗ worker )
```
Let the worker know it does not belong to the context and that it should stop poping from it

6.60.2.11 _starpu_worker_belongs_to_a_sched_ctx()

```
unsigned _starpu_worker_belongs_to_a_sched_ctx (
            int workerid,
            unsigned sched_ctx_id )
```
Check if the worker belongs to another sched_ctx

6.60.2.12 _starpu_sched_ctx_last_worker_awake()

```
unsigned _starpu_sched_ctx_last_worker_awake (
            struct _starpu_worker ∗ worker )
```
indicates wheather this worker should go to sleep or not (if it is the last one awake in a context he should better keep awake)

6.60.2.13 _starpu_sched_ctx_get_current_context()

```
unsigned _starpu_sched_ctx_get_current_context ( )
```
If starpu_sched_ctx_set_context() has been called, returns the context id set by its last call, or the id of the initial context

6.60.2.14 _starpu_workers_able_to_execute_task()

```
int _starpu_workers_able_to_execute_task (
            struct starpu_task ∗ task,
            struct _starpu_sched_ctx ∗ sched_ctx )
```
verify that some worker can execute a certain task

6.60.2.15 _starpu_sched_ctx_post_exec_task_cb()

```
void _starpu_sched_ctx_post_exec_task_cb (
            int workerid,
            struct starpu_task ∗ task,
            size_t data_size,
            uint32_t footprint )
```
Notifies the hypervisor that a tasks was poped from the workers' list

6.60.2.16 __starpu_sched_ctx_get_sched_ctx_for_worker_and_job()

```
struct _starpu_sched_ctx∗ __starpu_sched_ctx_get_sched_ctx_for_worker_and_job (
            struct _starpu_worker ∗ worker,
            struct _starpu_job ∗ j )
```
if the worker is the master of a parallel context, and the job is meant to be executed on this parallel context, return a pointer to the context

6.60.2.17 _starpu_worker_apply_deferred_ctx_changes()

```
void _starpu_worker_apply_deferred_ctx_changes (
            void  )
```

Go through the list of deferred ctx changes of the current worker and apply any ctx change operation found until the list is empty

6.61 sched_ctx_list.h File Reference

Data Structures

- struct [_starpu_sched_ctx_list](#)
- struct [_starpu_sched_ctx_elt](#)
- struct [_starpu_sched_ctx_list_iterator](#)

Functions

- struct [_starpu_sched_ctx_elt](#) ∗ [_starpu_sched_ctx_elt_find](#) (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- void **_starpu_sched_ctx_elt_ensure_consistency** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched↩
  _ctx)
- void **_starpu_sched_ctx_elt_init** (struct [_starpu_sched_ctx_elt](#) ∗elt, unsigned sched_ctx)
- struct [_starpu_sched_ctx_elt](#) ∗ **_starpu_sched_ctx_elt_add_after** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- struct [_starpu_sched_ctx_elt](#) ∗ **_starpu_sched_ctx_elt_add_before** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- struct [_starpu_sched_ctx_elt](#) ∗ **_starpu_sched_ctx_elt_add** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- void **_starpu_sched_ctx_elt_remove** (struct [_starpu_sched_ctx_list](#) ∗list, struct [_starpu_sched_ctx_elt](#) ∗elt)
- int **_starpu_sched_ctx_elt_exists** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- int **_starpu_sched_ctx_elt_get_priority** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- struct [_starpu_sched_ctx_list](#) ∗ [_starpu_sched_ctx_list_find](#) (struct [_starpu_sched_ctx_list](#) ∗list, unsigned prio)
- struct [_starpu_sched_ctx_elt](#) ∗ **_starpu_sched_ctx_list_add_prio** (struct [_starpu_sched_ctx_list](#) ∗∗list, unsigned prio, unsigned sched_ctx)
- int **_starpu_sched_ctx_list_add** (struct [_starpu_sched_ctx_list](#) ∗∗list, unsigned sched_ctx)
- void **_starpu_sched_ctx_list_remove_elt** (struct [_starpu_sched_ctx_list](#) ∗∗list, struct [_starpu_sched_ctx↩
  _elt](#) ∗rm)
- int **_starpu_sched_ctx_list_remove** (struct [_starpu_sched_ctx_list](#) ∗∗list, unsigned sched_ctx)
- int **_starpu_sched_ctx_list_move** (struct [_starpu_sched_ctx_list](#) ∗∗list, unsigned sched_ctx, unsigned prio_to)
- int **_starpu_sched_ctx_list_exists** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned prio)
- void **_starpu_sched_ctx_list_remove_all** (struct [_starpu_sched_ctx_list](#) ∗list)
- void **_starpu_sched_ctx_list_delete** (struct [_starpu_sched_ctx_list](#) ∗∗list)
- int [_starpu_sched_ctx_list_push_event](#) (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- int **_starpu_sched_ctx_list_pop_event** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- int **_starpu_sched_ctx_list_pop_all_event** (struct [_starpu_sched_ctx_list](#) ∗list, unsigned sched_ctx)
- int [_starpu_sched_ctx_list_iterator_init](#) (struct [_starpu_sched_ctx_list](#) ∗list, struct [_starpu_sched_ctx_list_↩
  iterator](#) ∗it)
- int **_starpu_sched_ctx_list_iterator_has_next** (struct [_starpu_sched_ctx_list_iterator](#) ∗it)
- struct [_starpu_sched_ctx_elt](#) ∗ **_starpu_sched_ctx_list_iterator_get_next** (struct [_starpu_sched_ctx_↩
  list_iterator](#) ∗it)

6.61.1 Data Structure Documentation

6.61.1.1 struct _starpu_sched_ctx_list

Data Fields

struct _starpu_sched_ctx_list ∗	prev	
struct _starpu_sched_ctx_list ∗	next	
struct _starpu_sched_ctx_elt ∗	head	
unsigned	priority	

6.61.1.2   struct _starpu_sched_ctx_elt

Represents a circular list of sched context.

Data Fields

struct _starpu_sched_ctx_elt ∗	prev	
struct _starpu_sched_ctx_elt ∗	next	
struct _starpu_sched_ctx_list ∗	parent	
unsigned	sched_ctx	
long	task_number	
unsigned	last_poped	

6.61.1.3   struct _starpu_sched_ctx_list_iterator

Data Fields

struct _starpu_sched_ctx_list ∗	list_head	
struct _starpu_sched_ctx_elt ∗	cursor	

6.61.2   Function Documentation

6.61.2.1   _starpu_sched_ctx_elt_find()

struct _starpu_sched_ctx_elt∗ _starpu_sched_ctx_elt_find (
            struct _starpu_sched_ctx_list ∗ *list,*
            unsigned *sched_ctx*)

Element (sched_ctx) level operations

6.61.2.2   _starpu_sched_ctx_list_find()

struct _starpu_sched_ctx_list∗ _starpu_sched_ctx_list_find (
            struct _starpu_sched_ctx_list ∗ *list,*
            unsigned *prio*)

List (priority) level operations

6.61.2.3   _starpu_sched_ctx_list_push_event()

int _starpu_sched_ctx_list_push_event (
            struct _starpu_sched_ctx_list ∗ *list,*
            unsigned *sched_ctx*)

Task number management

6.61.2.4   _starpu_sched_ctx_list_iterator_init()

int _starpu_sched_ctx_list_iterator_init (

```
            struct _starpu_sched_ctx_list * list,
            struct _starpu_sched_ctx_list_iterator * it )
```
Iterator operations

6.62 sched_policy.h File Reference

```
#include <starpu.h>
#include <signal.h>
#include <core/workers.h>
#include <core/sched_ctx.h>
#include <starpu_scheduler.h>
#include <core/simgrid.h>
```

Macros

- #define **_STARPU_SCHED_BEGIN**
- #define **_STARPU_SCHED_END**
- #define **_STARPU_TASK_BREAK_ON**(task, what)

Functions

- void **_starpu_sched_init** (void)
- struct starpu_sched_policy * **_starpu_get_sched_policy** (struct _starpu_sched_ctx *sched_ctx)
- void **_starpu_init_sched_policy** (struct _starpu_machine_config *config, struct _starpu_sched_ctx *sched_ctx, struct starpu_sched_policy *policy)
- void **_starpu_deinit_sched_policy** (struct _starpu_sched_ctx *sched_ctx)
- struct starpu_sched_policy * **_starpu_select_sched_policy** (struct _starpu_machine_config *config, const char *required_policy)
- void **_starpu_sched_task_submit** (struct starpu_task *task)
- void **_starpu_sched_do_schedule** (unsigned sched_ctx_id)
- int **_starpu_push_task** (struct _starpu_job *task)
- int **_starpu_repush_task** (struct _starpu_job *task)
- int **_starpu_push_task_to_workers** (struct starpu_task *task)
- struct starpu_task * **_starpu_pop_task** (struct _starpu_worker *worker)
- struct starpu_task * **_starpu_pop_every_task** (struct _starpu_sched_ctx *sched_ctx)
- void **_starpu_sched_post_exec_hook** (struct starpu_task *task)
- int **_starpu_pop_task_end** (struct starpu_task *task)
- void **_starpu_wait_on_sched_event** (void)
- struct starpu_task * **_starpu_create_conversion_task** (starpu_data_handle_t handle, unsigned int node) STARPU_ATTRIBUTE_MALLOC
- struct starpu_task * **_starpu_create_conversion_task_for_arch** (starpu_data_handle_t handle, enum starpu_node_kind node_kind) STARPU_ATTRIBUTE_MALLOC
- void **_starpu_sched_pre_exec_hook** (struct starpu_task *task)
- void **_starpu_print_idle_time** ()

Variables

- struct starpu_sched_policy **_starpu_sched_lws_policy**
- struct starpu_sched_policy **_starpu_sched_ws_policy**
- struct starpu_sched_policy **_starpu_sched_prio_policy**
- struct starpu_sched_policy **_starpu_sched_random_policy**
- struct starpu_sched_policy **_starpu_sched_dm_policy**
- struct starpu_sched_policy **_starpu_sched_dmda_policy**
- struct starpu_sched_policy **_starpu_sched_dmda_prio_policy**
- struct starpu_sched_policy **_starpu_sched_dmda_ready_policy**

- [struct](#) starpu_sched_policy **_starpu_sched_dmda_sorted_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_dmda_sorted_decision_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_eager_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_parallel_heft_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_peager_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_heteroprio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_eager_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_eager_prefetching_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_eager_prio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_gemm_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_prio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_prio_prefetching_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_random_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_random_prio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_random_prefetching_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_random_prio_prefetching_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_parallel_random_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_parallel_random_prio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_ws_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_heft_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_heft_prio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_heft2_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_heteroprio_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_heteroprio_heft_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_modular_parallel_heft_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_graph_test_policy**
- [struct](#) starpu_sched_policy **_starpu_sched_tree_heft_hierarchical_policy**
- long **_starpu_task_break_on_push**
- long **_starpu_task_break_on_sched**
- long **_starpu_task_break_on_pop**
- long **_starpu_task_break_on_exec**

6.62.1 Function Documentation

6.62.1.1 _starpu_push_task_to_workers()

```
int _starpu_push_task_to_workers (
            struct starpu_task * task )
```
actually pushes the tasks to the specific worker or to the scheduler

6.62.1.2 _starpu_pop_task()

```
struct starpu_task* _starpu_pop_task (
            struct _starpu_worker * worker )
```
pop a task that can be executed on the worker

6.62.1.3 _starpu_pop_every_task()

```
struct starpu_task* _starpu_pop_every_task (
            struct _starpu_sched_ctx * sched_ctx )
```
pop every task that can be executed on the worker

6.63 simgrid.h File Reference

```
#include <xbt/xbt_os_time.h>
```

Data Structures

- struct _starpu_pthread_args

Macros

- #define **MAX_TSD**
- #define **STARPU_MPI_AS_PREFIX**
- #define **_starpu_simgrid_running_smpi**()
- #define **_starpu_simgrid_cuda_malloc_cost**()
- #define **_starpu_simgrid_queue_malloc_cost**()
- #define **_starpu_simgrid_task_submit_cost**()
- #define **_starpu_simgrid_fetching_input_cost**()
- #define **_starpu_simgrid_sched_cost**()
- #define **_SIMGRID_TIMER_BEGIN**(cond)
- #define **_SIMGRID_TIMER_END**
- #define _starpu_simgrid_data_new(size)
- #define **_starpu_simgrid_data_increase**(size)
- #define **_starpu_simgrid_data_alloc**(size)
- #define **_starpu_simgrid_data_free**(size)
- #define **_starpu_simgrid_data_transfer**(size, src_node, dst_node)

Functions

- void **_starpu_start_simgrid** (int ∗argc, char ∗∗argv)
- void **_starpu_simgrid_init_early** (int ∗argc, char ∗∗∗argv)
- void **_starpu_simgrid_init** (void)
- void **_starpu_simgrid_deinit** (void)
- void **_starpu_simgrid_deinit_late** (void)
- void **_starpu_simgrid_actor_setup** (void)
- void **_starpu_simgrid_wait_tasks** (int workerid)
- void **_starpu_simgrid_submit_job** (int workerid, struct _starpu_job ∗job, struct starpu_perfmodel_arch ∗perf_arch, double length, unsigned ∗finished)
- int **_starpu_simgrid_transfer** (size_t size, unsigned src_node, unsigned dst_node, struct _starpu_data_↩ request ∗req)
- int **_starpu_simgrid_wait_transfer_event** (union _starpu_async_channel_event ∗event)
- int **_starpu_simgrid_test_transfer_event** (union _starpu_async_channel_event ∗event)
- void **_starpu_simgrid_sync_gpus** (void)
- int _starpu_simgrid_get_nbhosts (const char ∗prefix)
- unsigned long long **_starpu_simgrid_get_memsize** (const char ∗prefix, unsigned devid)
- starpu_sg_host_t **_starpu_simgrid_get_host_by_name** (const char ∗name)
- starpu_sg_host_t **_starpu_simgrid_get_memnode_host** (unsigned node)
- starpu_sg_host_t **_starpu_simgrid_get_host_by_worker** (struct _starpu_worker ∗worker)
- void **_starpu_simgrid_get_platform_path** (int version, char ∗path, size_t maxlen)
- msg_as_t **_starpu_simgrid_get_as_by_name** (const char ∗name)
- int **starpu_mpi_world_rank** (void)
- int **_starpu_mpi_simgrid_init** (int argc, char ∗argv[ ])
- void _starpu_simgrid_count_ngpus (void)
- void **_starpu_simgrid_xbt_thread_create** (const char ∗name, void_f_pvoid_t code, void ∗param)

Variables

- starpu_pthread_queue_t **_starpu_simgrid_transfer_queue** [STARPU_MAXNODES]
- starpu_pthread_queue_t **_starpu_simgrid_task_queue** [STARPU_NMAXWORKERS]

6.63.1   Macro Definition Documentation

6.63.1.1   _starpu_simgrid_data_new

```
#define _starpu_simgrid_data_new(
              size )
```
Experimental functions for OOC stochastic analysis

6.63.2   Function Documentation

6.63.2.1   _starpu_simgrid_get_nbhosts()

```
int _starpu_simgrid_get_nbhosts (
              const char * prefix )
```
Return the number of hosts prefixed by PREFIX

6.63.2.2   _starpu_simgrid_count_ngpus()

```
void _starpu_simgrid_count_ngpus (
              void  )
```
Called at initialization to count how many GPUs are interfering with each bus

6.64   sink_common.h File Reference

```
#include <common/config.h>
```

6.65   sort_data_handles.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <stdlib.h>
#include <stdarg.h>
#include <core/jobs.h>
#include <datawizard/coherency.h>
#include <datawizard/memalloc.h>
```

Functions

- void _starpu_sort_task_handles (struct _starpu_data_descr descr[ ], unsigned nbuffers)

6.65.1   Function Documentation

6.65.1.1 _starpu_sort_task_handles()

```
void _starpu_sort_task_handles (
            struct _starpu_data_descr descr[],
            unsigned nbuffers )
```
To avoid deadlocks, we reorder the different buffers accessed to by the task so that we always grab the rw-lock associated to the handles in the same order.

6.66 source_common.h File Reference

6.67 starpu_clusters_create.h File Reference

```
#include <starpu.h>
#include <core/workers.h>
#include <common/list.h>
#include <string.h>
#include <omp.h>
```

6.68 starpu_data_cpy.h File Reference

```
#include <starpu.h>
```

Functions

- int **_starpu_data_cpy** (starpu_data_handle_t dst_handle, starpu_data_handle_t src_handle, int asynchronous, void(∗callback_func)(void ∗), void ∗callback_arg, int reduction, struct starpu_task ∗reduction_↩ dep_task)

6.69 starpu_debug_helpers.h File Reference

```
#include <starpu.h>
#include <starpu_config.h>
#include <starpu_util.h>
```

Functions

- void _starpu_benchmark_ping_pong (starpu_data_handle_t handle, unsigned node0, unsigned node1, unsigned niter)
- void _starpu_debug_display_structures_size (FILE ∗stream)

6.69.1 Function Documentation

6.69.1.1 _starpu_benchmark_ping_pong()

```
void _starpu_benchmark_ping_pong (
            starpu_data_handle_t handle,
            unsigned node0,
            unsigned node1,
            unsigned niter )
```
Perform a ping pong between the two memory nodes

6.69.1.2 _starpu_debug_display_structures_size()

```
void _starpu_debug_display_structures_size (
             FILE * stream )
```
Display the size of different data structures

6.70 starpu_fxt.h File Reference

```
#include <starpu.h>
#include <starpu_config.h>
#include <common/config.h>
```

6.71 starpu_parameters.h File Reference

Macros

- #define **_STARPU_CPU_ALPHA**
- #define **_STARPU_CUDA_ALPHA**
- #define **_STARPU_OPENCL_ALPHA**
- #define **_STARPU_MIC_ALPHA**
- #define **_STARPU_MPI_MS_ALPHA**

6.72 starpu_spinlock.h File Reference

```
#include <errno.h>
#include <stdint.h>
#include <common/config.h>
#include <common/fxt.h>
#include <common/thread.h>
#include <starpu.h>
```

Data Structures

- struct _starpu_spinlock

Macros

- #define **_starpu_spin_destroy**(_lock)
- #define **_starpu_spin_checklocked**(_lock)
- #define **_starpu_spin_lock**(lock)
- #define **_starpu_spin_trylock**(lock)
- #define **_starpu_spin_unlock**(lock)
- #define **STARPU_SPIN_MAXTRY**

Functions

- static int **_starpu_spin_init** (struct _starpu_spinlock ∗lock)
- static int **__starpu_spin_lock** (struct _starpu_spinlock ∗lock, const char ∗file STARPU_ATTRIBUTE_UN↩
  USED, int line STARPU_ATTRIBUTE_UNUSED, const char ∗func STARPU_ATTRIBUTE_UNUSED)
- static int **__starpu_spin_trylock** (struct _starpu_spinlock ∗lock, const char ∗file STARPU_ATTRIBUTE_↩
  UNUSED, int line STARPU_ATTRIBUTE_UNUSED, const char ∗func STARPU_ATTRIBUTE_UNUSED)
- static int **__starpu_spin_unlock** (struct _starpu_spinlock ∗lock, const char ∗file STARPU_ATTRIBUTE_↩
  UNUSED, int line STARPU_ATTRIBUTE_UNUSED, const char ∗func STARPU_ATTRIBUTE_UNUSED)

6.72.1 Data Structure Documentation

6.72.1.1 struct _starpu_spinlock

Data Fields

starpu_pthread_spinlock_t	lock	

6.73 starpu_task_insert_utils.h File Reference

```
#include <stdlib.h>
#include <stdarg.h>
#include <starpu.h>
```

Typedefs

- typedef void(∗ **_starpu_callback_func_t**) (void ∗)

Functions

- int **_starpu_codelet_pack_args** (void ∗∗arg_buffer, size_t ∗arg_buffer_size, va_list varg_list)
- int **_starpu_task_insert_create** (struct starpu_codelet ∗cl, struct starpu_task ∗task, va_list varg_list)
- int **_fstarpu_task_insert_create** (struct starpu_codelet ∗cl, struct starpu_task ∗task, void ∗∗arglist)

6.74 tags.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/starpu_spinlock.h>
#include <core/dependencies/cg.h>
```

Data Structures

- struct _starpu_tag

Macros

- #define **_STARPU_TAG_SIZE**

Enumerations

- enum _starpu_tag_state {
  STARPU_INVALID_STATE, STARPU_ASSOCIATED, STARPU_BLOCKED, STARPU_READY,
  STARPU_DONE }

Functions

- void **_starpu_init_tags** (void)
- void **_starpu_notify_tag_dependencies** (struct _starpu_tag ∗tag)
- void **_starpu_notify_job_start_tag_dependencies** (struct _starpu_tag ∗tag, _starpu_notify_job_start_data ∗data)
- void **_starpu_tag_declare** (starpu_tag_t id, struct _starpu_job ∗job)
- void _starpu_tag_set_ready (struct _starpu_tag ∗tag)
- unsigned **_starpu_submit_job_enforce_task_deps** (struct _starpu_job ∗j)
- void **_starpu_tag_clear** (void)

6.74.1 Data Structure Documentation

6.74.1.1 struct _starpu_tag

Data Fields

struct _starpu_spinlock	lock	Lock for this structure. Locking order is in dependency order: a tag must not be locked before locking a tag it depends on
starpu_tag_t	id	an identifier for the task
enum _starpu_tag_state	state	
struct _starpu_cg_list	tag_successors	
struct _starpu_job *	job	which job is associated to the tag if any ?
unsigned	is_assigned	
unsigned	is_submitted	

6.74.2 Enumeration Type Documentation

6.74.2.1 _starpu_tag_state

enum _starpu_tag_state

Enumerator

STARPU_INVALID_STATE	this tag is not declared by any task
STARPU_ASSOCIATED	_starpu_tag_declare was called to associate the tag to a task
STARPU_BLOCKED	some task dependencies are not fulfilled yet
STARPU_READY	the task can be (or has been) submitted to the scheduler (all deps fulfilled)
STARPU_DONE	the task has been performed

6.74.3 Function Documentation

6.74.3.1 _starpu_tag_set_ready()

```
void _starpu_tag_set_ready (
            struct _starpu_tag * tag )
```
lock should be taken, and this releases it

6.75 task.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <core/jobs.h>
```

Macros

- #define **_STARPU_JOB_UNSET**
- #define **_STARPU_JOB_SETTING**
- #define **_STARPU_TASK_SET_INTERFACE**(task, interface, i)
- #define **_STARPU_TASK_GET_INTERFACES**(task)

Functions

- void _starpu_task_destroy (struct starpu_task ∗task)
- int _starpu_task_test_termination (struct starpu_task ∗task)
- void _starpu_task_init (void)
- void **_starpu_task_deinit** (void)
- void **_starpu_set_current_task** (struct starpu_task ∗task)
- int **_starpu_submit_job** (struct _starpu_job ∗j)
- int **_starpu_task_submit_nodeps** (struct starpu_task ∗task)
- void **_starpu_task_declare_deps_array** (struct starpu_task ∗task, unsigned ndeps, struct starpu_task ∗task_array[ ], int check)
- struct _starpu_job ∗ _starpu_get_job_associated_to_task_slow (struct starpu_task ∗task, struct _starpu_job ∗job)
- static struct _starpu_job ∗ **_starpu_get_job_associated_to_task** (struct starpu_task ∗task)
- int _starpu_task_submit_internally (struct starpu_task ∗task)
- int **_starpu_handle_needs_conversion_task** (starpu_data_handle_t handle, unsigned int node)
- int **_starpu_handle_needs_conversion_task_for_arch** (starpu_data_handle_t handle, enum starpu_↩ node_kind node_kind)
- void _starpu_task_prepare_for_continuation_ext (unsigned continuation_resubmit, void(∗continuation_↩ callback_on_sleep)(void ∗arg), void ∗continuation_callback_on_sleep_arg)
- void **_starpu_task_prepare_for_continuation** (void)
- void **_starpu_task_set_omp_cleanup_callback** (struct starpu_task ∗task, void(∗omp_cleanup_↩ callback)(void ∗arg), void ∗omp_cleanup_callback_arg)
- int **_starpu_task_uses_multiformat_handles** (struct starpu_task ∗task)
- int **_starpu_task_submit_conversion_task** (struct starpu_task ∗task, unsigned int workerid)
- void **_starpu_task_check_deprecated_fields** (struct starpu_task ∗task)
- void **_starpu_codelet_check_deprecated_fields** (struct starpu_codelet ∗cl)
- static starpu_cpu_func_t **_starpu_task_get_cpu_nth_implementation** (struct starpu_codelet ∗cl, unsigned nimpl)
- static starpu_cuda_func_t **_starpu_task_get_cuda_nth_implementation** (struct starpu_codelet ∗cl, unsigned nimpl)
- static starpu_opencl_func_t **_starpu_task_get_opencl_nth_implementation** (struct starpu_codelet ∗cl, unsigned nimpl)
- static starpu_mic_func_t **_starpu_task_get_mic_nth_implementation** (struct starpu_codelet ∗cl, unsigned nimpl)
- static starpu_mpi_ms_func_t **_starpu_task_get_mpi_ms_nth_implementation** (struct starpu_codelet ∗cl, unsigned nimpl)
- static const char ∗ **_starpu_task_get_cpu_name_nth_implementation** (struct starpu_codelet ∗cl, unsigned nimpl)
- void **_starpu_watchdog_init** (void)
- void **_starpu_watchdog_shutdown** (void)
- int **_starpu_task_wait_for_all_and_return_nb_waited_tasks** (void)
- int **_starpu_task_wait_for_all_in_ctx_and_return_nb_waited_tasks** (unsigned sched_ctx)

6.75.1 Function Documentation

6.75.1.1 _starpu_task_destroy()

```
void _starpu_task_destroy (
            struct starpu_task * task )
```
Internal version of starpu_task_destroy: don't check task->destroy flag

6.75.1.2 _starpu_task_test_termination()

```
int _starpu_task_test_termination (
            struct starpu_task * task )
```
Test for the termination of the task. Call starpu_task_destroy if required and the task is terminated.

6.75.1.3 _starpu_task_init()

```
void _starpu_task_init (
            void  )
```
A pthread key is used to store the task currently executed on the thread. _starpu_task_init initializes this pthread key and _starpu_set_current_task updates its current value.

6.75.1.4 _starpu_get_job_associated_to_task_slow()

```
struct _starpu_job* _starpu_get_job_associated_to_task_slow (
            struct starpu_task * task,
            struct _starpu_job * job )
```
Returns the job structure (which is the internal data structure associated to a task).

6.75.1.5 _starpu_task_submit_internally()

```
int _starpu_task_submit_internally (
            struct starpu_task * task )
```
Submits starpu internal tasks to the initial context

6.75.1.6 _starpu_task_prepare_for_continuation_ext()

```
void _starpu_task_prepare_for_continuation_ext (
            unsigned continuation_resubmit,
            void(*)(void *arg) continuation_callback_on_sleep,
            void * continuation_callback_on_sleep_arg )
```
Prepare the current task for accepting new dependencies before becoming a continuation.

6.76 task_bundle.h File Reference

```
#include <starpu_thread.h>
```

Data Structures

- struct _starpu_task_bundle_entry
- struct _starpu_task_bundle
- struct _starpu_handle_list

Functions

- void _starpu_task_bundle_destroy (starpu_task_bundle_t bundle)
- void _starpu_insertion_handle_sorted (struct _starpu_handle_list ∗∗listp, starpu_data_handle_t handle, enum starpu_data_access_mode mode)

6.76.1 Data Structure Documentation

6.76.1.1 struct _starpu_task_bundle_entry

struct _starpu_task_bundle_entry ================================ Purpose ======= Structure used to describe a linked list containing tasks in _starpu_task_bundle.
Fields ====== task Pointer to the task structure.
next Pointer to the next element in the linked list.

Data Fields

struct starpu_task ∗	task	
struct _starpu_task_bundle_entry ∗	next	

6.76.1.2   struct _starpu_task_bundle

struct _starpu_task_bundle =========================== Purpose ======= Structure describing a list of tasks that should be scheduled on the same worker whenever it's possible. It must be considered as a hint given to the scheduler as there is no guarantee that they will be executed on the same worker.
Fields ====== mutex Mutex protecting the structure.
list Array of tasks included in the bundle.
closed Used to know if the user is still willing to add/remove some tasks in the bundle. Especially useful for the runtime to know whether it is safe to destroy a bundle.

Data Fields

starpu_pthread_mutex_t	mutex	Mutex protecting the bundle
struct _starpu_task_bundle_entry ∗	list	
int	closed	

6.76.1.3   struct _starpu_handle_list

struct _starpu_handle_list =========================== Purpose ======= Structure describing a list of handles sorted by address to speed-up when looking for an element. The list cannot containes duplicate handles.
Fields ====== handle Pointer to the handle structure.
access_mode Total access mode over the whole bundle.
next Pointer to the next element in the linked list.

Data Fields

starpu_data_handle_t	handle	
enum starpu_data_access_mode	mode	
struct _starpu_handle_list ∗	next	

6.76.2   Function Documentation

6.76.2.1   _starpu_task_bundle_destroy()

```
void _starpu_task_bundle_destroy (
            starpu_task_bundle_t bundle )
```
_starpu_task_bundle_destroy =========================== Purpose ======= Destroy and deinitialize a bundle, memory previoulsy allocated is freed.
Arguments ========= bundle (input) Bundle to destroy.

6.76.2.2   _starpu_insertion_handle_sorted()

```
void _starpu_insertion_handle_sorted (
            struct _starpu_handle_list ** listp,
            starpu_data_handle_t handle,
            enum starpu_data_access_mode mode )
```

_starpu_insertion_handle_sorted ======================== Purpose ======= Insert an handle in a [_↩](#)
[starpu_handle_list](#), elements are sorted in increasing order, considering their physical address. As the list doesn't
accept duplicate elements, a handle with the same address as an handle contained in the list is not inserted, but its
mode access is merged with the one of the latter.

Arguments ========= listp (input, output) Pointer to the first element of the list. In the case of an empty list or an
inserted handle with small address, it should have changed when the call returns.

handle (input) Handle to insert in the list.

mode (input) Access mode of the handle.

6.77 thread.h File Reference

```
#include <common/utils.h>
```

Macros

- #define **starpu_pthread_spin_init**
- #define **starpu_pthread_spin_destroy**
- #define **starpu_pthread_spin_lock**
- #define **starpu_pthread_spin_trylock**
- #define **starpu_pthread_spin_unlock**

Functions

- static int **_starpu_pthread_spin_init** (starpu_pthread_spinlock_t *lock, int pshared STARPU_ATTRIBUT↩
  E_UNUSED)
- static int **_starpu_pthread_spin_destroy** (starpu_pthread_spinlock_t *lock STARPU_ATTRIBUTE_UNU↩
  SED)
- static int **_starpu_pthread_spin_lock** (starpu_pthread_spinlock_t *lock)
- static void **_starpu_pthread_spin_checklocked** (starpu_pthread_spinlock_t *lock STARPU_ATTRIBUT↩
  E_UNUSED)
- static int **_starpu_pthread_spin_trylock** (starpu_pthread_spinlock_t *lock)
- static int **_starpu_pthread_spin_unlock** (starpu_pthread_spinlock_t *lock)

6.78 timing.h File Reference

```
#include <stdint.h>
#include <common/config.h>
#include <starpu.h>
#include <starpu_util.h>
```

Functions

- void [_starpu_timing_init](#) (void)
- void **_starpu_clock_gettime** ([struct](#) timespec *ts)

6.78.1 Function Documentation

6.78.1.1 _starpu_timing_init()

```
void _starpu_timing_init (
            void )
```

_starpu_timing_init must be called prior to using any of these timing functions.

6.79 topology.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/list.h>
#include <common/fxt.h>
```

Macros

- #define **STARPU_NOWORKERID**
- #define **STARPU_ACTIVETHREAD**
- #define **STARPU_NONACTIVETHREAD**

Functions

- int _starpu_build_topology (struct _starpu_machine_config ∗config, int no_mp_config)
- void _starpu_destroy_machine_config (struct _starpu_machine_config ∗config)
- void _starpu_destroy_topology (struct _starpu_machine_config ∗config)
- unsigned _starpu_topology_get_nhwcpu (struct _starpu_machine_config ∗config)
- unsigned _starpu_topology_get_nhwpu (struct _starpu_machine_config ∗config)
- unsigned _starpu_topology_get_nnumanodes (struct _starpu_machine_config ∗config)
- unsigned _starpu_get_nhyperthreads ()
- void _starpu_topology_filter (hwloc_topology_t topology)
- int _starpu_bind_thread_on_cpu (int cpuid, int workerid, const char ∗name)
- void _starpu_bind_thread_on_cpus (struct _starpu_combined_worker ∗combined_worker)
- struct _starpu_worker ∗ **_starpu_get_worker_from_driver** (struct starpu_driver ∗d)
- int **starpu_memory_nodes_get_numa_count** (void)
- int **starpu_memory_nodes_numa_id_to_hwloclogid** (unsigned id)
- int _starpu_task_data_get_node_on_node (struct starpu_task ∗task, unsigned index, unsigned target_node)
- int **_starpu_task_data_get_node_on_worker** (struct starpu_task ∗task, unsigned index, unsigned worker)

6.79.1 Function Documentation

6.79.1.1 _starpu_build_topology()

```
int _starpu_build_topology (
            struct _starpu_machine_config * config,
            int no_mp_config )
```

Detect the number of memory nodes and where to bind the different workers.

6.79.1.2 _starpu_destroy_machine_config()

```
void _starpu_destroy_machine_config (
            struct _starpu_machine_config * config )
```

Should be called instead of _starpu_destroy_topology when _starpu_build_topology returns a non zero value.

6.79.1.3 _starpu_destroy_topology()

```
void _starpu_destroy_topology (
            struct _starpu_machine_config * config )
```

Destroy all resources used to store the topology of the machine.

6.79.1.4 _starpu_topology_get_nhwcpu()

```
unsigned _starpu_topology_get_nhwcpu (
            struct _starpu_machine_config * config )
```
returns the number of physical cpus

6.79.1.5 _starpu_topology_get_nhwpu()

```
unsigned _starpu_topology_get_nhwpu (
            struct _starpu_machine_config * config )
```
returns the number of logical cpus

6.79.1.6 _starpu_topology_get_nnumanodes()

```
unsigned _starpu_topology_get_nnumanodes (
            struct _starpu_machine_config * config )
```
returns the number of NUMA nodes

6.79.1.7 _starpu_get_nhyperthreads()

```
unsigned _starpu_get_nhyperthreads ( )
```
returns the number of hyperthreads per core

6.79.1.8 _starpu_topology_filter()

```
void _starpu_topology_filter (
            hwloc_topology_t topology )
```
Small convenient function to filter hwloc topology depending on HWLOC API version

6.79.1.9 _starpu_bind_thread_on_cpu()

```
int _starpu_bind_thread_on_cpu (
            int cpuid,
            int workerid,
            const char * name )
```
Bind the current thread on the CPU logically identified by "cpuid". The logical ordering of the processors is either that of hwloc (if available), or the ordering exposed by the OS.

6.79.1.10 _starpu_bind_thread_on_cpus()

```
void _starpu_bind_thread_on_cpus (
            struct _starpu_combined_worker * combined_worker )
```
Bind the current thread on the set of CPUs for the given combined worker.

6.79.1.11 _starpu_task_data_get_node_on_node()

```
int _starpu_task_data_get_node_on_node (
            struct starpu_task * task,
            unsigned index,
            unsigned target_node )
```
Get the memory node for data number i when task is to be executed on memory node target_node

6.80 utils.h File Reference

```
#include <common/config.h>
#include <starpu.h>
#include <sys/stat.h>
#include <string.h>
#include <stdlib.h>
```

```
#include <math.h>
```

Macros

- #define **DO_CREQ_v_WW**(_creqF, _ty1F, _arg1F, _ty2F, _arg2F)
- #define **DO_CREQ_v_W**(_creqF, _ty1F, _arg1F)
- #define **ANNOTATE_HAPPENS_BEFORE**(obj)
- #define **ANNOTATE_HAPPENS_BEFORE_FORGET_ALL**(obj)
- #define **ANNOTATE_HAPPENS_AFTER**(obj)
- #define **VALGRIND_HG_DISABLE_CHECKING**(start, len)
- #define **VALGRIND_HG_ENABLE_CHECKING**(start, len)
- #define **VALGRIND_STACK_REGISTER**(stackbottom, stacktop)
- #define **VALGRIND_STACK_DEREGISTER**(id)
- #define **RUNNING_ON_VALGRIND**
- #define **STARPU_RUNNING_ON_VALGRIND**
- #define **STARPU_HG_DISABLE_CHECKING**(variable)
- #define **STARPU_HG_ENABLE_CHECKING**(variable)
- #define **STARPU_DEBUG_PREFIX**
- #define **_STARPU_UYIELD**()
- #define **STARPU_VALGRIND_YIELD**()
- #define **STARPU_UYIELD**()
- #define **_STARPU_DEBUG**(fmt, ...)
- #define **_STARPU_DEBUG_NO_HEADER**(fmt, ...)
- #define **_STARPU_EXTRA_DEBUG**(fmt, ...)
- #define **_STARPU_LOG_IN**()
- #define **_STARPU_LOG_OUT**()
- #define **_STARPU_LOG_OUT_TAG**(outtag)
- #define **_STARPU_MSG**(fmt, ...)
- #define **_STARPU_DISP**(fmt, ...)
- #define **_STARPU_ERROR**(fmt, ...)
- #define **_STARPU_DECLTYPE**(x)
- #define **_STARPU_MALLOC**(ptr, size)
- #define **_STARPU_CALLOC**(ptr, nmemb, size)
- #define **_STARPU_REALLOC**(ptr, size)
- #define **_STARPU_IS_ZERO**(a)

Functions

- char ∗ **_starpu_mkdtemp_internal** (char ∗tmpl)
- char ∗ **_starpu_mkdtemp** (char ∗tmpl)
- int **_starpu_mkpath** (const char ∗s, mode_t mode)
- void **_starpu_mkpath_and_check** (const char ∗s, mode_t mode)
- char ∗ **_starpu_mktemp** (const char ∗directory, int flags, int ∗fd)
- char ∗ _starpu_mktemp_many (const char ∗directory, int depth, int flags, int ∗fd)
- void **_starpu_rmtemp_many** (char ∗path, int depth)
- void **_starpu_rmdir_many** (char ∗path, int depth)
- int **_starpu_fftruncate** (FILE ∗file, size_t length)
- int **_starpu_ftruncate** (int fd, size_t length)
- int **_starpu_frdlock** (FILE ∗file)
- int **_starpu_frdunlock** (FILE ∗file)
- int **_starpu_fwrlock** (FILE ∗file)
- int **_starpu_fwrunlock** (FILE ∗file)
- char ∗ **_starpu_get_home_path** (void)
- void **_starpu_gethostname** (char ∗hostname, size_t size)

- void _starpu_drop_comments (FILE ∗f)
- const char ∗ _starpu_job_get_model_name (struct _starpu_job ∗j)
- const char ∗ _starpu_job_get_task_name (struct _starpu_job ∗j)
- const char ∗ _starpu_codelet_get_model_name (struct starpu_codelet ∗cl)
- int **_starpu_check_mutex_deadlock** (starpu_pthread_mutex_t ∗mutex)
- void **_starpu_util_init** (void)

6.80.1 Function Documentation

6.80.1.1 _starpu_mktemp_many()

```
char* _starpu_mktemp_many (
            const char * directory,
            int depth,
            int flags,
            int * fd )
```
This version creates a hierarchy of n temporary directories, useful when creating a lot of temporary files to be stored in the same place

6.80.1.2 _starpu_drop_comments()

```
void _starpu_drop_comments (
            FILE * f )
```
If FILE is currently on a comment line, eat it.

6.80.1.3 _starpu_job_get_model_name()

```
const char* _starpu_job_get_model_name (
            struct _starpu_job * j )
```
Returns the symbol associated to that job if any.

6.80.1.4 _starpu_job_get_task_name()

```
const char* _starpu_job_get_task_name (
            struct _starpu_job * j )
```
Returns the name associated to that job if any.

6.80.1.5 _starpu_codelet_get_model_name()

```
const char* _starpu_codelet_get_model_name (
            struct starpu_codelet * cl )
```
Returns the symbol associated to that job if any.

6.81 uthash.h File Reference

```
#include <string.h>
#include <stddef.h>
#include <inttypes.h>
```

Data Structures

- struct UT_hash_bucket
- struct UT_hash_table
- struct UT_hash_handle

Macros

- #define **DECLTYPE**(x)
- #define **DECLTYPE_ASSIGN**(dst, src)
- #define **UTHASH_VERSION**
- #define **uthash_fatal**(msg)
- #define **uthash_malloc**(sz)
- #define **uthash_free**(ptr, sz)
- #define **uthash_noexpand_fyi**(tbl)
- #define **uthash_expand_fyi**(tbl)
- #define **HASH_INITIAL_NUM_BUCKETS**
- #define **HASH_INITIAL_NUM_BUCKETS_LOG2**
- #define **HASH_BKT_CAPACITY_THRESH**
- #define **ELMT_FROM_HH**(tbl, hhp)
- #define **HASH_FIND**(hh, head, keyptr, keylen, out)
- #define **HASH_BLOOM_MAKE**(tbl)
- #define **HASH_BLOOM_FREE**(tbl)
- #define **HASH_BLOOM_ADD**(tbl, hashv)
- #define **HASH_BLOOM_TEST**(tbl, hashv)
- #define **HASH_MAKE_TABLE**(hh, head)
- #define **HASH_ADD**(hh, head, fieldname, keylen_in, add)
- #define **HASH_CHECK_KEY**(hh, head, keyptr, keylen, out)
- #define **HASH_ADD_KEYPTR**(hh, head, keyptr, keylen_in, add)
- #define **HASH_TO_BKT**(hashv, num_bkts, bkt)
- #define **HASH_DELETE**(hh, head, delptr)
- #define **HASH_FIND_STR**(head, findstr, out)
- #define **HASH_ADD_STR**(head, strfield, add)
- #define **HASH_FIND_INT**(head, findint, out)
- #define **HASH_ADD_INT**(head, intfield, add)
- #define **HASH_FIND_PTR**(head, findptr, out)
- #define **HASH_ADD_PTR**(head, ptrfield, add)
- #define **HASH_DEL**(head, delptr)
- #define **HASH_FSCK**(hh, head)
- #define **HASH_EMIT_KEY**(hh, head, keyptr, fieldlen)
- #define **HASH_FCN**
- #define **HASH_BER**(key, keylen, num_bkts, hashv, bkt)
- #define **HASH_SAX**(key, keylen, num_bkts, hashv, bkt)
- #define **HASH_FNV**(key, keylen, num_bkts, hashv, bkt)
- #define **HASH_OAT**(key, keylen, num_bkts, hashv, bkt)
- #define **HASH_JEN_MIX**(a, b, c)
- #define **HASH_JEN**(key, keylen, num_bkts, hashv, bkt)
- #define **get16bits**(d)
- #define **HASH_SFH**(key, keylen, num_bkts, hashv, bkt)
- #define **HASH_KEYCMP**(a, b, len)
- #define **HASH_FIND_IN_BKT**(tbl, hh, head, keyptr, keylen_in, out)
- #define **HASH_ADD_TO_BKT**(head, addhh)
- #define **HASH_DEL_IN_BKT**(hh, head, hh_del)
- #define **HASH_EXPAND_BUCKETS**(tbl)
- #define **HASH_SORT**(head, cmpfcn)
- #define **HASH_SRT**(hh, head, cmpfcn)
- #define **HASH_SELECT**(hh_dst, dst, hh_src, src, cond)
- #define **HASH_CLEAR**(hh, head)
- #define **HASH_ITER**(hh, head, el, tmp)
- #define **HASH_COUNT**(head)
- #define **HASH_CNT**(hh, head)
- #define **HASH_SIGNATURE**
- #define **HASH_BLOOM_SIGNATURE**

Typedefs

- typedef struct UT_hash_bucket **UT_hash_bucket**
- typedef struct UT_hash_table **UT_hash_table**
- typedef struct UT_hash_handle **UT_hash_handle**

6.81.1 Data Structure Documentation

6.81.1.1 struct UT_hash_bucket

Data Fields

struct UT_hash_handle ∗	hh_head	
unsigned	count	
unsigned	expand_mult	

6.81.1.2 struct UT_hash_table

Data Fields

UT_hash_bucket ∗	buckets	
unsigned	num_buckets	
unsigned	log2_num_buckets	
unsigned	num_items	
struct UT_hash_handle ∗	tail	
ptrdiff_t	hho	
unsigned	ideal_chain_maxlen	
unsigned	nonideal_items	
unsigned	ineff_expands	
unsigned	noexpand	
uint32_t	signature	

6.81.1.3 struct UT_hash_handle

Data Fields

struct UT_hash_table ∗	tbl	
void ∗	prev	
void ∗	next	
struct UT_hash_handle ∗	hh_prev	
struct UT_hash_handle ∗	hh_next	
void ∗	key	
unsigned	keylen	
unsigned	hashv	

6.82 write_back.h File Reference

```
#include <starpu.h>
#include <datawizard/coherency.h>
```

Functions

- void _starpu_write_through_data (starpu_data_handle_t handle, unsigned requesting_node, uint32_t write↩
  _through_mask)

6.82.1 Function Documentation

6.82.1.1 _starpu_write_through_data()

```
void _starpu_write_through_data (
            starpu_data_handle_t handle,
            unsigned requesting_node,
            uint32_t write_through_mask )
```

If a write-through mask is associated to that data handle, this propagates the the current value of the data onto the different memory nodes in the write_through_mask.

Chapter 7

StarPU MPI File Documentation

7.1 starpu_mpi_cache.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
```

Functions

- void **_starpu_mpi_cache_init** (MPI_Comm comm)
- void **_starpu_mpi_cache_shutdown** ()
- void **_starpu_mpi_cache_data_init** (starpu_data_handle_t data_handle)
- void **_starpu_mpi_cache_data_clear** (starpu_data_handle_t data_handle)
- void **_starpu_mpi_cache_flush** (starpu_data_handle_t data_handle)

Variables

- int **_starpu_cache_enabled**

7.2 starpu_mpi_driver.h File Reference

```
#include <starpu.h>
```

Functions

- void **_starpu_mpi_driver_init** (struct starpu_conf ∗conf)
- void **_starpu_mpi_driver_shutdown** ()

7.3 starpu_mpi_init.h File Reference

```
#include <starpu.h>
#include <starpu_mpi.h>
```

Functions

- void **_starpu_mpi_do_initialize** (struct _starpu_mpi_argc_argv ∗argc_argv)

7.4 starpu_mpi_nmad_backend.h File Reference

```
#include <common/config.h>
#include <nm_sendrecv_interface.h>
#include <nm_session_interface.h>
#include <nm_mpi_nmad.h>
```

Data Structures

- struct _starpu_mpi_req_backend

7.4.1 Data Structure Documentation

7.4.1.1 struct _starpu_mpi_req_backend

Data Fields

MPI_Request	data_request	
starpu_pthread_mutex_t	req_mutex	
starpu_pthread_cond_t	req_cond	
starpu_pthread_cond_t	posted_cond	
struct _starpu_mpi_req ∗	other_request	In the case of a Wait/Test request, we are going to post a request to test the completion of another request
MPI_Request	size_req	
struct _starpu_mpi_envelope ∗	envelope	
unsigned	is_internal_req:1	
unsigned	to_destroy:1	
struct _starpu_mpi_req ∗	internal_req	
struct _starpu_mpi_early_data_handle ∗	early_data_handle	
UT_hash_handle	hh	
nm_gate_t	gate	
nm_session_t	session	
nm_sr_request_t	data_request	
int	waited	
piom_cond_t	req_cond	
nm_sr_request_t	size_req	

7.5 starpu_mpi_stats.h File Reference

```
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
```

Functions

- void **_starpu_mpi_comm_amounts_init** (MPI_Comm comm)
- void **_starpu_mpi_comm_amounts_shutdown** ()
- void **_starpu_mpi_comm_amounts_inc** (MPI_Comm comm, unsigned dst, MPI_Datatype datatype, int count)
- void **_starpu_mpi_comm_amounts_display** (FILE ∗stream, int node)

7.6 starpu_mpi_cache_stats.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
```

Macros

- #define **_starpu_mpi_cache_stats_inc**(dst, data_handle)
- #define **_starpu_mpi_cache_stats_dec**(dst, data_handle)

Functions

- void **_starpu_mpi_cache_stats_init** ()
- void **_starpu_mpi_cache_stats_shutdown** ()
- void **_starpu_mpi_cache_stats_update** (unsigned dst, starpu_data_handle_t data_handle, int count)

7.7 starpu_mpi_early_data.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
#include <common/config.h>
#include <common/list.h>
#include <common/uthash.h>
#include <starpu_mpi_private.h>
```

Data Structures

- struct _starpu_mpi_early_data_handle
- struct _starpu_mpi_early_data_handle_tag_hashlist

Functions

- void **_starpu_mpi_early_data_init** (void)
- void **_starpu_mpi_early_data_check_termination** (void)
- void **_starpu_mpi_early_data_shutdown** (void)
- struct _starpu_mpi_early_data_handle ∗ **_starpu_mpi_early_data_create** (struct _starpu_mpi_envelope ∗envelope, int source, MPI_Comm comm) STARPU_ATTRIBUTE_MALLOC
- struct _starpu_mpi_early_data_handle ∗ **_starpu_mpi_early_data_find** (struct _starpu_mpi_node_tag ∗node_tag)
- void **_starpu_mpi_early_data_add** (struct _starpu_mpi_early_data_handle ∗early_data_handle)
- struct _starpu_mpi_early_data_handle_tag_hashlist ∗ **_starpu_mpi_early_data_extract** (struct _starpu_↩ mpi_node_tag ∗node_tag)

7.7.1 Data Structure Documentation

7.7.1.1 struct _starpu_mpi_early_data_handle

Data Fields

starpu_data_handle_t	handle	
struct _starpu_mpi_req ∗	req	
void ∗	buffer	

Data Fields

size_t	size	
struct _starpu_mpi_node_tag	node_tag	
starpu_pthread_mutex_t	req_mutex	
starpu_pthread_cond_t	req_cond	

7.7.1.2 struct _starpu_mpi_early_data_handle_tag_hashlist

Data Fields

struct _starpu_mpi_early_data_handle_list	list	
UT_hash_handle	hh	
starpu_mpi_tag_t	data_tag	

7.8 starpu_mpi_sync_data.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
#include <common/config.h>
#include <common/list.h>
```

Functions

- void **_starpu_mpi_sync_data_init** (void)
- void **_starpu_mpi_sync_data_check_termination** (void)
- void **_starpu_mpi_sync_data_shutdown** (void)
- struct _starpu_mpi_req ∗ **_starpu_mpi_sync_data_find** (starpu_mpi_tag_t data_tag, int source, MPI_←
  Comm comm)
- void **_starpu_mpi_sync_data_add** (struct _starpu_mpi_req ∗req)
- int **_starpu_mpi_sync_data_count** ()

7.9 starpu_mpi_comm.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
#include <mpi/starpu_mpi_mpi_backend.h>
```

Functions

- void **_starpu_mpi_comm_init** (MPI_Comm comm)
- void **_starpu_mpi_comm_shutdown** ()
- void **_starpu_mpi_comm_register** (MPI_Comm comm)
- void **_starpu_mpi_comm_post_recv** ()
- int **_starpu_mpi_comm_test_recv** (MPI_Status ∗status, struct _starpu_mpi_envelope ∗∗envelope, MPI_←
  Comm ∗comm)
- void **_starpu_mpi_comm_cancel_recv** ()

7.10 starpu_mpi_early_request.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
#include <common/config.h>
#include <common/list.h>
```

Data Structures

- struct _starpu_mpi_early_request_tag_hashlist

Functions

- void **_starpu_mpi_early_request_init** (void)
- void **_starpu_mpi_early_request_shutdown** (void)
- int **_starpu_mpi_early_request_count** (void)
- void **_starpu_mpi_early_request_check_termination** (void)
- void **_starpu_mpi_early_request_enqueue** (struct _starpu_mpi_req ∗req)
- struct _starpu_mpi_req ∗ **_starpu_mpi_early_request_dequeue** (starpu_mpi_tag_t data_tag, int source, MPI_Comm comm)
- struct _starpu_mpi_early_request_tag_hashlist ∗ **_starpu_mpi_early_request_extract** (starpu_mpi_tag_t data_tag, int source, MPI_Comm comm)

7.10.1 Data Structure Documentation

7.10.1.1 struct _starpu_mpi_early_request_tag_hashlist

Data Fields

struct _starpu_mpi_req_list	list	
UT_hash_handle	hh	
starpu_mpi_tag_t	data_tag	

7.11 starpu_mpi_mpi_backend.h File Reference

```
#include <common/config.h>
#include <common/uthash.h>
```

Data Structures

- struct _starpu_mpi_envelope
- struct _starpu_mpi_req_backend

Macros

- #define **_STARPU_MPI_TAG_ENVELOPE**
- #define **_STARPU_MPI_TAG_DATA**
- #define **_STARPU_MPI_TAG_SYNC_DATA**

Enumerations

- enum **_starpu_envelope_mode** { **_STARPU_MPI_ENVELOPE_DATA**, **_STARPU_MPI_ENVELOPE_S**↩ **YNC_READY** }

Variables

- int **_starpu_mpi_tag**

7.11.1 Data Structure Documentation

7.11.1.1 struct _starpu_mpi_envelope

Data Fields

enum _starpu_envelope_mode	mode	
starpu_ssize_t	size	
starpu_mpi_tag_t	data_tag	
unsigned	sync	

7.11.1.2 struct _starpu_mpi_req_backend

Data Fields

MPI_Request	data_request	
starpu_pthread_mutex_t	req_mutex	
starpu_pthread_cond_t	req_cond	
starpu_pthread_cond_t	posted_cond	
struct _starpu_mpi_req *	other_request	In the case of a Wait/Test request, we are going to post a request to test the completion of another request
MPI_Request	size_req	
struct _starpu_mpi_envelope *	envelope	
unsigned	is_internal_req:1	
unsigned	to_destroy:1	
struct _starpu_mpi_req *	internal_req	
struct _starpu_mpi_early_data_handle *	early_data_handle	
UT_hash_handle	hh	
nm_gate_t	gate	
nm_session_t	session	
nm_sr_request_t	data_request	
int	waited	
piom_cond_t	req_cond	
nm_sr_request_t	size_req	

7.12 starpu_mpi_private.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/uthash.h>
#include <starpu_mpi.h>
#include <starpu_mpi_fxt.h>
#include <common/list.h>
#include <common/prio_list.h>
#include <common/starpu_spinlock.h>
#include <core/simgrid.h>
```

Data Structures

- struct _starpu_simgrid_mpi_req
- struct _starpu_mpi_node
- struct _starpu_mpi_node_tag
- struct _starpu_mpi_coop_sends
- struct _starpu_mpi_data
- struct _starpu_mpi_req
- struct _starpu_mpi_argc_argv
- struct _starpu_mpi_backend

Macros

- #define **STARPU_MPI_ASSERT_MSG**(x, msg, ...)
- #define **_STARPU_MPI_MALLOC**(ptr, size)
- #define **_STARPU_MPI_CALLOC**(ptr, nmemb, size)
- #define **_STARPU_MPI_REALLOC**(ptr, size)
- #define **_STARPU_MPI_COMM_DEBUG**(ptr, count, datatype, node, tag, utag, comm, way)
- #define **_STARPU_MPI_COMM_TO_DEBUG**(ptr, count, datatype, dest, tag, utag, comm)
- #define **_STARPU_MPI_COMM_FROM_DEBUG**(ptr, count, datatype, source, tag, utag, comm)
- #define **_STARPU_MPI_DEBUG**(level, fmt, ...)
- #define **_STARPU_MPI_DISP**(fmt, ...)
- #define **_STARPU_MPI_MSG**(fmt, ...)
- #define **_STARPU_MPI_LOG_IN**()
- #define **_STARPU_MPI_LOG_OUT**()

Enumerations

- enum **_starpu_mpi_request_type** {
  SEND_REQ, **RECV_REQ**, **WAIT_REQ**, **TEST_REQ**,
  BARRIER_REQ, **PROBE_REQ**, **UNKNOWN_REQ** }

Functions

- int **_starpu_mpi_simgrid_mpi_test** (unsigned ∗done, int ∗flag)
- void **_starpu_mpi_simgrid_wait_req** (MPI_Request ∗request, MPI_Status ∗status, starpu_pthread_↩ queue_t ∗queue, unsigned ∗done)
- char ∗ **_starpu_mpi_get_mpi_error_code** (int code)
- void **_starpu_mpi_env_init** (void)
- struct _starpu_mpi_data ∗ **_starpu_mpi_data_get** (starpu_data_handle_t data_handle)
- void **_starpu_mpi_submit_ready_request** (void ∗arg)
- void **_starpu_mpi_release_req_data** (struct _starpu_mpi_req ∗req)
- void **_starpu_mpi_coop_sends_build_tree** (struct _starpu_mpi_coop_sends ∗coop_sends)
- void **_starpu_mpi_coop_send** (starpu_data_handle_t data_handle, struct _starpu_mpi_req ∗req, enum starpu_data_access_mode mode, int sequential_consistency)
- void **_starpu_mpi_submit_coop_sends** (struct _starpu_mpi_coop_sends ∗coop_sends, int submit_control, int submit_data)
- void **_starpu_mpi_submit_ready_request_inc** (struct _starpu_mpi_req ∗req)
- void **_starpu_mpi_request_init** (struct _starpu_mpi_req ∗∗req)
- struct _starpu_mpi_req ∗ **_starpu_mpi_request_fill** (starpu_data_handle_t data_handle, int sr-cdst, starpu_mpi_tag_t data_tag, MPI_Comm comm, unsigned detached, unsigned sync, int prio, void(∗callback)(void ∗), void ∗arg, enum _starpu_mpi_request_type request_type, void(∗func)(struct _↩ starpu_mpi_req ∗), int sequential_consistency, int is_internal_req, starpu_ssize_t count)
- void **_starpu_mpi_request_destroy** (struct _starpu_mpi_req ∗req)
- void **_starpu_mpi_isend_size_func** (struct _starpu_mpi_req ∗req)
- void **_starpu_mpi_irecv_size_func** (struct _starpu_mpi_req ∗req)

- int **_starpu_mpi_wait** (starpu_mpi_req ∗public_req, MPI_Status ∗status)
- int **_starpu_mpi_test** (starpu_mpi_req ∗public_req, int ∗flag, MPI_Status ∗status)
- int **_starpu_mpi_barrier** (MPI_Comm comm)
- void **_starpu_mpi_progress_shutdown** (void ∗∗value)
- int **_starpu_mpi_progress_init** (struct _starpu_mpi_argc_argv ∗argc_argv)
- void **_starpu_mpi_wait_for_initialization** ()
- void **_starpu_mpi_data_flush** (starpu_data_handle_t data_handle)

Variables

- starpu_pthread_wait_t **_starpu_mpi_thread_wait**
- starpu_pthread_queue_t **_starpu_mpi_thread_dontsleep**
- int **_starpu_debug_rank**
- int **_starpu_mpi_comm_debug**
- int **_starpu_mpi_fake_world_size**
- int **_starpu_mpi_fake_world_rank**
- int **_starpu_mpi_use_prio**
- int **_starpu_mpi_nobind**
- int **_starpu_mpi_thread_cpuid**
- int **_starpu_mpi_use_coop_sends**
- PRIO_struct **_starpu_mpi_req**
- struct _starpu_mpi_backend **_mpi_backend**

7.12.1 Data Structure Documentation

7.12.1.1 struct _starpu_simgrid_mpi_req

Data Fields

MPI_Request ∗	request	
MPI_Status ∗	status	
starpu_pthread_queue_t ∗	queue	
unsigned ∗	done	

7.12.1.2 struct _starpu_mpi_node

Data Fields

MPI_Comm	comm	
int	rank	

7.12.1.3 struct _starpu_mpi_node_tag

Data Fields

struct _starpu_mpi_node	node	
starpu_mpi_tag_t	data_tag	

7.12.1.4 struct _starpu_mpi_coop_sends

Data Fields

struct	

Data Fields

_starpu_mpi_req_multilist_coop_sends	reqs	
struct _starpu_mpi_data *	mpi_data	
struct _starpu_spinlock	lock	
struct _starpu_mpi_req **	reqs_array	
unsigned	n	
unsigned	redirects_sent	

7.12.1.5 struct _starpu_mpi_data

Data Fields

int	magic	
struct _starpu_mpi_node_tag	node_tag	
char *	cache_sent	
int	cache_received	
struct _starpu_spinlock	coop_lock	
struct _starpu_mpi_coop_sends *	coop_sends	

7.12.1.6 struct _starpu_mpi_argc_argv

Data Fields

int	initialize_mpi	
int *	argc	
char ***	argv	
MPI_Comm	comm	
int	fargc	Fortran argc
char **	fargv	Fortran argv
int	rank	
int	world_size	

7.13 starpu_mpi_tag.h File Reference

```
#include <starpu.h>
#include <stdlib.h>
#include <mpi.h>
```

Functions

- void **_starpu_mpi_tag_init** (void)
- void **_starpu_mpi_tag_shutdown** (void)
- void **_starpu_mpi_tag_data_register** (starpu_data_handle_t handle, starpu_mpi_tag_t data_tag)
- int **_starpu_mpi_tag_data_release** (starpu_data_handle_t handle)
- starpu_data_handle_t **_starpu_mpi_tag_get_data_handle_from_tag** (starpu_mpi_tag_t data_tag)

7.14 starpu_mpi_datatype.h File Reference

```
#include <starpu_mpi.h>
#include <starpu_mpi_private.h>
```

Functions

- void **_starpu_mpi_datatype_init** (void)
- void **_starpu_mpi_datatype_shutdown** (void)
- void **_starpu_mpi_datatype_allocate** (starpu_data_handle_t data_handle, struct _starpu_mpi_req ∗req)
- void **_starpu_mpi_datatype_free** (starpu_data_handle_t data_handle, MPI_Datatype ∗datatype)
- MPI_Datatype **_starpu_mpi_datatype_get_user_defined_datatype** (starpu_data_handle_t data_handle)

7.15 starpu_mpi_fxt.h File Reference

```
#include <starpu.h>
#include <common/config.h>
#include <common/fxt.h>
```

Macros

- #define **_STARPU_MPI_FUT_START**
- #define **_STARPU_MPI_FUT_STOP**
- #define **_STARPU_MPI_FUT_BARRIER**
- #define **_STARPU_MPI_FUT_ISEND_SUBMIT_BEGIN**
- #define **_STARPU_MPI_FUT_ISEND_SUBMIT_END**
- #define **_STARPU_MPI_FUT_IRECV_SUBMIT_BEGIN**
- #define **_STARPU_MPI_FUT_IRECV_SUBMIT_END**
- #define **_STARPU_MPI_FUT_ISEND_COMPLETE_BEGIN**
- #define **_STARPU_MPI_FUT_ISEND_COMPLETE_END**
- #define **_STARPU_MPI_FUT_DATA_SET_RANK**
- #define **_STARPU_MPI_FUT_IRECV_TERMINATED**
- #define **_STARPU_MPI_FUT_ISEND_TERMINATED**
- #define **_STARPU_MPI_FUT_TESTING_DETACHED_BEGIN**
- #define **_STARPU_MPI_FUT_TESTING_DETACHED_END**
- #define **_STARPU_MPI_FUT_TEST_BEGIN**
- #define **_STARPU_MPI_FUT_TEST_END**
- #define **_STARPU_MPI_FUT_IRECV_COMPLETE_BEGIN**
- #define **_STARPU_MPI_FUT_IRECV_COMPLETE_END**
- #define **_STARPU_MPI_FUT_SLEEP_BEGIN**
- #define **_STARPU_MPI_FUT_SLEEP_END**
- #define **_STARPU_MPI_FUT_DTESTING_BEGIN**
- #define **_STARPU_MPI_FUT_DTESTING_END**
- #define **_STARPU_MPI_FUT_UTESTING_BEGIN**
- #define **_STARPU_MPI_FUT_UTESTING_END**
- #define **_STARPU_MPI_FUT_UWAIT_BEGIN**
- #define **_STARPU_MPI_FUT_UWAIT_END**
- #define **_STARPU_MPI_FUT_POLLING_BEGIN**
- #define **_STARPU_MPI_FUT_POLLING_END**
- #define **_STARPU_MPI_FUT_DRIVER_RUN_BEGIN**
- #define **_STARPU_MPI_FUT_DRIVER_RUN_END**
- #define **_STARPU_MPI_FUT_DATA_SET_TAG**
- #define **_STARPU_MPI_TRACE_START**(a, b)
- #define **_STARPU_MPI_TRACE_STOP**(a, b)
- #define **_STARPU_MPI_TRACE_BARRIER**(a, b, c)
- #define **_STARPU_MPI_TRACE_ISEND_SUBMIT_BEGIN**(a, b, c)
- #define **_STARPU_MPI_TRACE_ISEND_SUBMIT_END**(a, b, c, d)

- #define **_STARPU_MPI_TRACE_IRECV_SUBMIT_BEGIN**(a, b)
- #define **_STARPU_MPI_TRACE_IRECV_SUBMIT_END**(a, b)
- #define **_STARPU_MPI_TRACE_ISEND_COMPLETE_BEGIN**(a, b, c)
- #define **_STARPU_MPI_TRACE_COMPLETE_BEGIN**(a, b, c)
- #define **_STARPU_MPI_TRACE_COMPLETE_END**(a, b, c)
- #define **_STARPU_MPI_TRACE_TERMINATED**(a, b, c)
- #define **_STARPU_MPI_TRACE_ISEND_COMPLETE_END**(a, b, c)
- #define **_STARPU_MPI_TRACE_IRECV_COMPLETE_BEGIN**(a, b)
- #define **_STARPU_MPI_TRACE_IRECV_COMPLETE_END**(a, b)
- #define **_STARPU_MPI_TRACE_SLEEP_BEGIN**()
- #define **_STARPU_MPI_TRACE_SLEEP_END**()
- #define **_STARPU_MPI_TRACE_DTESTING_BEGIN**()
- #define **_STARPU_MPI_TRACE_DTESTING_END**()
- #define **_STARPU_MPI_TRACE_UTESTING_BEGIN**(a, b)
- #define **_STARPU_MPI_TRACE_UTESTING_END**(a, b)
- #define **_STARPU_MPI_TRACE_UWAIT_BEGIN**(a, b)
- #define **_STARPU_MPI_TRACE_UWAIT_END**(a, b)
- #define **_STARPU_MPI_TRACE_DATA_SET_RANK**(a, b)
- #define **_STARPU_MPI_TRACE_DATA_SET_TAG**(a, b)
- #define **_STARPU_MPI_TRACE_TESTING_DETACHED_BEGIN**()
- #define **_STARPU_MPI_TRACE_TESTING_DETACHED_END**()
- #define **_STARPU_MPI_TRACE_TEST_BEGIN**(peer, data_tag)
- #define **_STARPU_MPI_TRACE_TEST_END**(peer, data_tag)
- #define **_STARPU_MPI_TRACE_POLLING_BEGIN**()
- #define **_STARPU_MPI_TRACE_POLLING_END**()
- #define **_STARPU_MPI_TRACE_DRIVER_RUN_BEGIN**()
- #define **_STARPU_MPI_TRACE_DRIVER_RUN_END**()

7.16 starpu_mpi_select_node.h File Reference

```
#include <mpi.h>
```

Macros

- #define **_STARPU_MPI_NODE_SELECTION_MAX_POLICY**

Functions

- void **_starpu_mpi_select_node_init** ()
- int **_starpu_mpi_select_node** (int me, int nb_nodes, struct starpu_data_descr ∗descr, int nb_data, int policy)

7.17 starpu_mpi_task_insert.h File Reference

Functions

- int **_starpu_mpi_find_executee_node** (starpu_data_handle_t data, enum starpu_data_access_mode mode, int me, int ∗do_execute, int ∗inconsistent_execute, int ∗xrank)
- void **_starpu_mpi_exchange_data_before_execution** (starpu_data_handle_t data, enum starpu_data_↩ access_mode mode, int me, int xrank, int do_execute, int prio, MPI_Comm comm)
- int **_starpu_mpi_task_postbuild_v** (MPI_Comm comm, int xrank, int do_execute, struct starpu_data_descr ∗descrs, int nb_data, int prio)

7.18 load_balancer_policy.h File Reference

```
#include <starpu_mpi_lb.h>
```

Data Structures

- struct load_balancer_policy

Variables

- struct load_balancer_policy **load_heat_propagation_policy**

7.19 load_data_interface.h File Reference

```
#include <starpu.h>
```

Data Structures

- struct load_data_interface

Macros

- #define **LOAD_DATA_GET_NSUBMITTED_TASKS**(interface)
- #define **LOAD_DATA_GET_SLEEP_THRESHOLD**(interface)
- #define **LOAD_DATA_GET_WAKEUP_THRESHOLD**(interface)

Functions

- void **load_data_data_register** (starpu_data_handle_t ∗handle, unsigned home_node, int sleep_task_↩
  threshold, double wakeup_ratio)
- int **load_data_get_sleep_threshold** (starpu_data_handle_t handle)
- int **load_data_get_wakeup_threshold** (starpu_data_handle_t handle)
- int **load_data_get_current_phase** (starpu_data_handle_t handle)
- int **load_data_get_nsubmitted_tasks** (starpu_data_handle_t handle)
- int **load_data_get_nfinished_tasks** (starpu_data_handle_t handle)
- int **load_data_inc_nsubmitted_tasks** (starpu_data_handle_t handle)
- int **load_data_inc_nfinished_tasks** (starpu_data_handle_t handle)
- int **load_data_next_phase** (starpu_data_handle_t handle)
- int **load_data_update_elapsed_time** (starpu_data_handle_t handle)
- double **load_data_get_elapsed_time** (starpu_data_handle_t handle)
- int **load_data_update_wakeup_cond** (starpu_data_handle_t handle)
- int **load_data_wakeup_cond** (starpu_data_handle_t handle)

7.19.1 Data Structure Documentation

7.19.1.1 struct load_data_interface

interface for load_data

Data Fields

double	start	Starting time of the execution
double	elapsed_time	Elapsed time until the start time and the time when event "launch a load balancing phase" is triggered

Data Fields

int	phase	Current submission phase, i.e how many balanced steps have already happened so far.
int	nsubmitted_tasks	Number of currently submitted tasks
int	nfinished_tasks	Number of currently finished tasks
int	sleep_task_threshold	Task threshold to sleep the submission thread
int	wakeup_task_threshold	Task threshold to wake-up the submission thread
double	wakeup_ratio	Ratio of submitted tasks to wait for completion before waking up the submission thread

7.20 data_movements_interface.h File Reference

```
#include <starpu.h>
```

Data Structures

- struct data_movements_interface

Macros

- #define **DATA_MOVEMENTS_GET_SIZE_TABLES**(interface)
- #define **DATA_MOVEMENTS_GET_TAGS_TABLE**(interface)
- #define **DATA_MOVEMENTS_GET_RANKS_TABLE**(interface)

Functions

- void **data_movements_data_register** (starpu_data_handle_t ∗handle, unsigned home_node, int ∗ranks, starpu_mpi_tag_t ∗tags, int size)
- starpu_mpi_tag_t ∗∗ **data_movements_get_ref_tags_table** (starpu_data_handle_t handle)
- int ∗∗ **data_movements_get_ref_ranks_table** (starpu_data_handle_t handle)
- int **data_movements_reallocate_tables** (starpu_data_handle_t handle, int size)
- starpu_mpi_tag_t ∗ **data_movements_get_tags_table** (starpu_data_handle_t handle)
- int ∗ **data_movements_get_ranks_table** (starpu_data_handle_t handle)
- int **data_movements_get_size_tables** (starpu_data_handle_t handle)

7.20.1 Data Structure Documentation

7.20.1.1 struct data_movements_interface

interface for data_movements

Data Fields

starpu_mpi_tag_t ∗	tags	Data tags table
int ∗	ranks	Ranks table (where to move the corresponding data)
int	size	Size of the tables

Chapter 8

StarPU Resource Manager File Documentation

8.1 starpurm_private.h File Reference

Data Structures

- struct s_starpurm

Enumerations

- enum **e_state** { **state_uninitialized**, **state_init** }
- enum **e_starpurm_unit_type** {
  starpurm_unit_cpu, **starpurm_unit_opencl**, **starpurm_unit_cuda**, **starpurm_unit_mic**,
  starpurm_unit_ntypes }

8.1.1 Data Structure Documentation

8.1.1.1 struct s_starpurm

Data Fields

hwloc_topology_t	topology	Machine topology as detected by hwloc.
unsigned	max_ncpus	Current upper bound on the number of CPU cores selectable for computing with the runtime system.
unsigned	selected_ncpus	Number of currently selected CPU workers
unsigned	selected_nworkers	Number of currently selected workers (CPU+devices)
int	state	Initialization state of the RM instance.
int	dynamic_resource_sharing	Boolean indicating the state of the dynamic resource sharing layer. !0 indicates that dynamic resource sharing is enabled. 0 indicates that dynamic resource sharing is disabled.
unsigned	sched_ctx_id	Id of the StarPU's sched_ctx used by the RM instance.
int	unit_ntypes	Number of unit types supported by this RM instance.
int ∗	nunits_by_type	Number of unitss available for each type.
int	nunits	Number of units.
int ∗	unit_offsets_by_type	Offset of unit numbering for each type.
struct s_starpurm_unit ∗	units	Array of units.

Data Fields

hwloc_cpuset_t	global_cpuset	Cpuset of all the StarPU's workers (CPU+devices.
hwloc_cpuset_t	all_cpu_workers_cpuset	Cpuset of all StarPU CPU workers.
hwloc_cpuset_t	all_opencl_device_workers_cpuset	Cpuset of all StarPU OpenCL workers.
hwloc_cpuset_t	all_cuda_device_workers_cpuset	Cpuset of all StarPU CUDA workers.
hwloc_cpuset_t	all_mic_device_workers_cpuset	Cpuset of all StarPU MIC workers.
hwloc_cpuset_t	all_device_workers_cpuset	Cpuset of all StarPU device workers.
hwloc_cpuset_t	selected_cpuset	Cpuset of all selected workers (CPU+devices).
hwloc_cpuset_t	initially_owned_cpuset_mask	Cpuset mask of initially owned cpuset or full if not used.
int	max_worker_id	maximum value among worker ids
int ∗	worker_unit_ids	worker id to unit id table
unsigned int	max_temporary_ctxs	Temporary contexts accounting.
unsigned int	avail_temporary_ctxs	
pthread_mutex_t	temporary_ctxs_mutex	
pthread_cond_t	temporary_ctxs_cond	
int	starpu_in_pause	Global StarPU pause state
pthread_t	event_thread	Event list.
pthread_mutex_t	event_list_mutex	
pthread_cond_t	event_list_cond	
pthread_cond_t	event_processing_cond	
int	event_processing_enabled	
int	event_processing_ended	
struct s_starpurm_event ∗	event_list_head	
struct s_starpurm_event ∗	event_list_tail	